



CS545 – Machine Learning for Signal Processing

Graphs in Signal Processing and Machine Learning

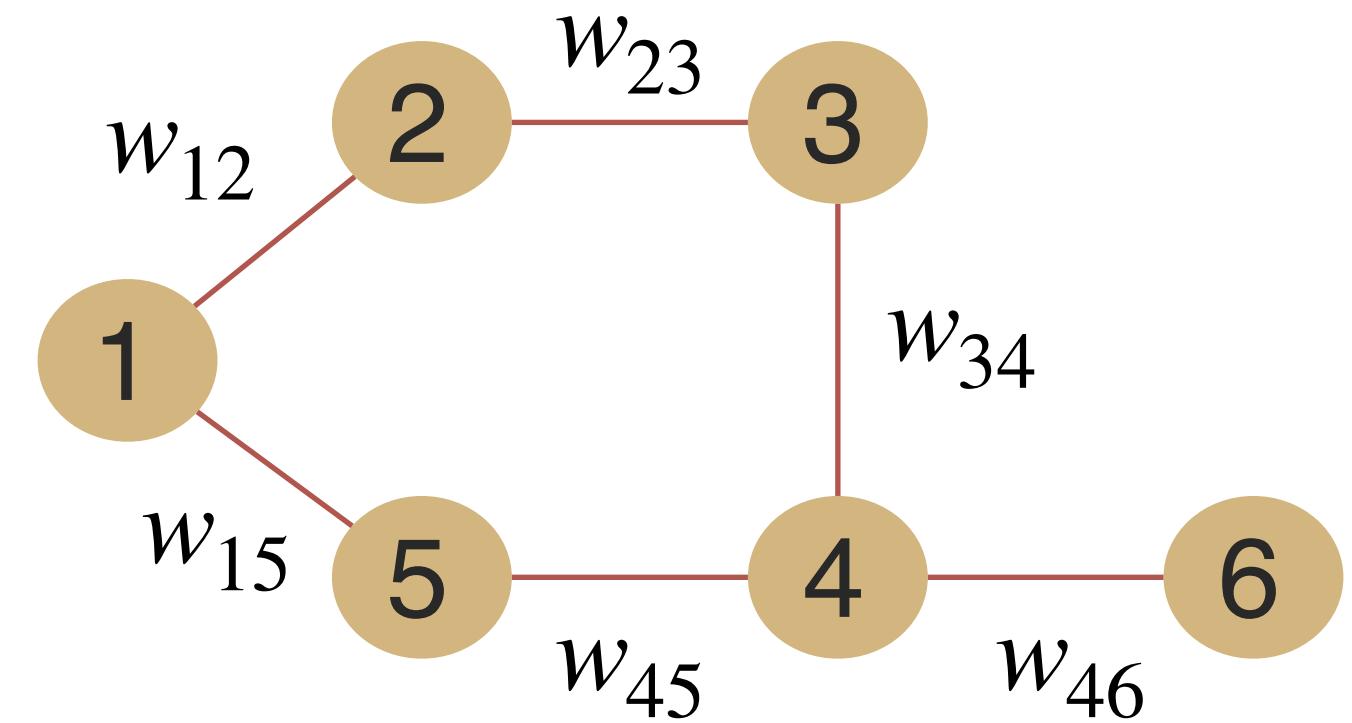
November 13 2023

Today's lecture

- Using graphs instead of vectors/matrices
- Graph signal processing basics
- Graph machine learning

A Graph

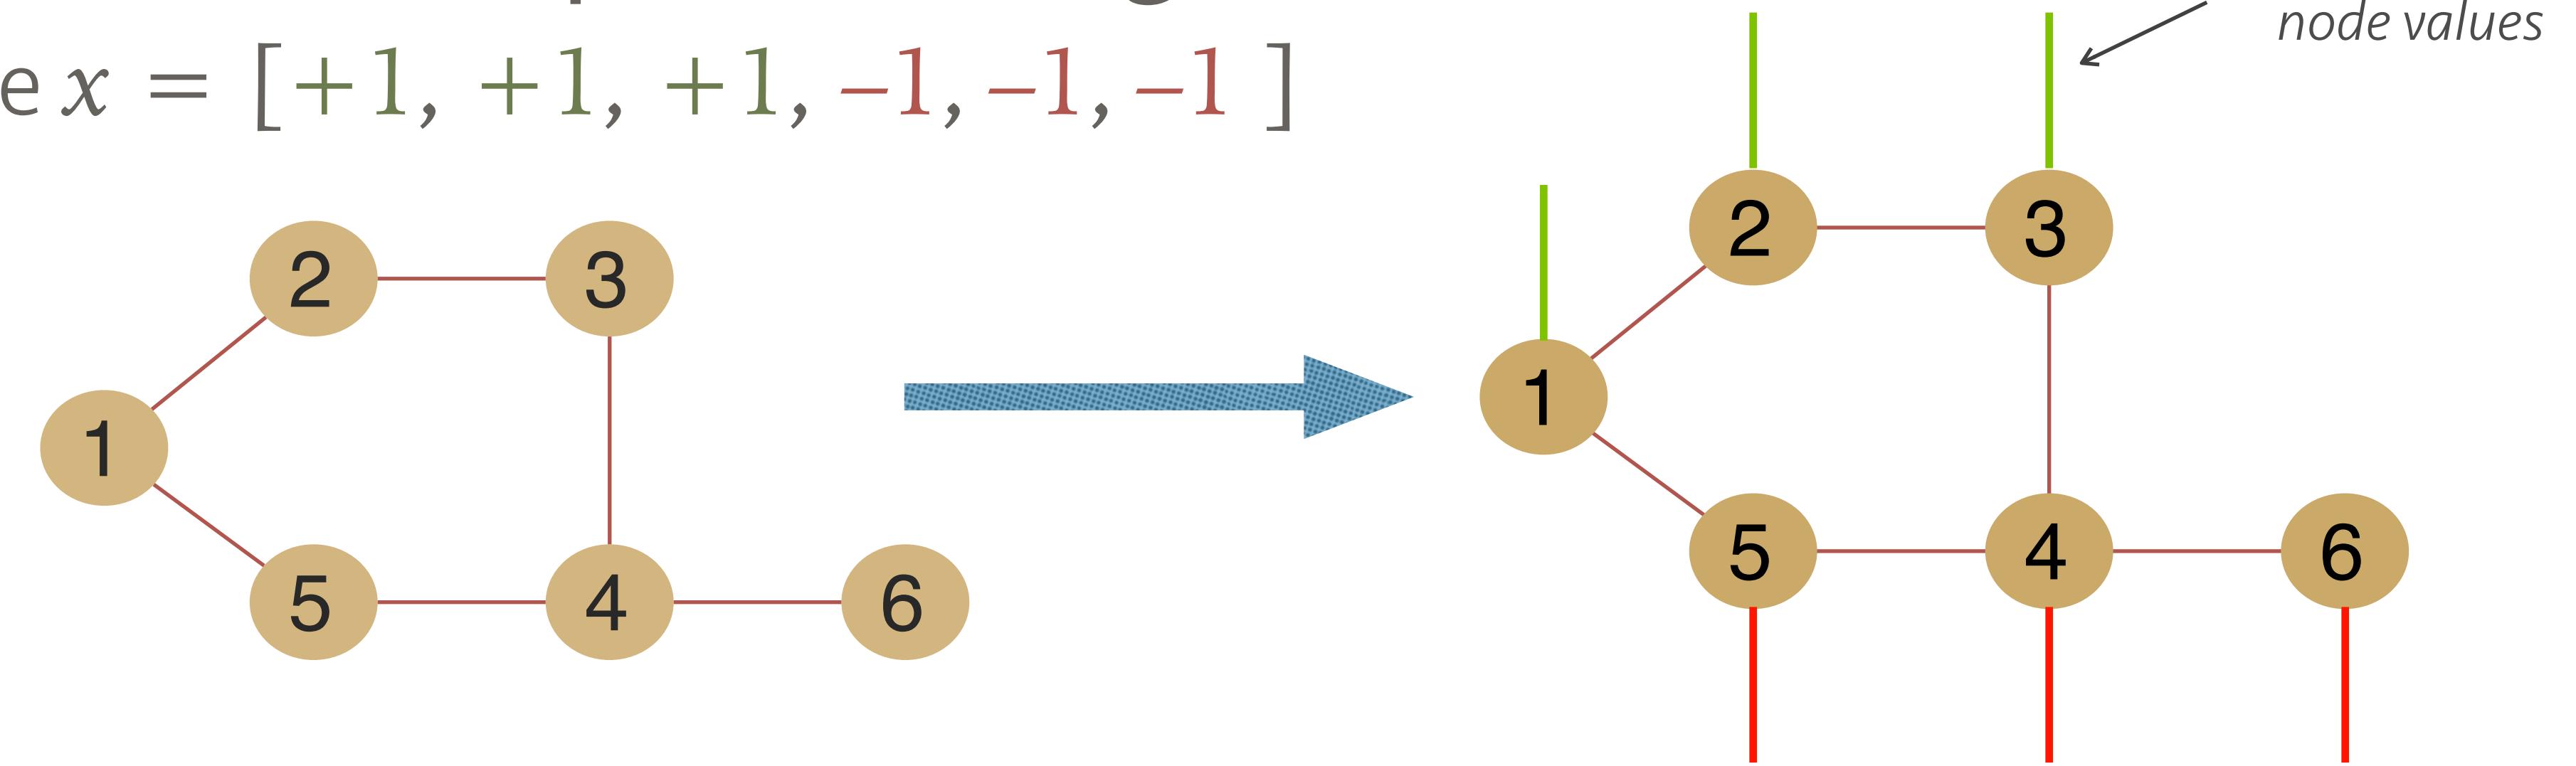
- A Graph is a set of vertices (nodes) and edges
- Each node has a feature or *message*
 - Messages move through nodes to distribute information
- Edges may have associated weights
 - Messages are weighted as they move through the edge
- Graphs can be directed or undirected
 - *Undirected*: $1 \rightarrow 2$ is equivalent to $2 \rightarrow 1$
 - *Directed*: $1 \rightarrow 2$ and $2 \rightarrow 1$ have different weights
- Good for representing pairwise relationships



$\begin{matrix} \text{Graph} \\ \downarrow \\ G = \{V, E\} \\ \uparrow \\ \text{Vertices or Nodes} \end{matrix}$

Representing a signal as a graph

- Associate each sample of the signal to a node
 - Example $x = [+1, +1, +1, -1, -1, -1]$



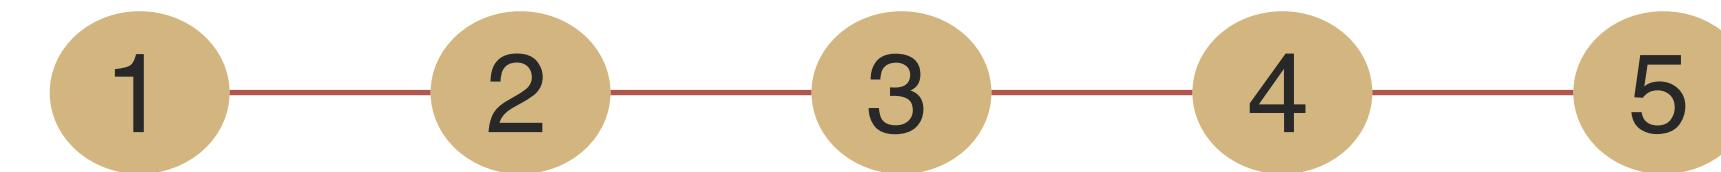
- The signal sample values become node values or features
 - The sample relationships come from the graph
- Graphs are a good representation of irregular signals
 - Edges give us structure, nodes hold samples

Graph Representations

- Symbolically, $G = \{V, E\}$ completely defines a graph
 - But it's an inconvenient representation for us
- Representing a graph as a familiar data structure
 - Use an *adjacency matrix* \mathbf{W} of size $N \times N$
 - N is the number of vertices
 - The elements of \mathbf{W} give the weights of the edges
 - Undirected graphs have a symmetric adjacency matrix
 - Pack all node values into an N -dimensional vector \mathbf{v}
 - i -th's node value is in v_i

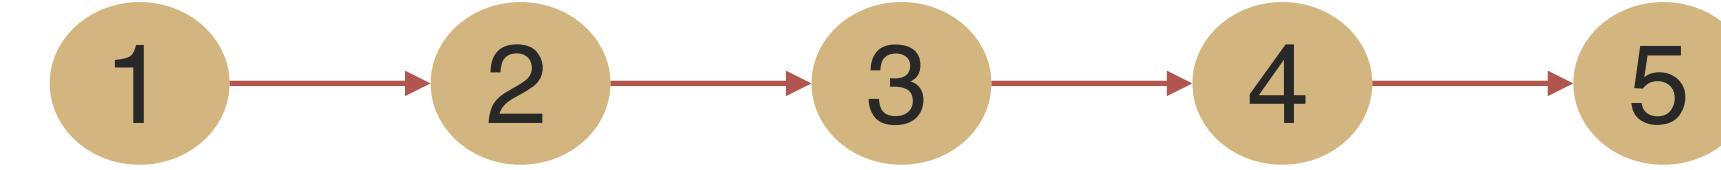
Examples

Undirected Path Graph



$$\mathbf{W} = \begin{bmatrix} & & & & & \text{to} \\ 0 & 1 & 0 & 0 & 0 & \\ 1 & 0 & 1 & 0 & 0 & \\ 0 & 1 & 0 & 1 & 0 & \\ 0 & 0 & 1 & 0 & 1 & \\ 0 & 0 & 0 & 1 & 0 & \\ & & & & & \text{from} \end{bmatrix}$$

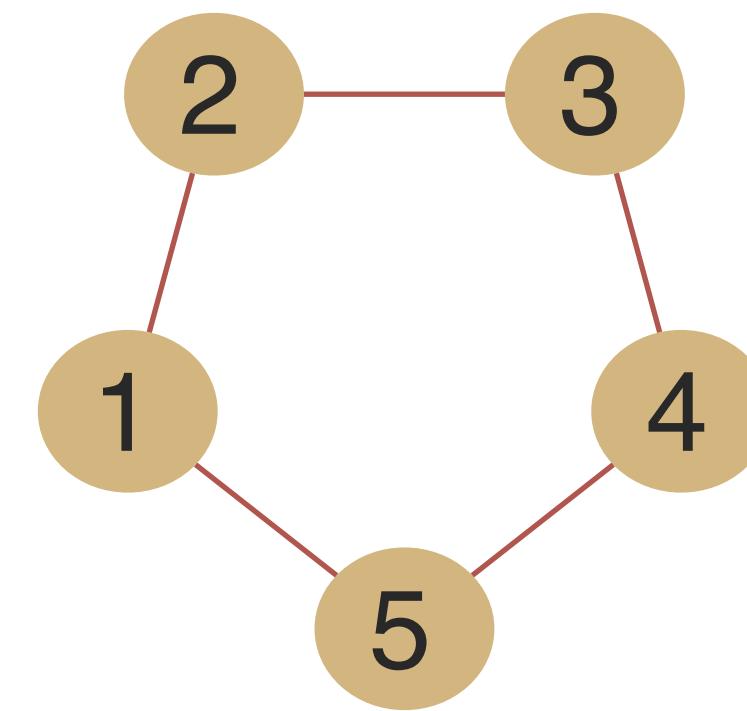
Directed Path Graph



$$\mathbf{W} = \begin{bmatrix} & & & & & \text{to} \\ 0 & 1 & 0 & 0 & 0 & \\ 0 & 0 & 1 & 0 & 0 & \\ 0 & 0 & 0 & 1 & 0 & \\ 0 & 0 & 0 & 0 & 1 & \\ 0 & 0 & 0 & 0 & 0 & \\ & & & & & \text{from} \end{bmatrix}$$

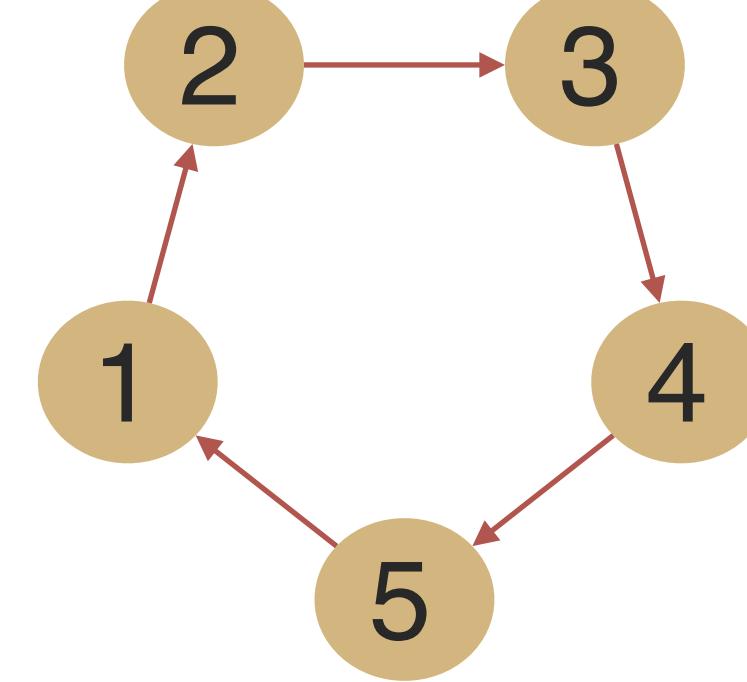
Examples

Undirected Ring Graph



$$\mathbf{W} = \begin{bmatrix} 0 & 1 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 & 0 \\ 0 & 1 & 0 & 1 & 0 \\ 0 & 0 & 1 & 0 & 1 \\ 1 & 0 & 0 & 1 & 0 \end{bmatrix}$$

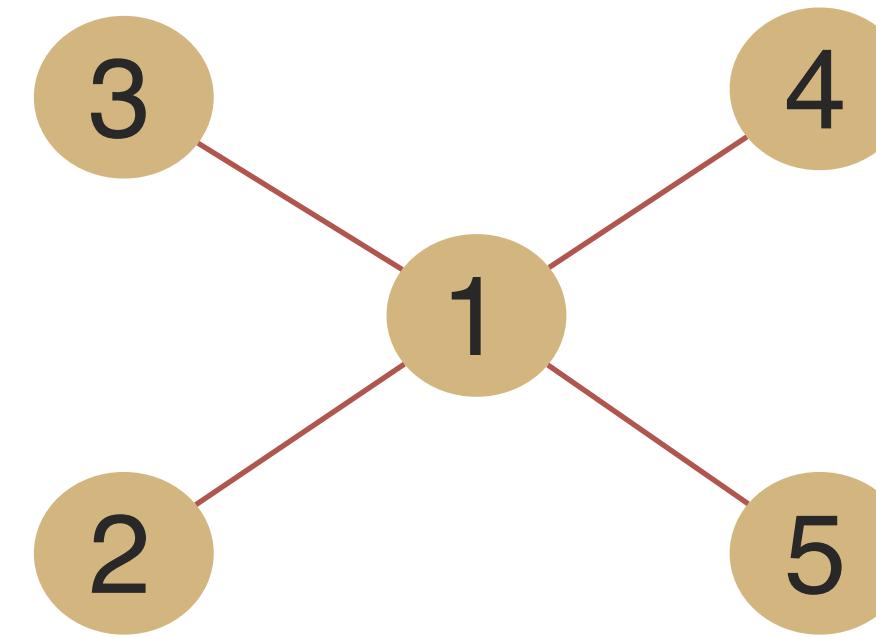
Directed Ring Graph



$$\mathbf{W} = \begin{bmatrix} 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 & 0 \end{bmatrix}$$

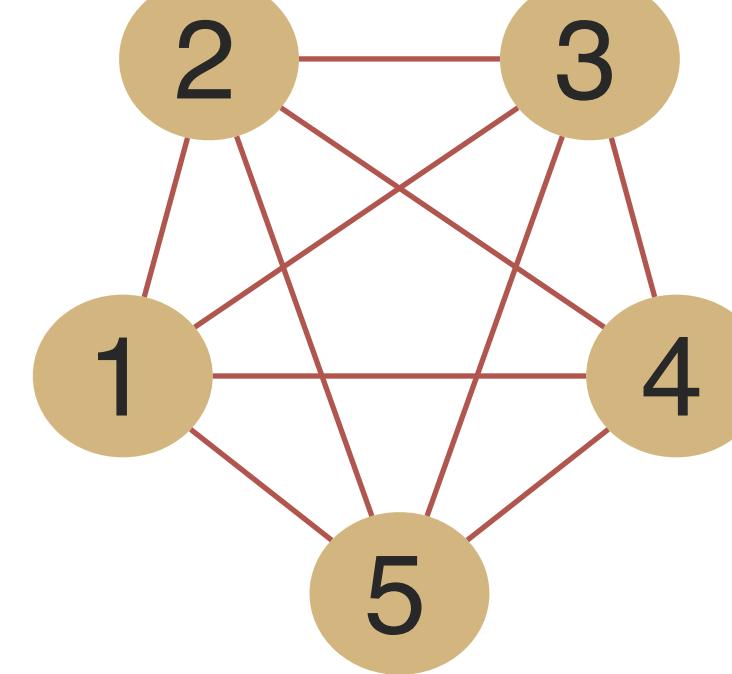
More Examples

Undirected Star Graph



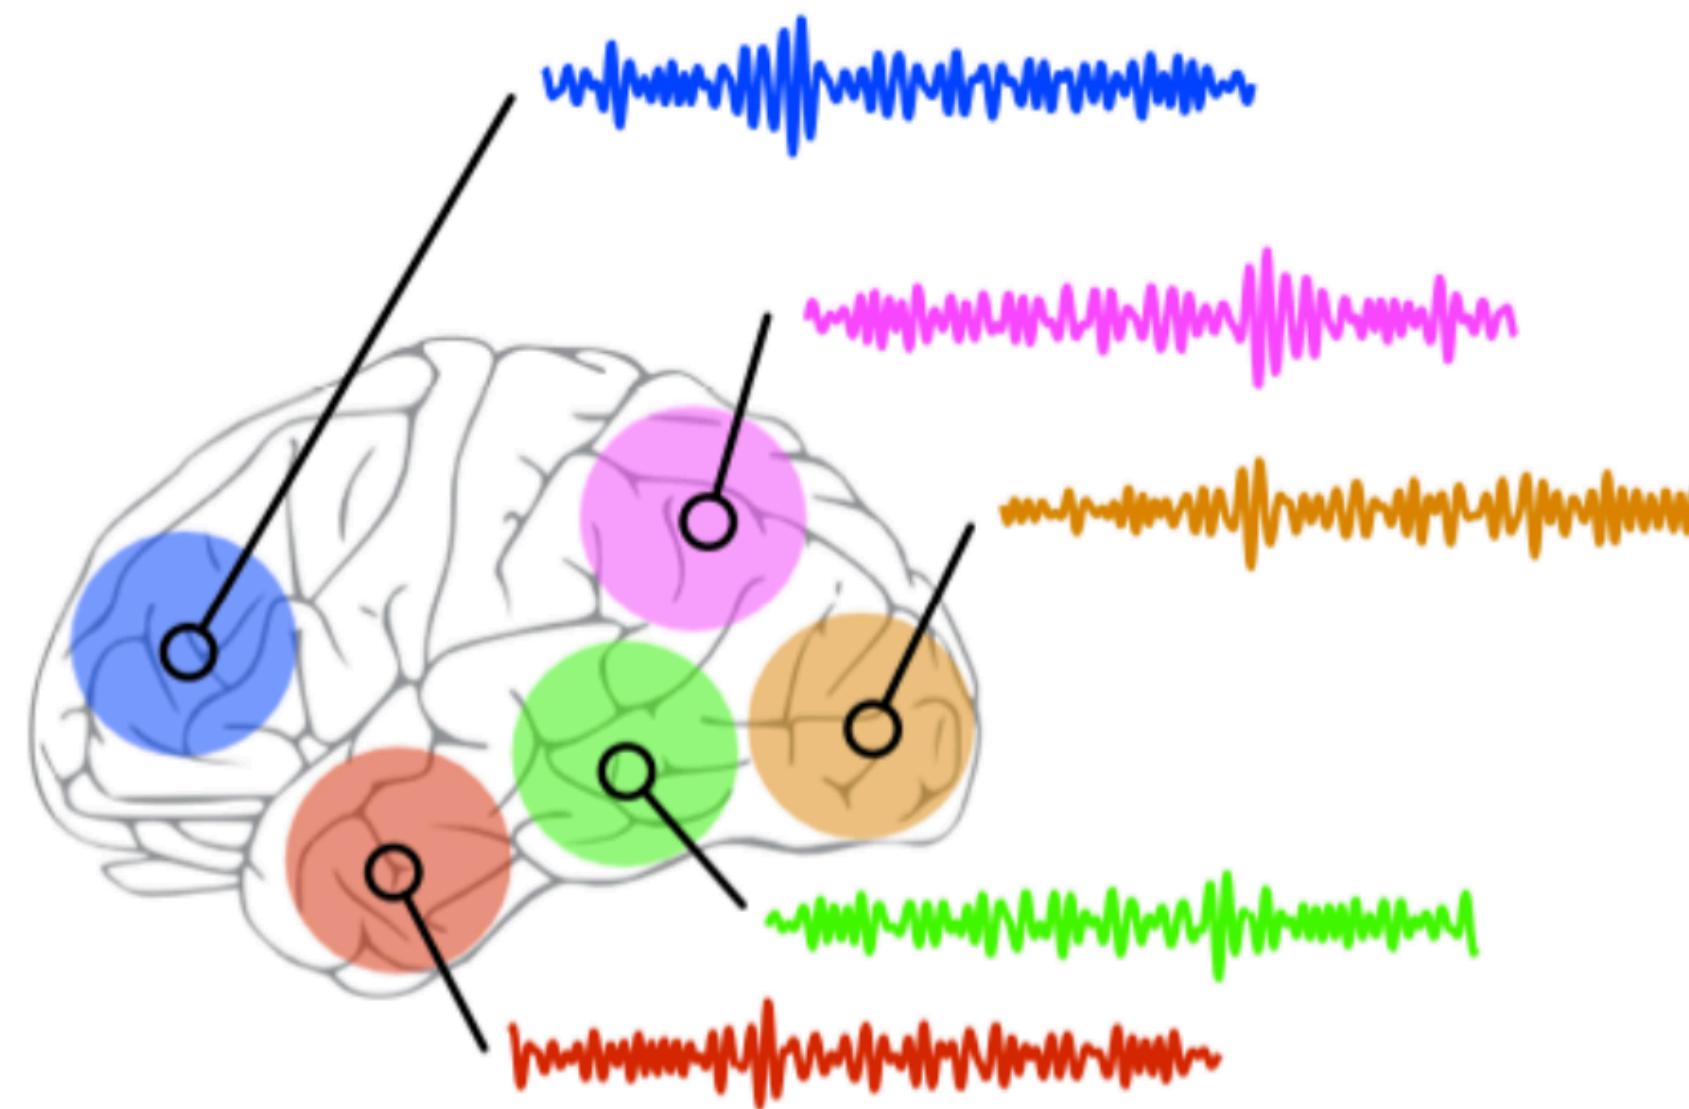
$$\mathbf{W} = \begin{bmatrix} 0 & 1 & 1 & 1 & 1 \\ 1 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 \end{bmatrix}$$

Undirected Fully Connected Graph



$$\mathbf{W} = \begin{bmatrix} 0 & 1 & 1 & 1 & 1 \\ 1 & 0 & 1 & 1 & 1 \\ 1 & 1 & 0 & 1 & 1 \\ 1 & 1 & 1 & 0 & 1 \\ 1 & 1 & 1 & 1 & 0 \end{bmatrix}$$

Real-world Examples

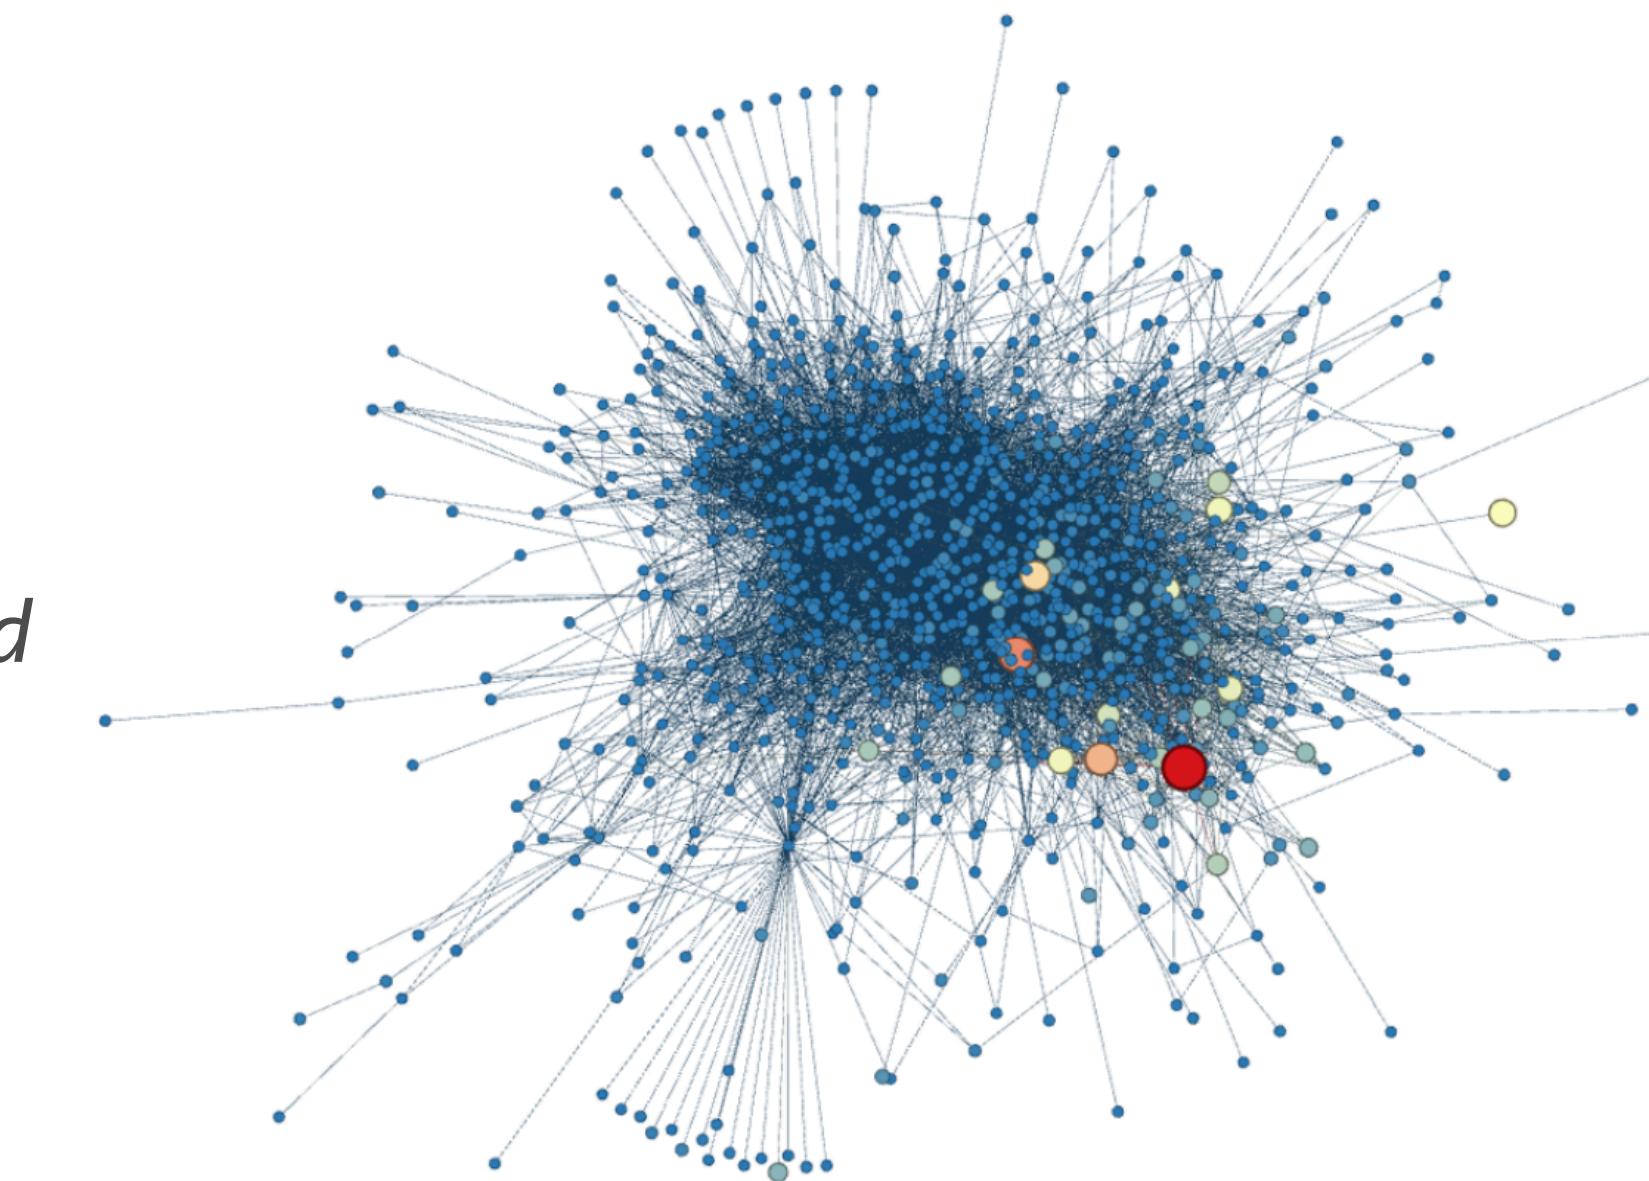


Social Network Graph

Vertices: 1,000 Twitter users

Edges: Relationship between users

Signal: Apple-related hash-tags in a 6 week period



Brain Imaging

Vertices: Different brain regions

Edges: Structural connectivity between regions

Signal: Blood-oxygen-level series

Graph Signal Processing

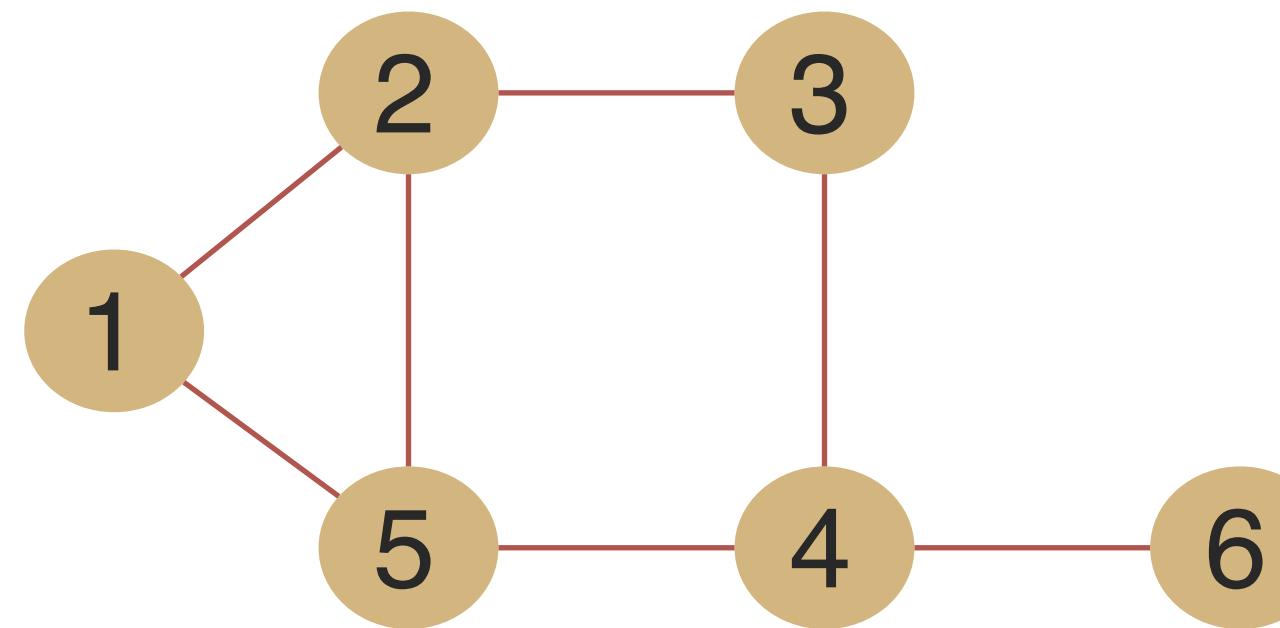
- How do we process such representations?
 - *Graph Signal Processing*
- Classical DSP assumes “regular” data
 - GSP generalizes classical DSP to graph signals
- GSP also helps define learning on graphs
 - Make signal processing parameters trainable

New domains to worry about

- Classical signal processing
 - Time domain →
 - Natural representation (e.g. pixels)
 - Simple relationship between samples
 - Frequency domain →
 - Fourier Transforms, Cosine Transforms etc.
 - Signal properties can be easily analyzed
- Graph signal processing
 - Vertex (spatial) domain
 - Described using the *adjacency matrix*
 - Arbitrary sample relations
 - Graph Spectral domain
 - *Graph Fourier Transform*
 - Important for analysis of signal properties
- How do we move from spatial to spectral domain?
 - We need the *Graph Laplacian matrix L*

The Graph Laplacian

- Consider an undirected graph
 - Assume that all the weights are equal to 1

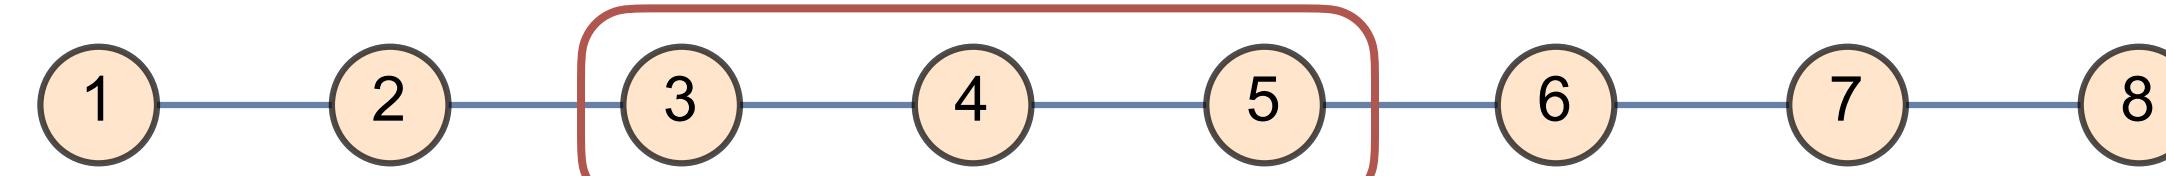


$$\begin{array}{c}
 \text{Graph Laplacian} \\
 \left[\begin{array}{cccccc}
 2 & -1 & 0 & 0 & -1 & 0 \\
 -1 & 3 & -1 & 0 & -1 & 0 \\
 0 & -1 & 2 & -1 & 0 & 0 \\
 0 & 0 & -1 & 3 & -1 & -1 \\
 -1 & -1 & 0 & -1 & 3 & 0 \\
 0 & 0 & 0 & -1 & 0 & 1
 \end{array} \right]
 \end{array}
 =
 \begin{array}{c}
 \text{Degree Matrix} \\
 \left[\begin{array}{cccccc}
 2 & 0 & 0 & 0 & 0 & 0 \\
 0 & 3 & 0 & 0 & 0 & 0 \\
 0 & 0 & 2 & 0 & 0 & 0 \\
 0 & 0 & 0 & 3 & 0 & 0 \\
 0 & 0 & 0 & 0 & 3 & 0 \\
 0 & 0 & 0 & 0 & 0 & 1
 \end{array} \right]
 \end{array}
 -
 \begin{array}{c}
 \text{Graph Adjacency} \\
 \left[\begin{array}{cccccc}
 0 & 1 & 0 & 0 & 1 & 0 \\
 1 & 0 & 1 & 0 & 1 & 0 \\
 0 & 1 & 0 & 1 & 0 & 0 \\
 0 & 0 & 1 & 0 & 1 & 1 \\
 1 & 1 & 0 & 1 & 0 & 0 \\
 0 & 0 & 0 & 1 & 0 & 0
 \end{array} \right]
 \end{array}$$

 L
 $D = \text{diag}(W \cdot \mathbf{1})$
(diag of sum or W's rows)
 W

For a simple graph

- Consider a very simple signal
 - 1d regularly sampled time series



Graph Laplacian

$$\begin{pmatrix} 1 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ -1 & 2 & -1 & 0 & 0 & 0 & 0 & 0 \\ 0 & -1 & 2 & -1 & 0 & 0 & 0 & 0 \\ 0 & 0 & -1 & 2 & -1 & 0 & 0 & 0 \\ 0 & 0 & 0 & -1 & 2 & -1 & 0 & 0 \\ 0 & 0 & 0 & 0 & -1 & 2 & -1 & 0 \\ 0 & 0 & 0 & 0 & 0 & -1 & 2 & -1 \\ 0 & 0 & 0 & 0 & 0 & 0 & -1 & 1 \end{pmatrix}$$

=

Degree Matrix

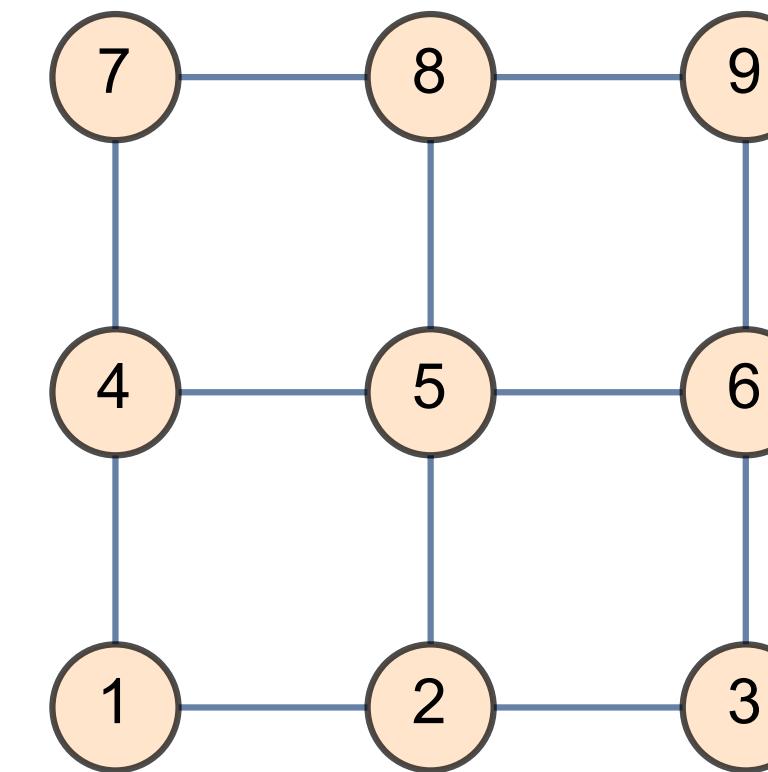
$$\begin{pmatrix} 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 2 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 2 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 2 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 2 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 2 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 2 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \end{pmatrix}$$

Graph Adjacency

$$\begin{pmatrix} 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \end{pmatrix}$$

Which generalizes

- For a regular 2D grid we see something similar
 - It's an edge filter!
 - Technically, the Laplace operator



Graph Laplacian

$$\begin{pmatrix} 2 & -1 & 0 & -1 & 0 & 0 & 0 & 0 & 0 \\ -1 & 3 & -1 & 0 & -1 & 0 & 0 & 0 & 0 \\ 0 & -1 & 2 & 0 & 0 & -1 & 0 & 0 & 0 \\ -1 & 0 & 0 & 3 & -1 & 0 & -1 & 0 & 0 \\ 0 & -1 & 0 & -1 & 4 & -1 & 0 & -1 & 0 \\ 0 & 0 & -1 & 0 & -1 & 3 & 0 & 0 & -1 \\ 0 & 0 & 0 & -1 & 0 & 0 & 2 & -1 & 0 \\ 0 & 0 & 0 & 0 & -1 & 0 & -1 & 3 & -1 \\ 0 & 0 & 0 & 0 & 0 & -1 & 0 & -1 & 2 \end{pmatrix}$$

=

Degree Matrix

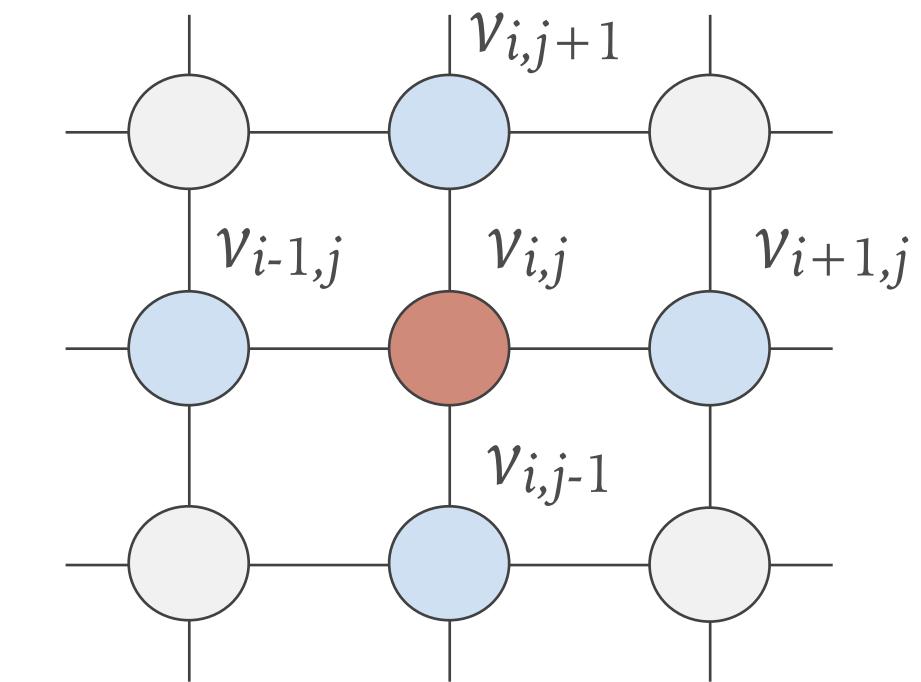
$$\begin{pmatrix} 2 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 3 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 2 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 3 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 4 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 3 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 2 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 3 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 2 \end{pmatrix}$$

Graph Adjacency

$$\begin{pmatrix} 0 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 1 & 0 & 1 & 0 & 1 & 0 & 1 & 0 \\ 0 & 0 & 1 & 0 & 1 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 \end{pmatrix} - \begin{pmatrix} 0 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 1 & 0 & 1 & 0 & 1 & 0 & 1 & 0 \\ 0 & 0 & 1 & 0 & 1 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 \end{pmatrix}$$

The Graph Laplacian: Properties

- Approximates the Laplace operator on the graph
 - Captures local geometry
- For undirected graphs, \mathbf{L} is symmetric
- Off-diagonal entries are non-positive
- Rows add up to zero
 - This is due to the degree matrix \mathbf{D}
- Normalized Laplacian
 - We will stick to using \mathbf{L}



$$\begin{bmatrix} 2 & -1 & 0 & 0 & -1 & 0 \\ -1 & 3 & -1 & 0 & -1 & 0 \\ 0 & -1 & 2 & -1 & 0 & 0 \\ 0 & 0 & -1 & 3 & -1 & -1 \\ -1 & -1 & 0 & -1 & 3 & 0 \\ 0 & 0 & 0 & -1 & 0 & 1 \end{bmatrix}$$

$$\mathbf{L}_{\text{norm}} = \mathbf{D}^{-1/2} \cdot (\mathbf{D} - \mathbf{W}) \cdot \mathbf{D}^{-1/2}$$

The Graph Fourier Transform

- Given a graph containing a signal s
 - Get graph Laplacian L
 - L has a complete set of orthonormal eigenvectors
 - Sorted by the magnitude of the eigenvalues

$$L = V \cdot \Lambda \cdot V^\top$$

Graph equivalent of Fourier bases *Graph Fourier spectral coefficients*

- Transforming the graph signal

$$\xrightarrow{\text{Signal Fourier coefficients}} S = V^\top \cdot s \xleftarrow{\text{Signal spatial coefficients}}$$

Graph Fourier Transform

What does this mean?

- For the k -th eigenvalue/vector pair:

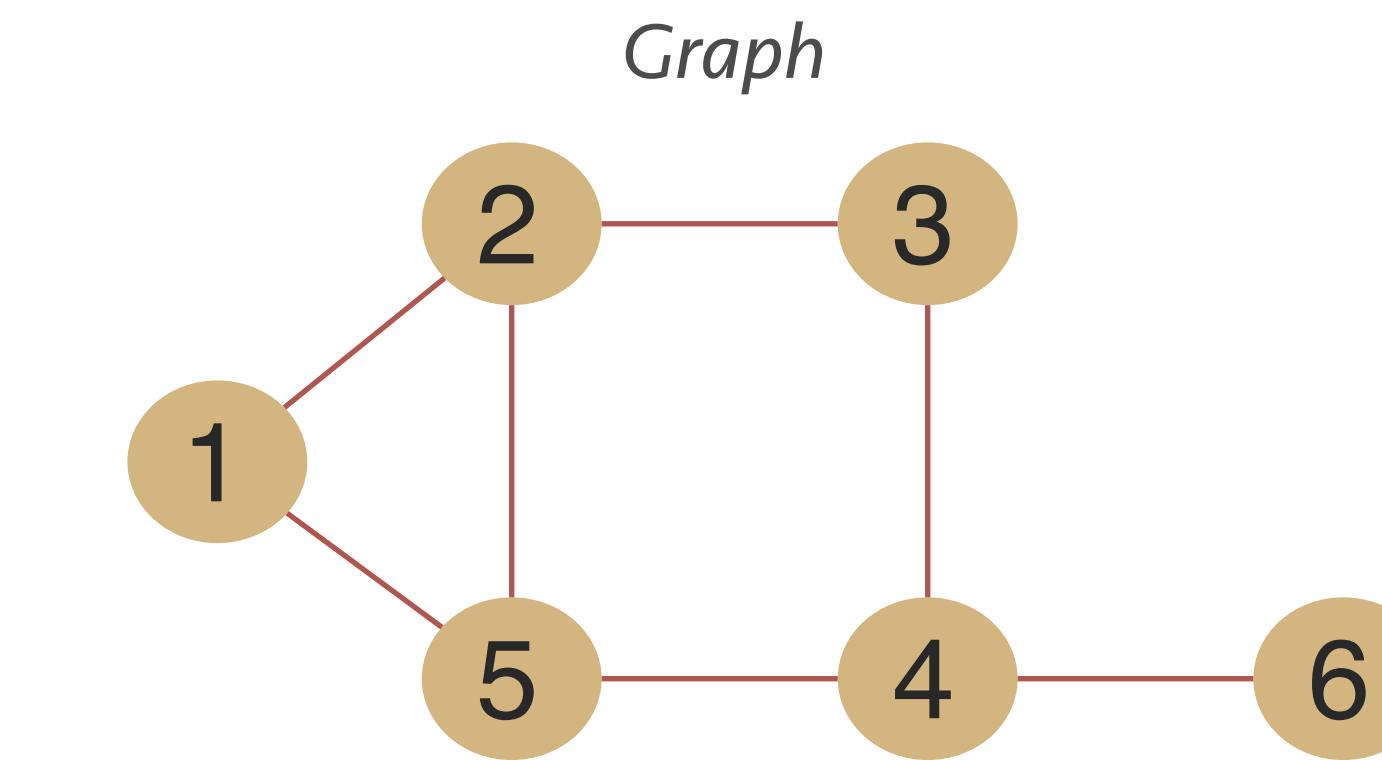
$$\lambda_k = \mathbf{u}_k^\top \cdot \underbrace{\mathbf{L} \cdot \mathbf{u}_k}_{\mathbf{L} \text{ is a difference filter, this product measures local change in } \mathbf{u}} \text{ with } \mathbf{u}_k^\top \cdot \mathbf{u}_k = 1$$



\mathbf{L} is a difference filter, this product measures local change in \mathbf{u}

- Bases \mathbf{u} capture “local variation”
 - “How much” does a signal change in a neighborhood
 - Lower eigenvalue → smoother signal → lower “frequency”
 - \mathbf{u} ordering will be in terms of local variation
 - i.e. amount of high frequency

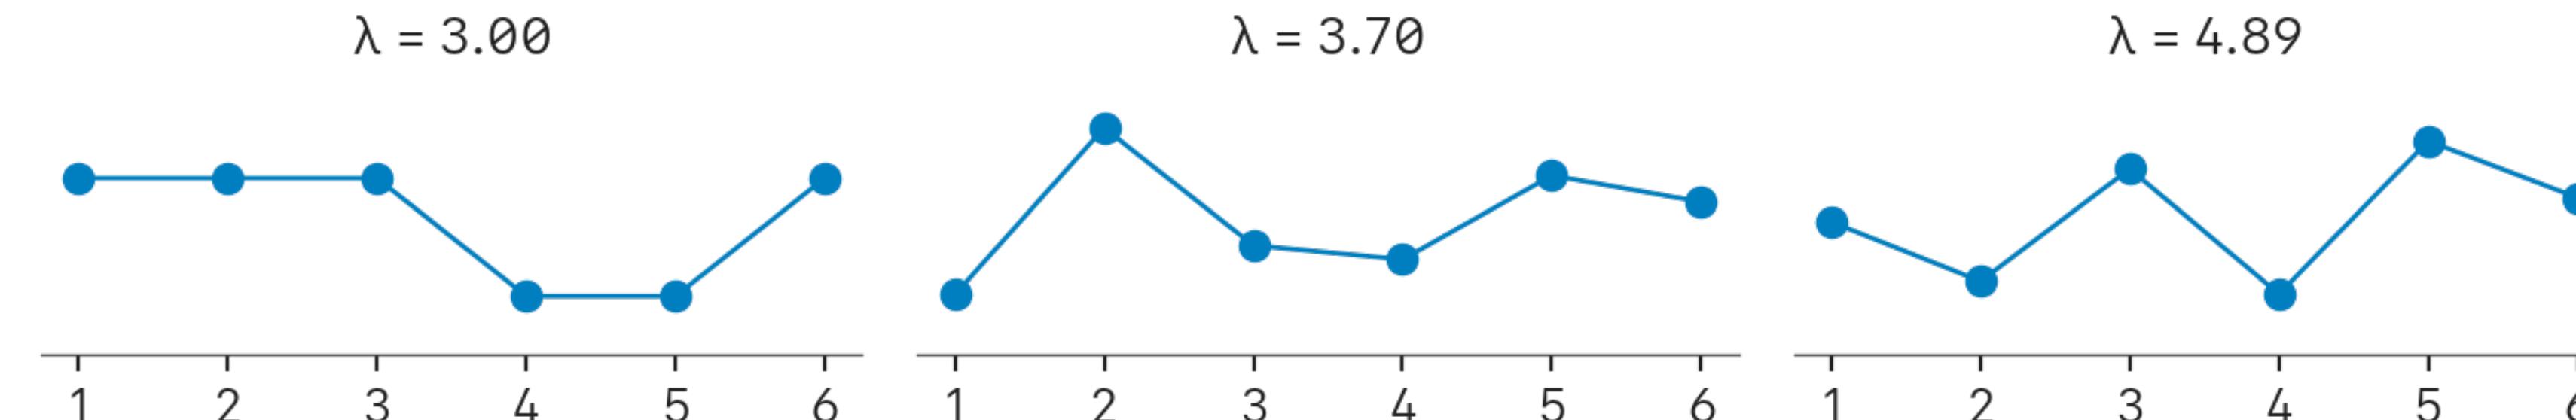
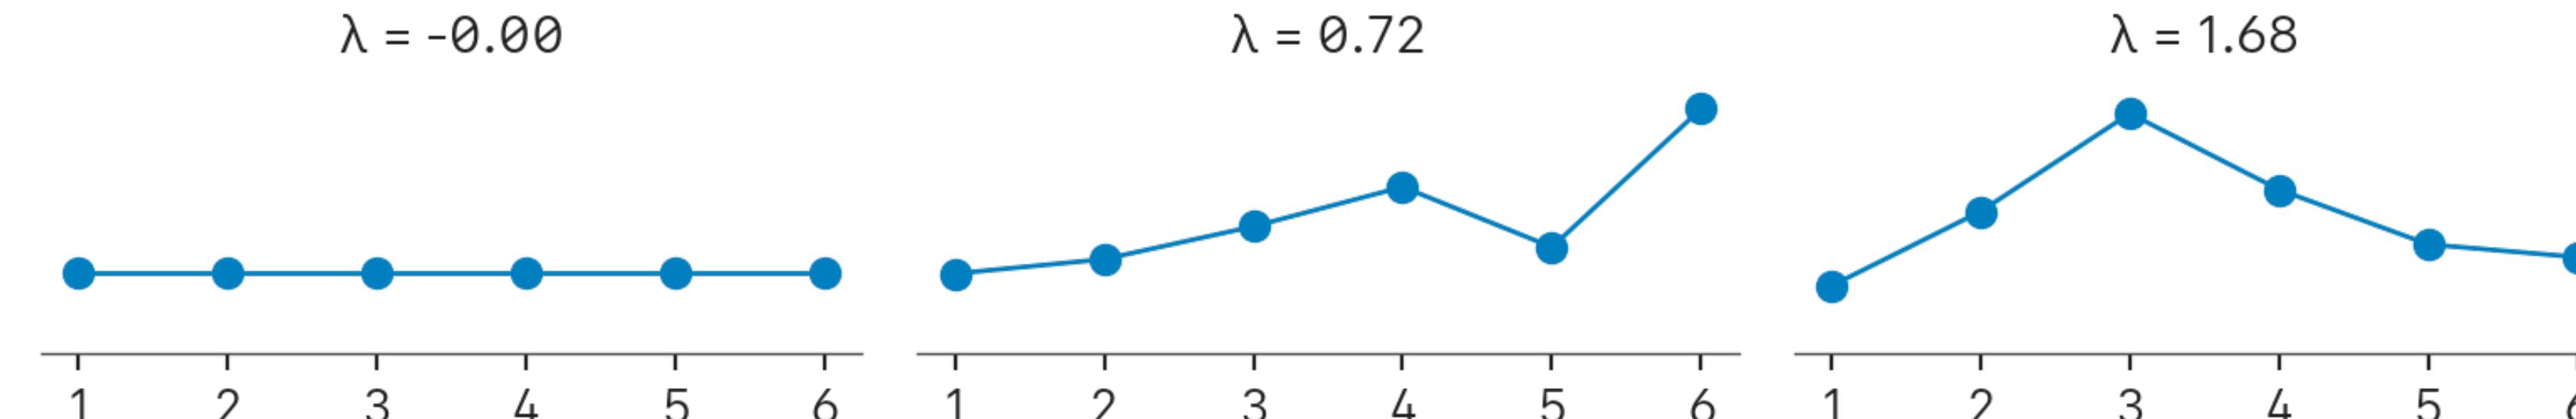
Example



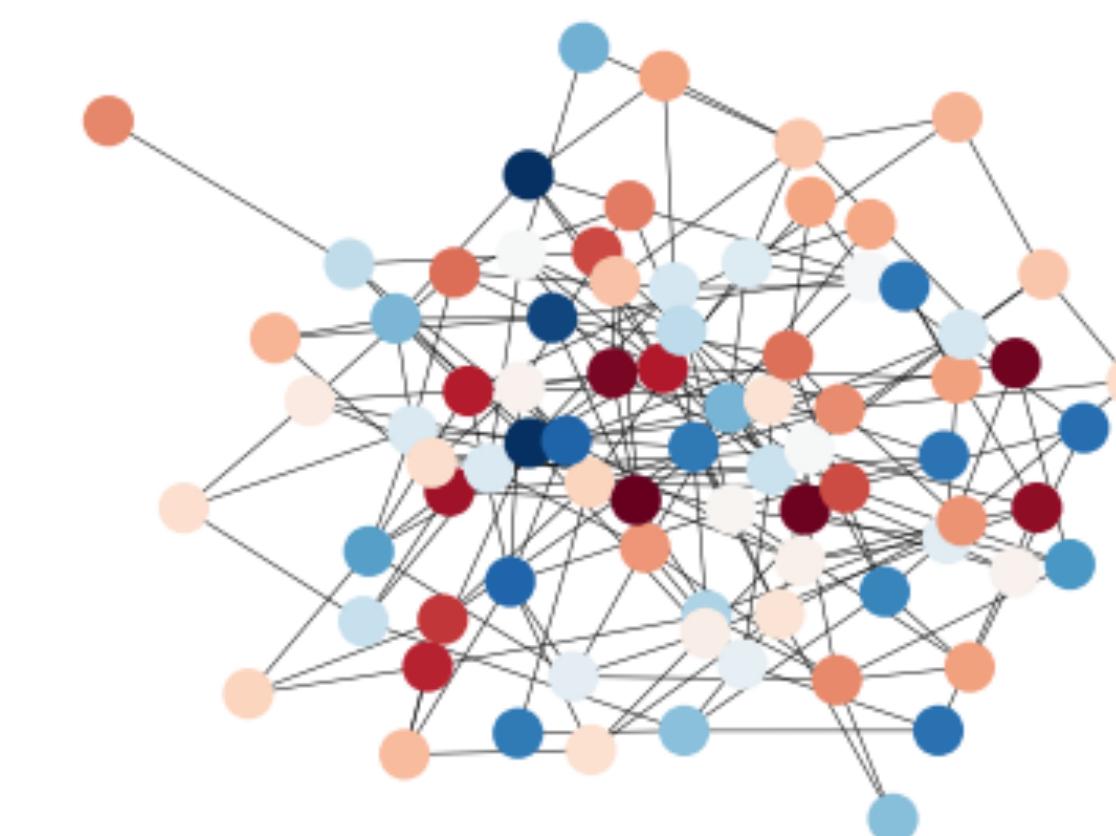
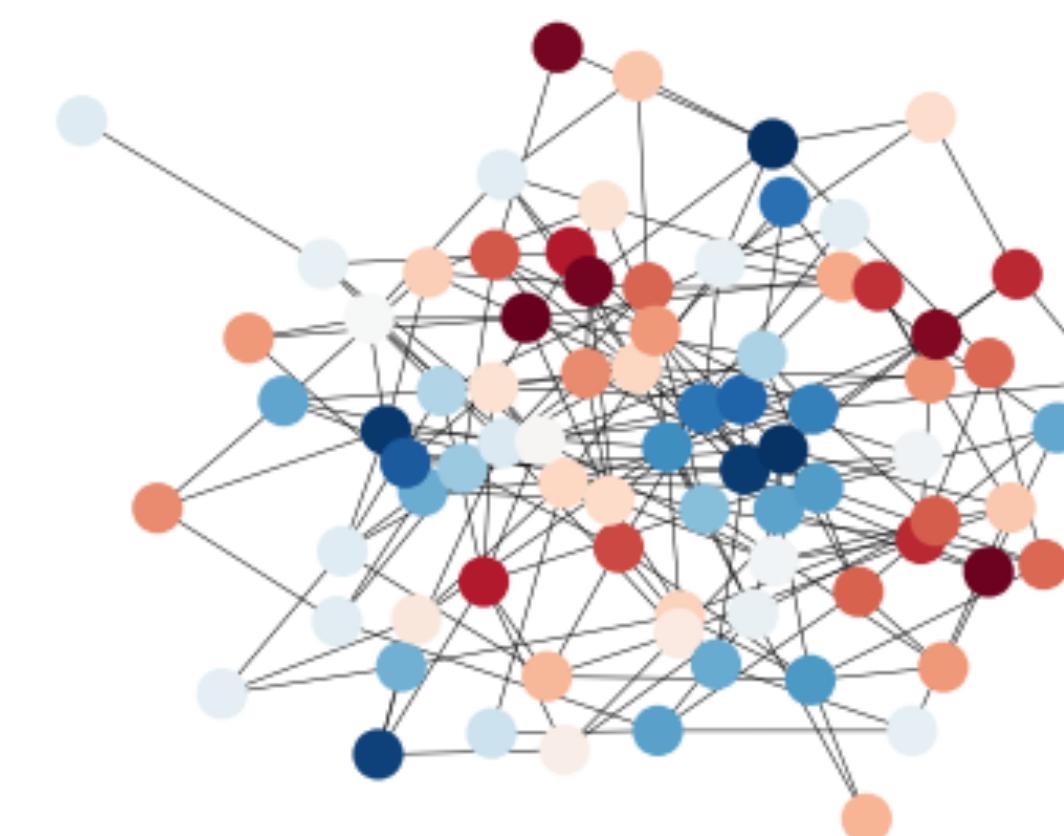
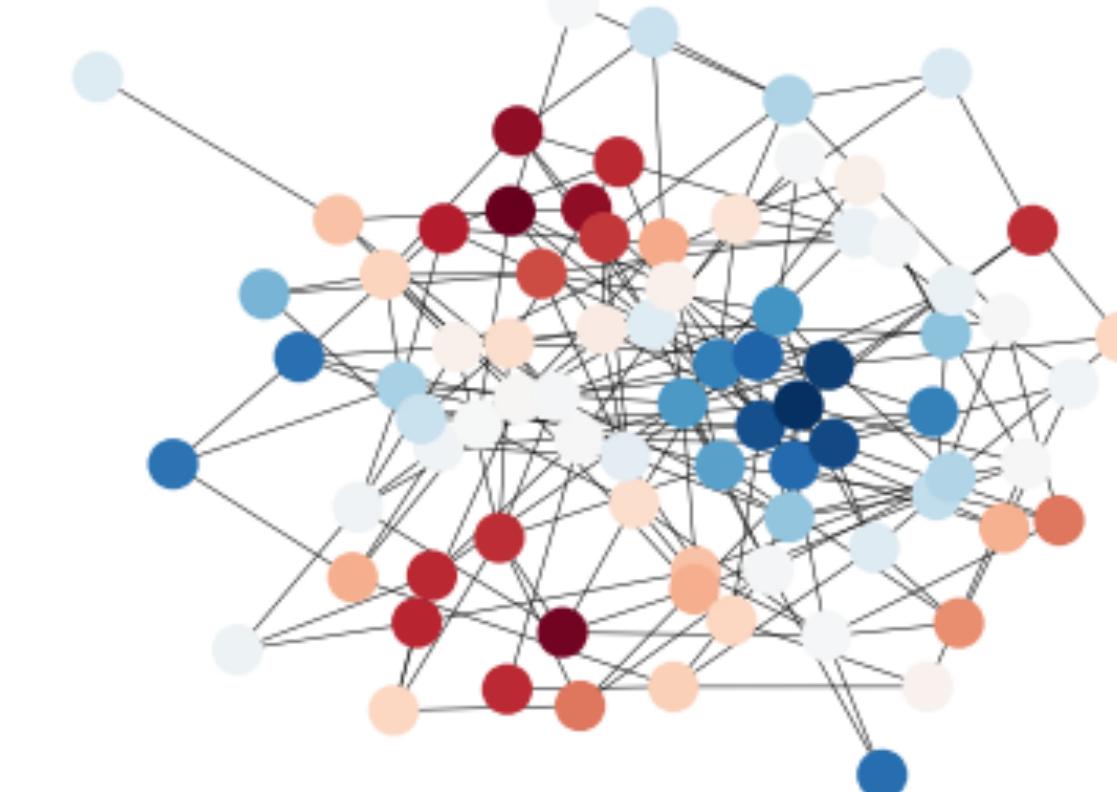
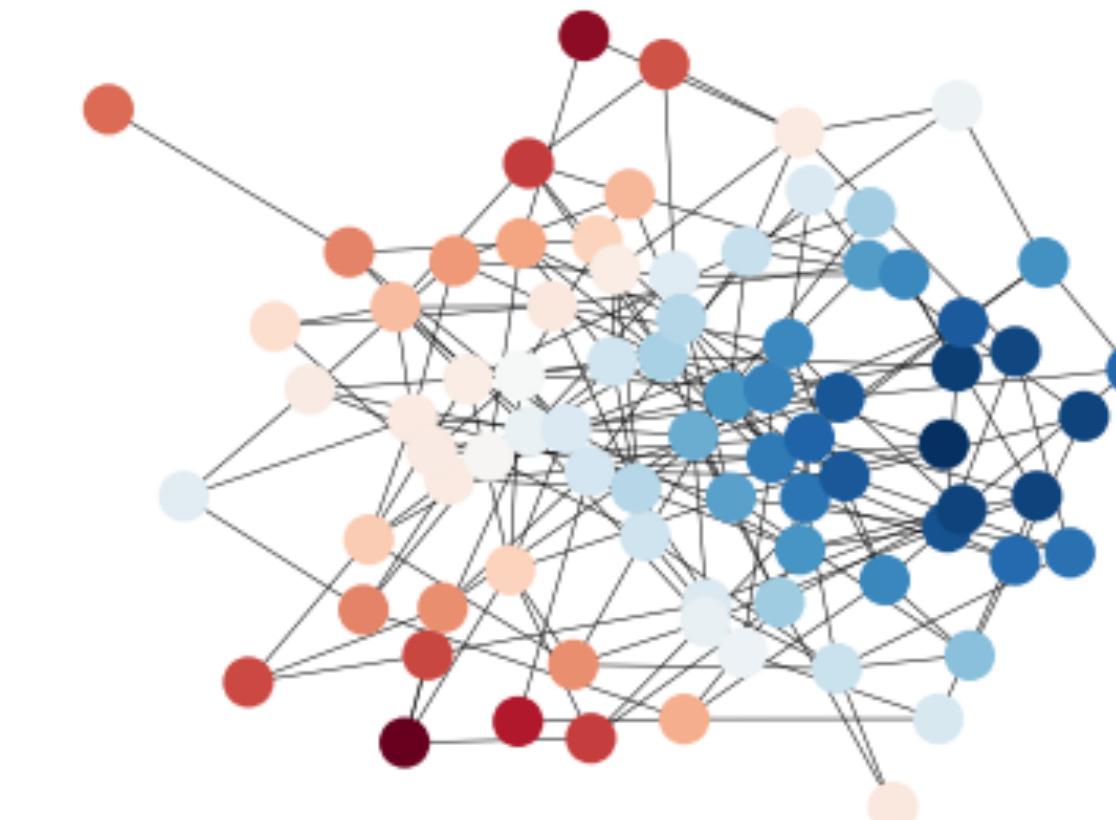
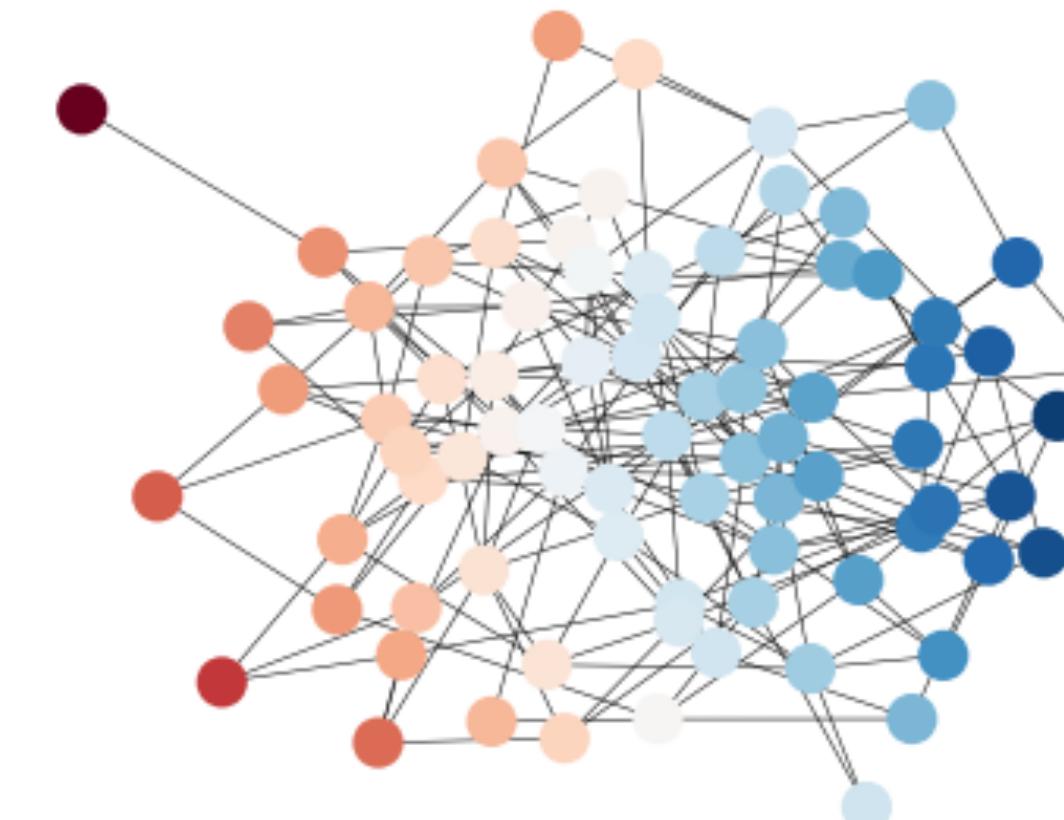
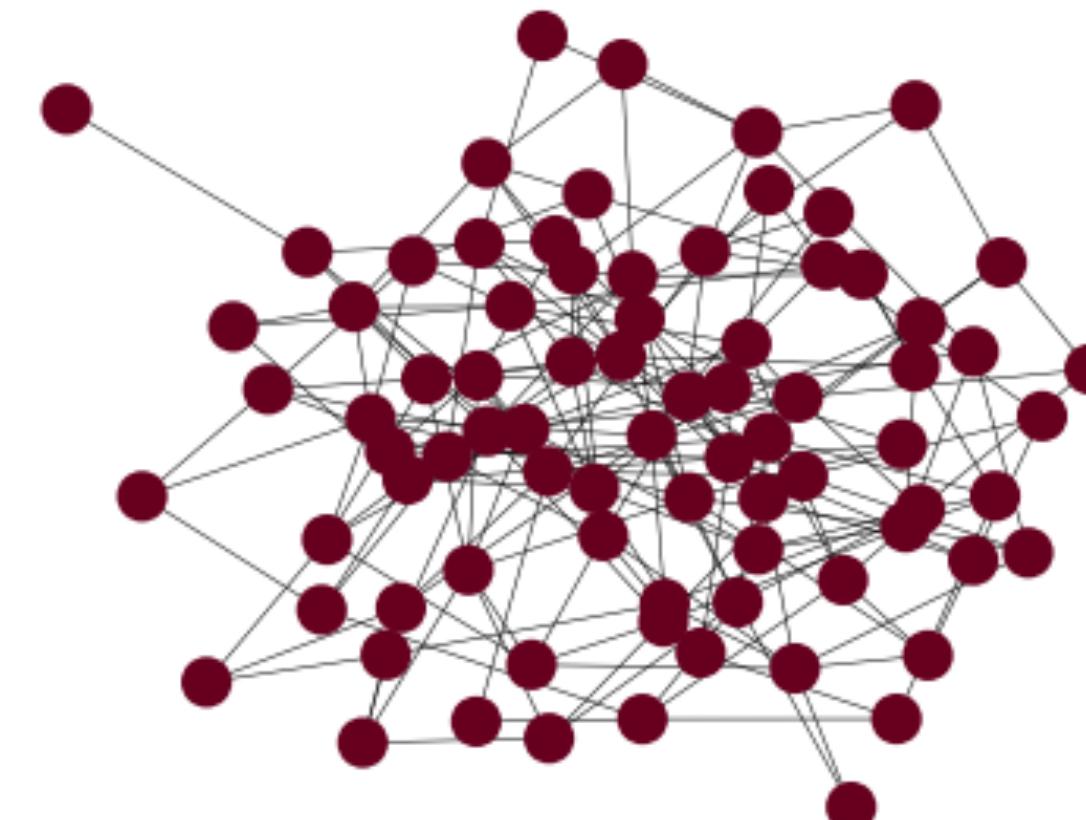
Laplacian

$$\mathbf{L} = \begin{bmatrix} 2 & -1 & 0 & 0 & -1 & 0 \\ -1 & 3 & -1 & 0 & -1 & 0 \\ 0 & -1 & 2 & -1 & 0 & 0 \\ 0 & 0 & -1 & 3 & -1 & -1 \\ -1 & -1 & 0 & -1 & 3 & 0 \\ 0 & 0 & 0 & -1 & 0 & 1 \end{bmatrix}$$

Eigenvectors/eigenvalues

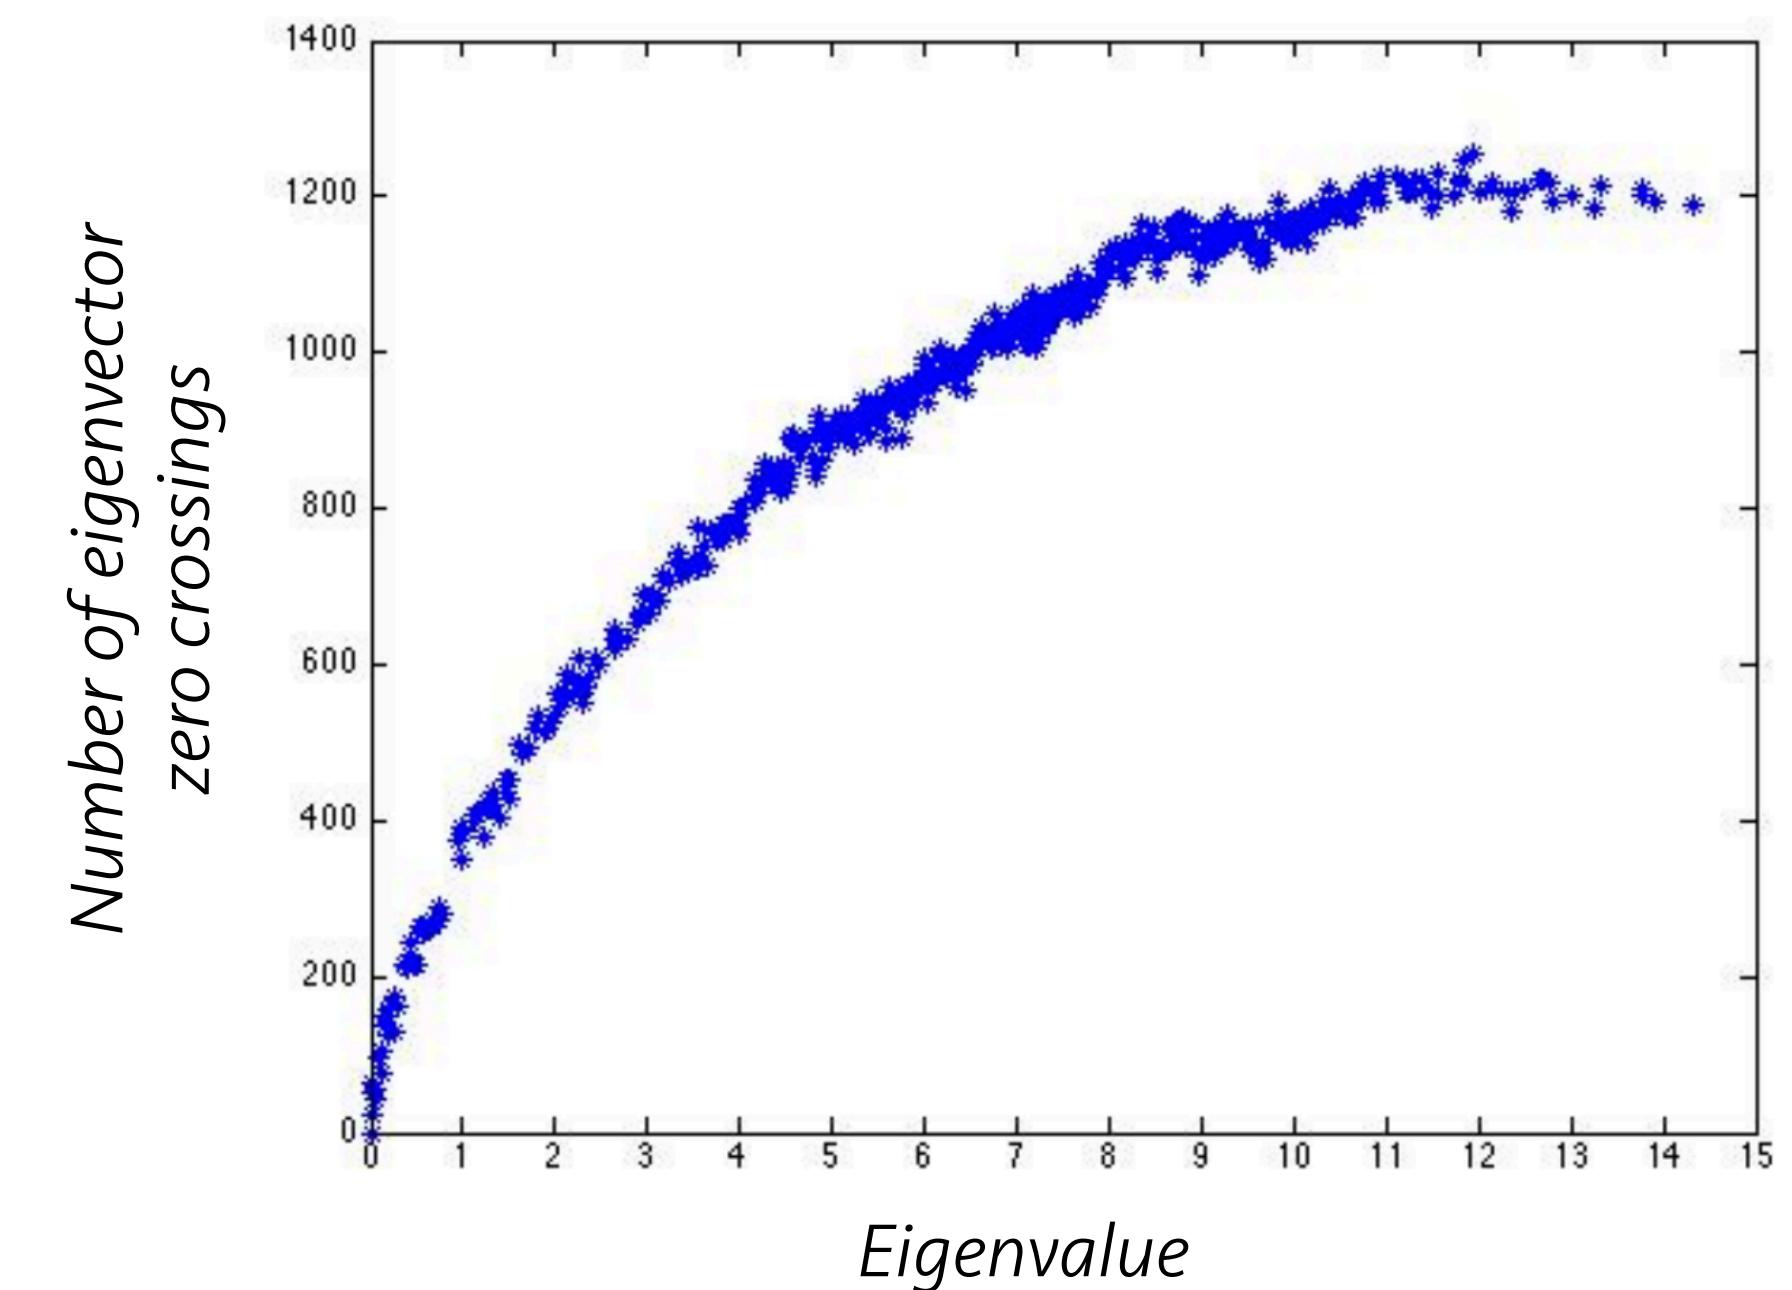


Low and high frequency graph signals



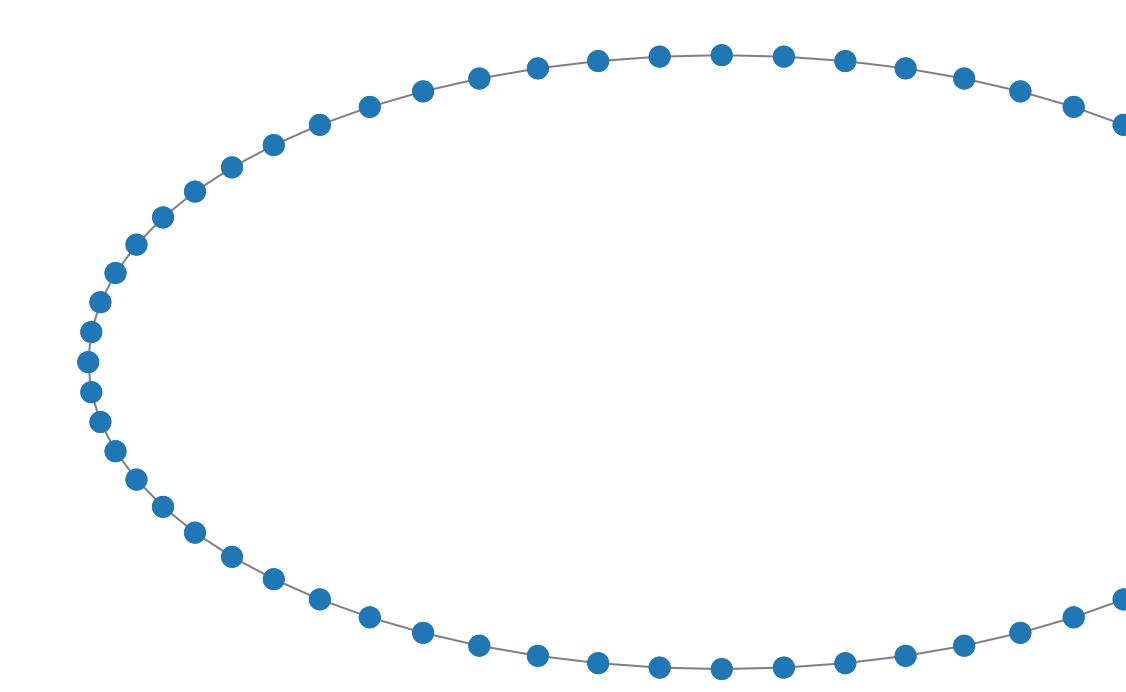
A new notion of “frequency”

- The eigenvalues measure graph “frequency” content
 - How much the corresponding samples change
- Eigenvectors are ranked by their eigenvalues
 - The initial eigenvectors are smooth and change slowly
 - Later eigenvectors change quickly
 - Similar to classical DSP frequency

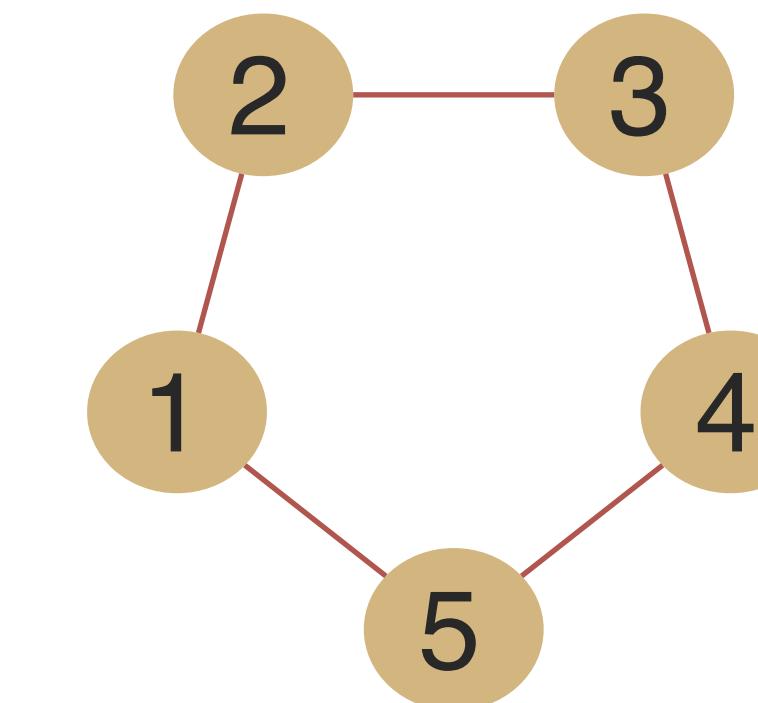


Special Case I: DFT

- Consider an undirected Ring graph
 - Assume all the weights are equal to unity



Ring graph
64 Vertices



Ring graph
5 Vertices

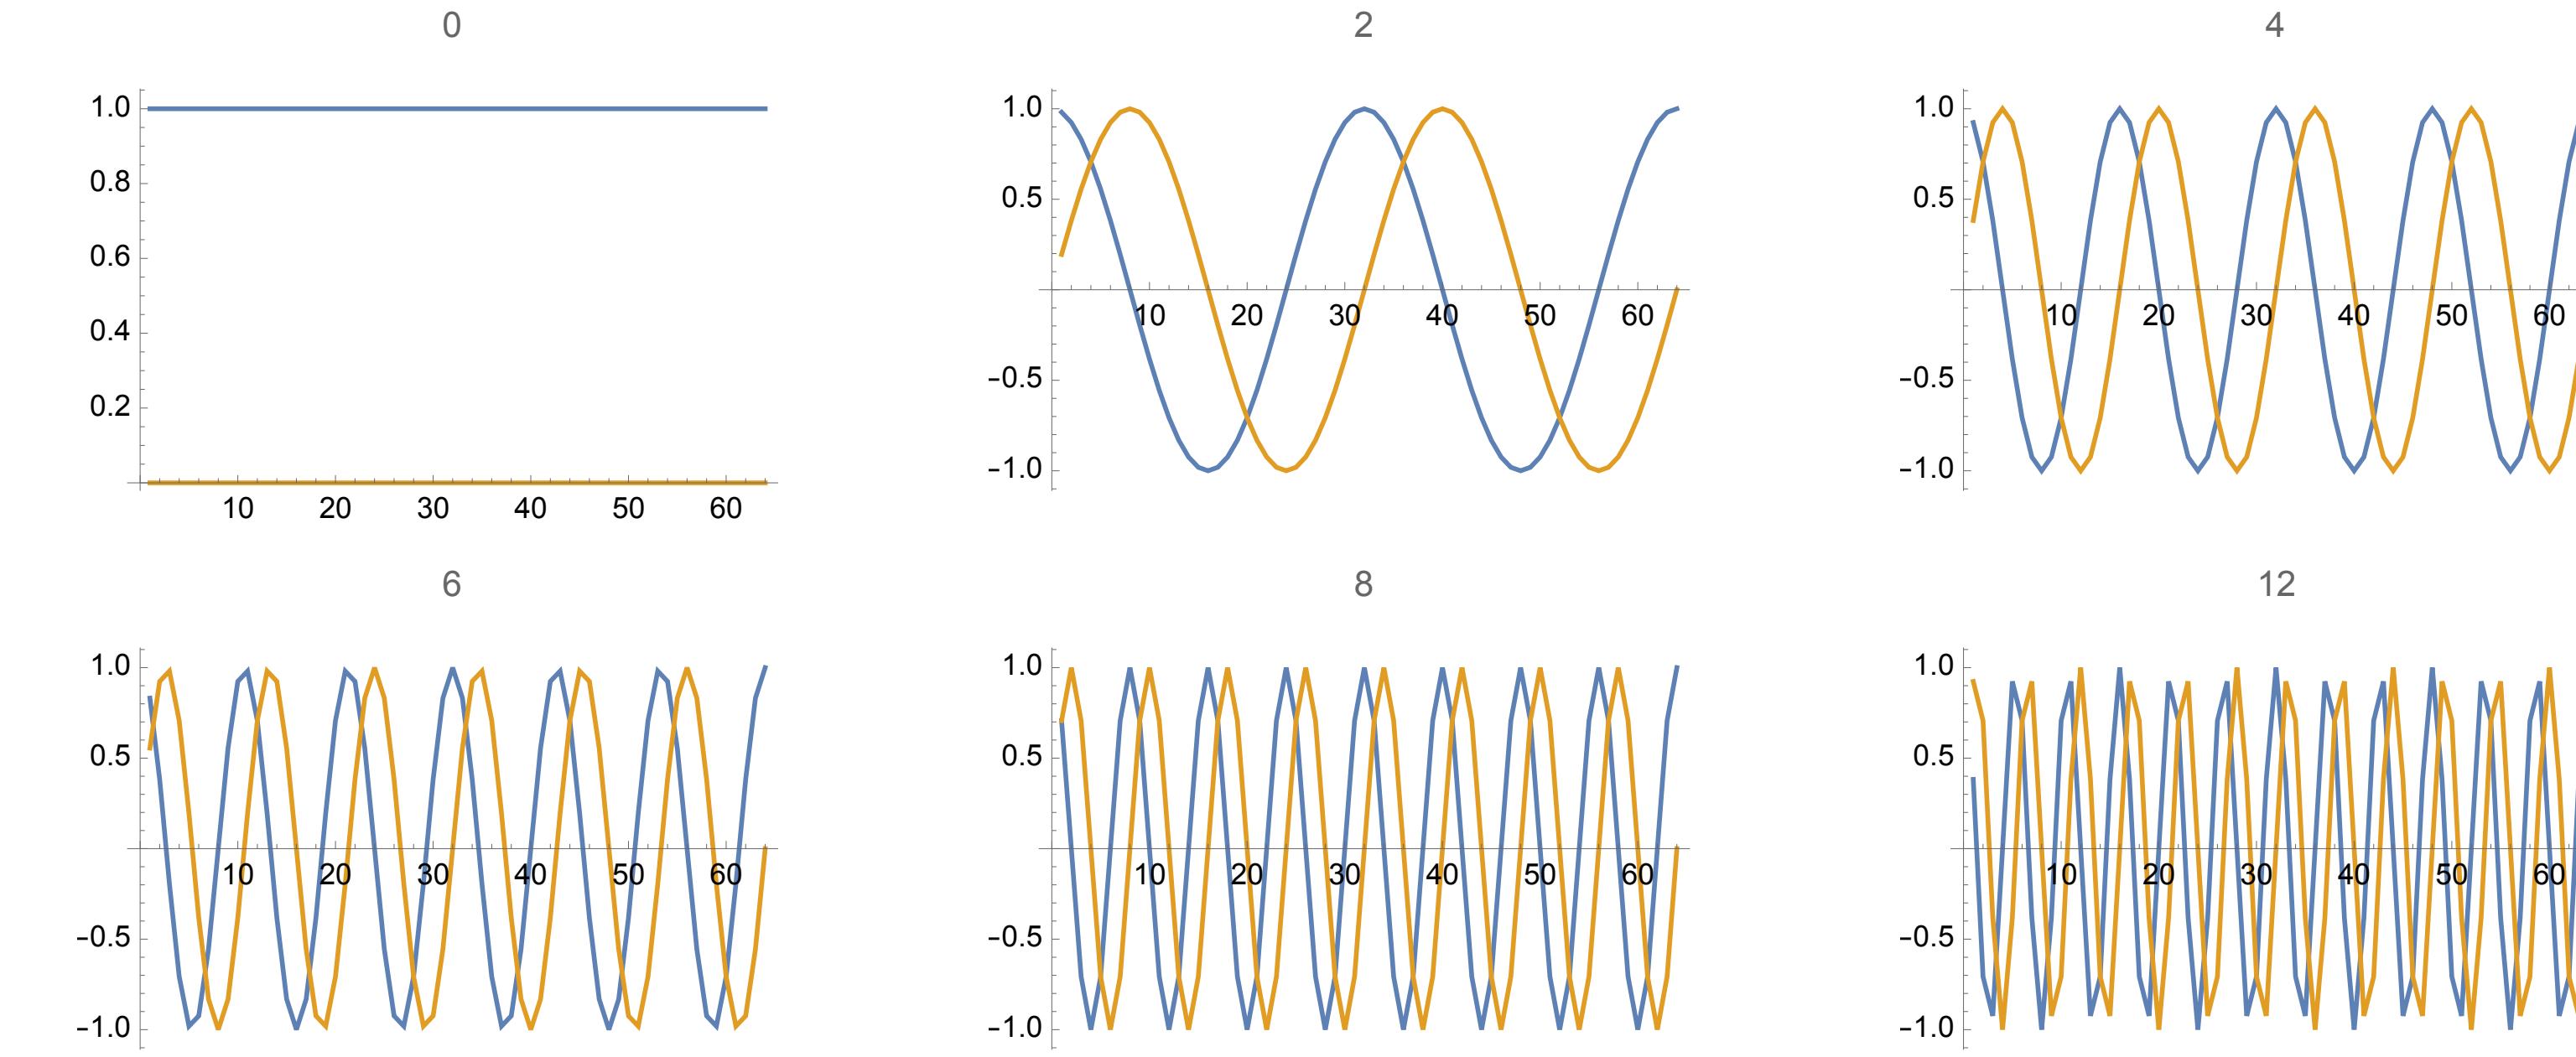
$$\mathbf{W} = \begin{bmatrix} 0 & 1 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 & 0 \\ 0 & 1 & 0 & 1 & 0 \\ 0 & 0 & 1 & 0 & 1 \\ 1 & 0 & 0 & 1 & 0 \end{bmatrix}$$

Adjacency
5 Vertices

- The adjacency matrix is *circulant*
 - And it's eigenvectors are DFT's complex exponentials

Special Case I: DFT

- DFT bases/eigenvectors: $u_k(n) = e^{i \frac{2\pi k n}{N}}$



- Paired sine and cosine terms
 - Similar local variation
 - Higher frequency for larger eigenvalues

Special Case II: DCT

- Consider an undirected Path graph
 - Assume all the weights are equal to unity



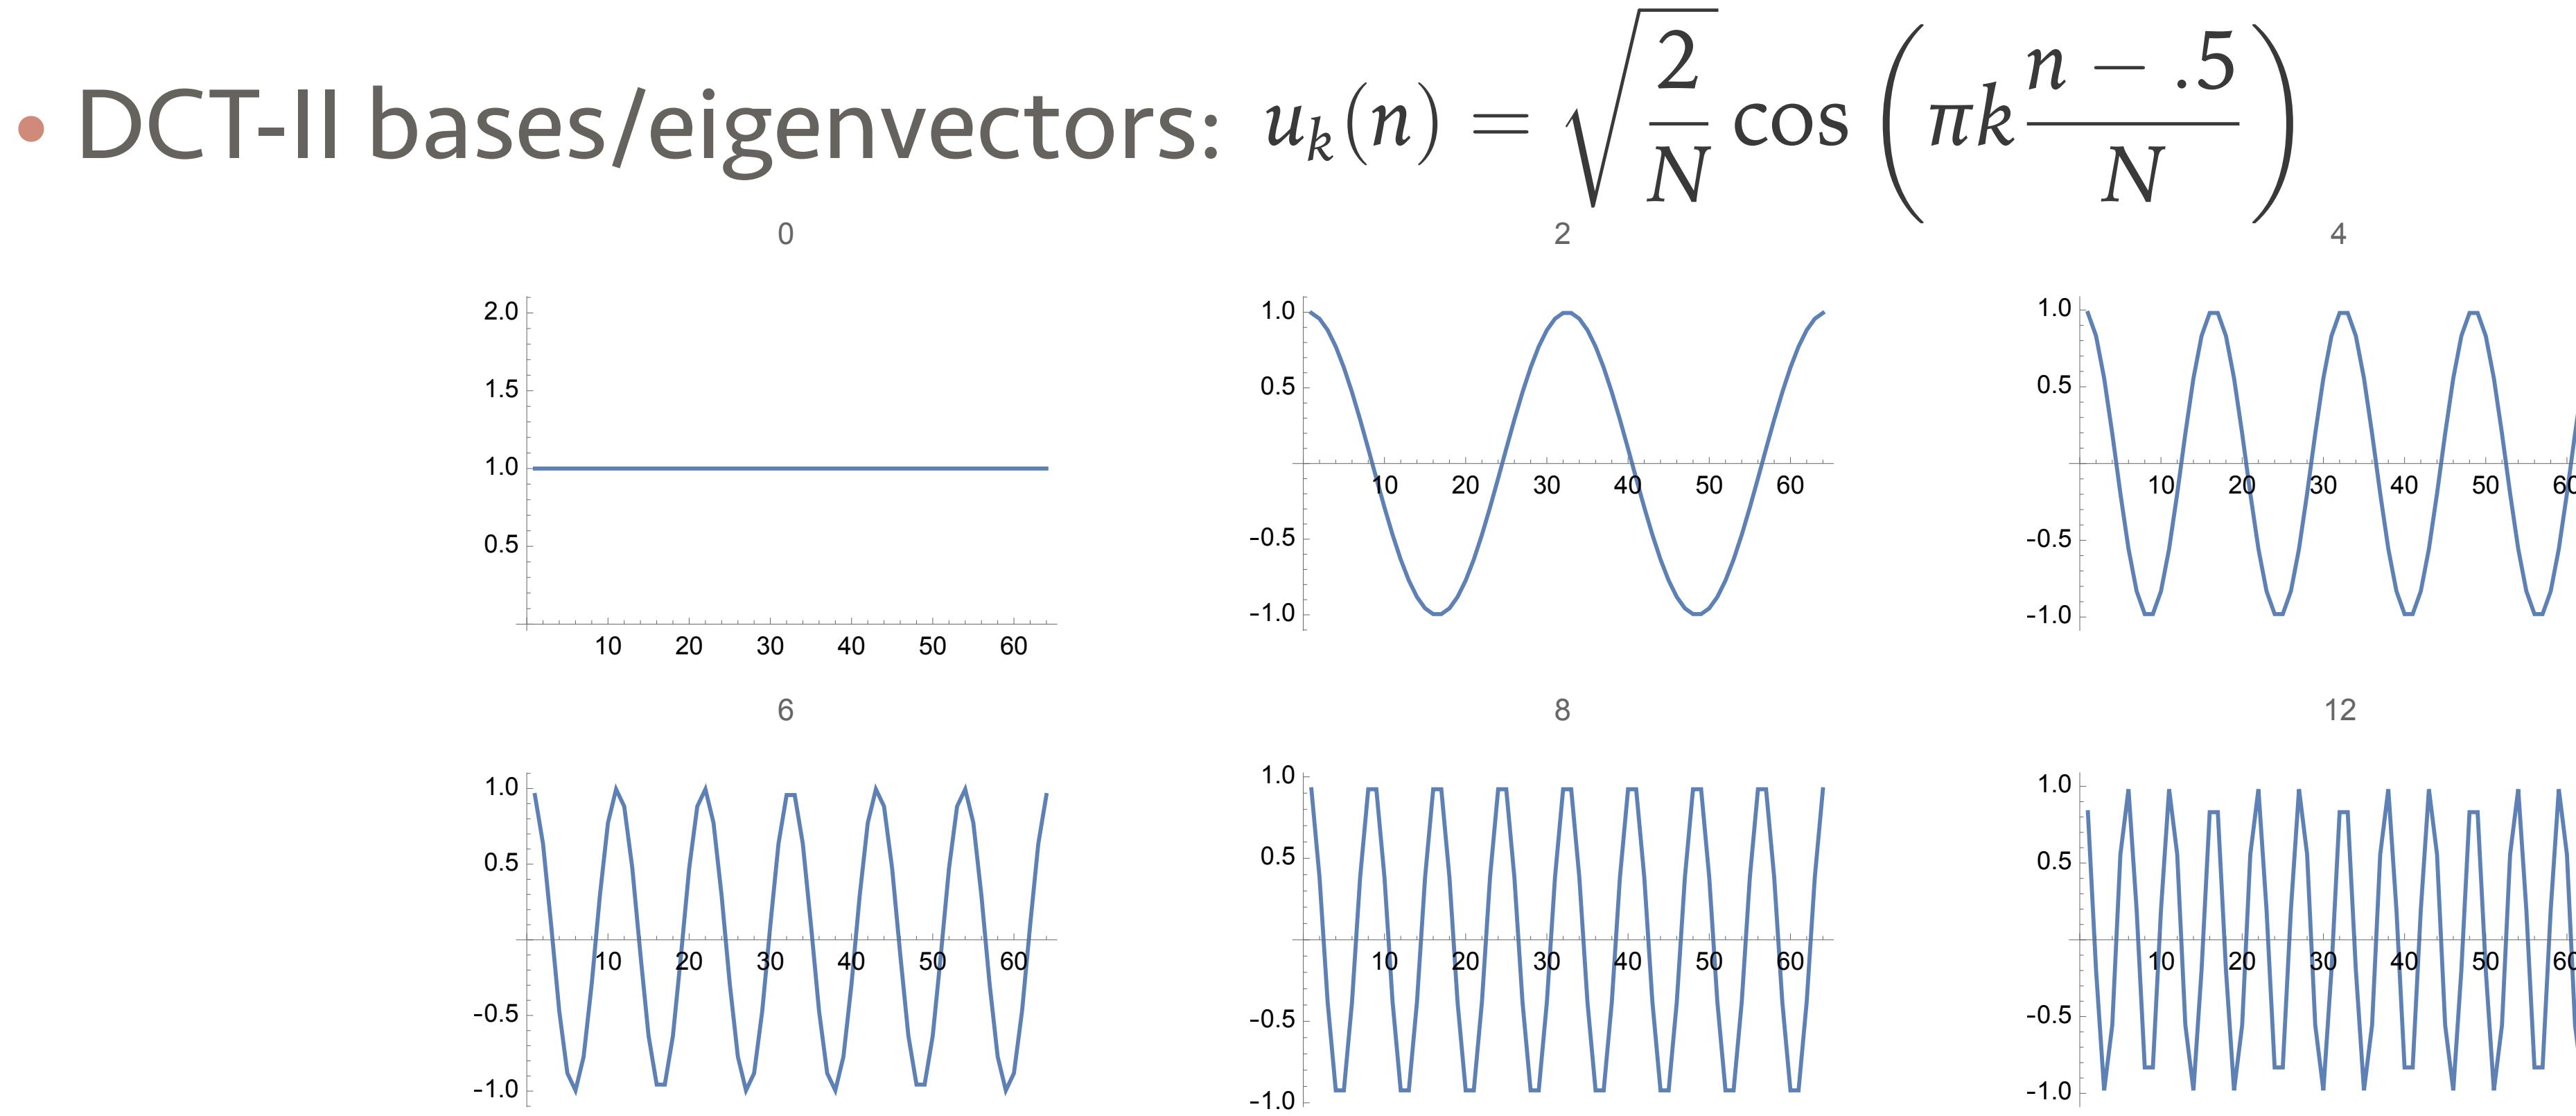
*Path graph
5 Vertices*

$$\mathbf{W} = \begin{bmatrix} 0 & 1 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 \\ 0 & 1 & 0 & 1 & 0 \\ 0 & 0 & 1 & 0 & 1 \\ 0 & 0 & 0 & 1 & 0 \end{bmatrix}$$

*Adjacency
5 Vertices*

- The adjacency matrix is *Toeplitz*
 - Eigenvectors will be sinusoids

Special Case II: DCT



- Type 2 Discrete Cosine Transform
 - Real-valued alternative to Fourier Transforms
 - What PCA tends to give us!

Imposing a Graph

- So far
 - Graphs are a collection of nodes and edges
 - Graphs are defined by their adjacency matrices \mathbf{W}
 - The graph Laplacian \mathbf{L} can be computed from \mathbf{W}
 - Used to compute the Graph Fourier Transform
 - Generalizes classical Fourier Transform theory
- We assumed that the data lived on an underlying graph
 - Can we learn a graph from the data itself?
 - Can audio / images be interpreted as a graph?
- Learning a Graph = Learning an adjacency matrix \mathbf{W}
 - Easy to learn adjacency matrices from the data
 - Use a suitable distance metric to get \mathbf{W}

Defining a graph from a signal

- Adjacency based on spatial location

$$w_t(i, j) = e^{-\alpha|t_i - t_j|}$$

- Can be time-indices or spatial coordinates

- Adjacency based on sample values

$$w_\nu(i, j) = e^{-\beta|\nu_i - \nu_j|}$$

- Can be features, sample-values, image colors etc.

- Overall Adjacency matrix

$$\mathbf{W} = \gamma \mathbf{W}_t + (1 - \gamma) \mathbf{W}_\nu$$

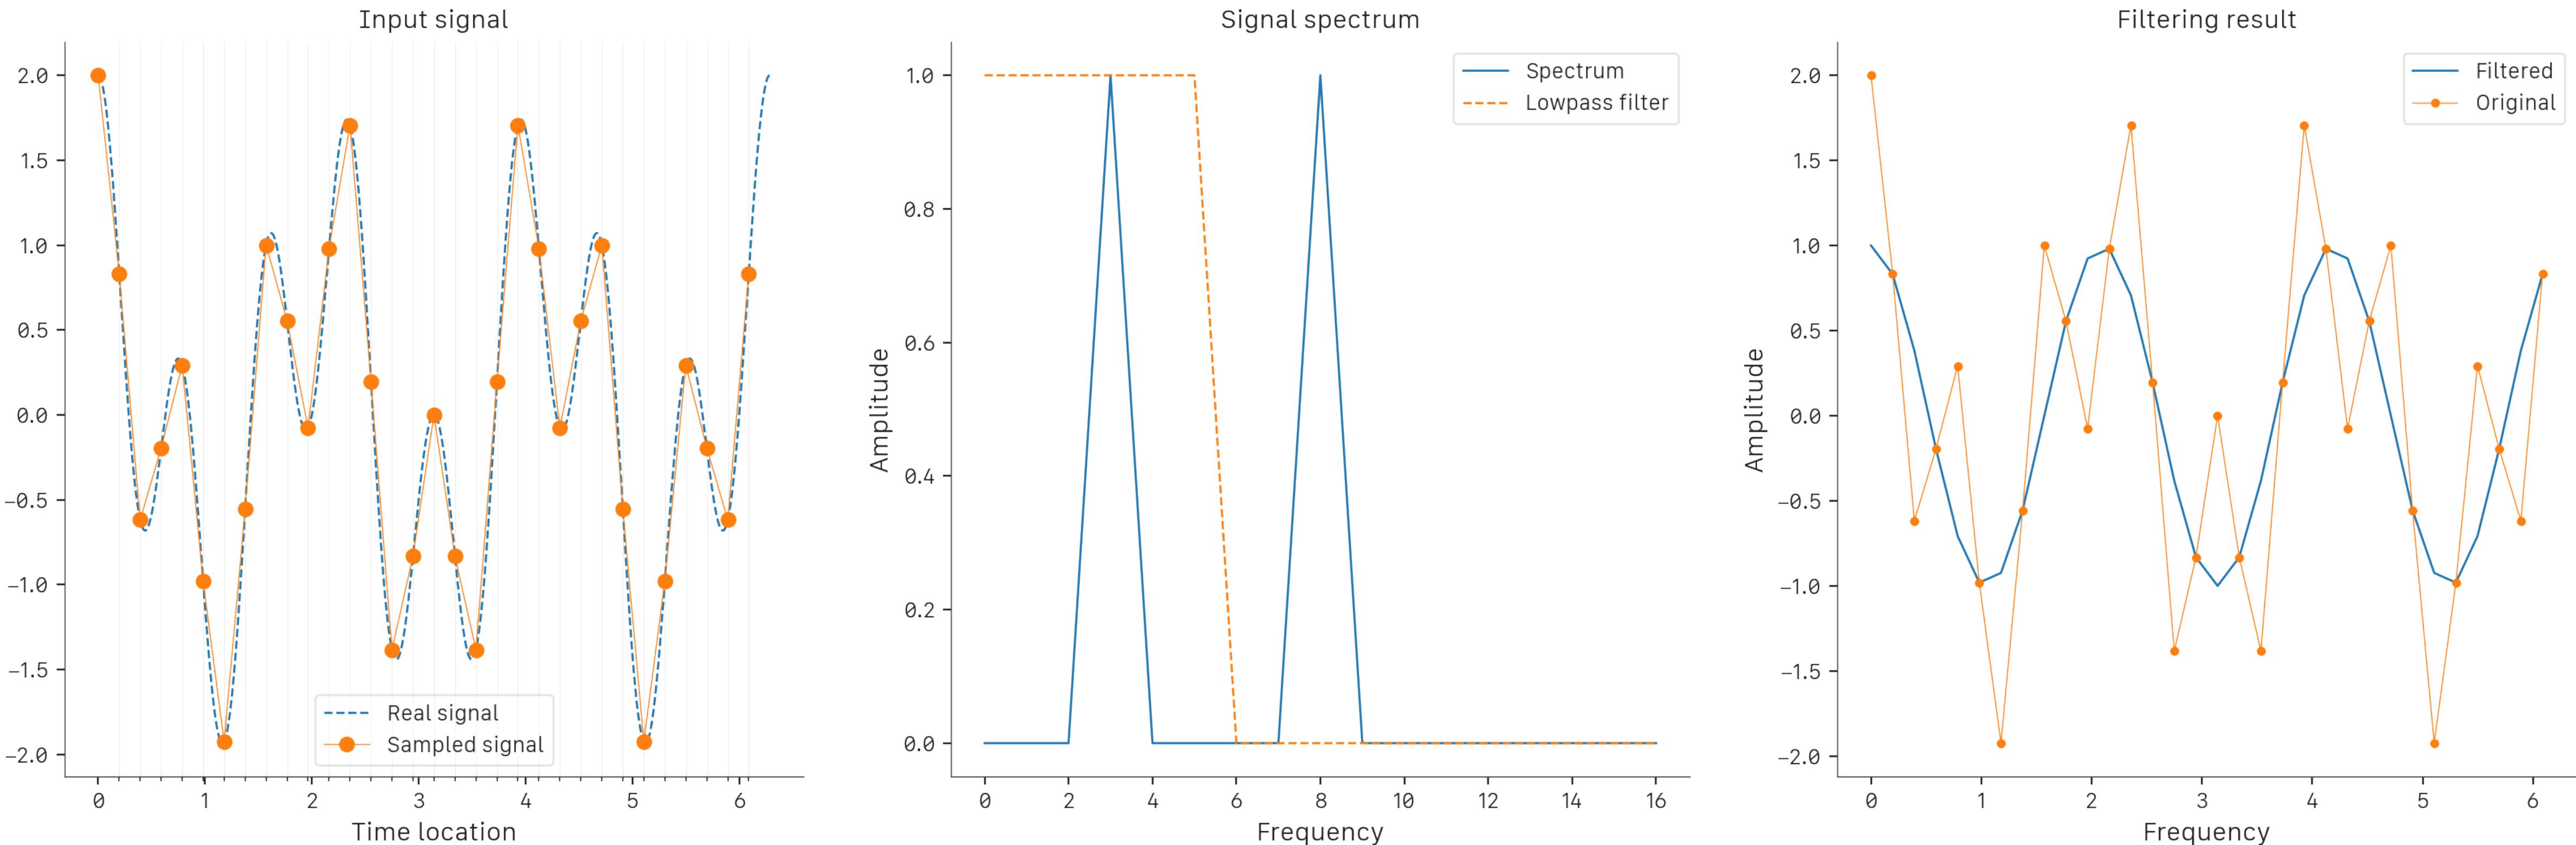
Classical filtering example

- Applying a filter in the Fourier domain

$$x(t) = \cos(3t) + \cos(8t)$$

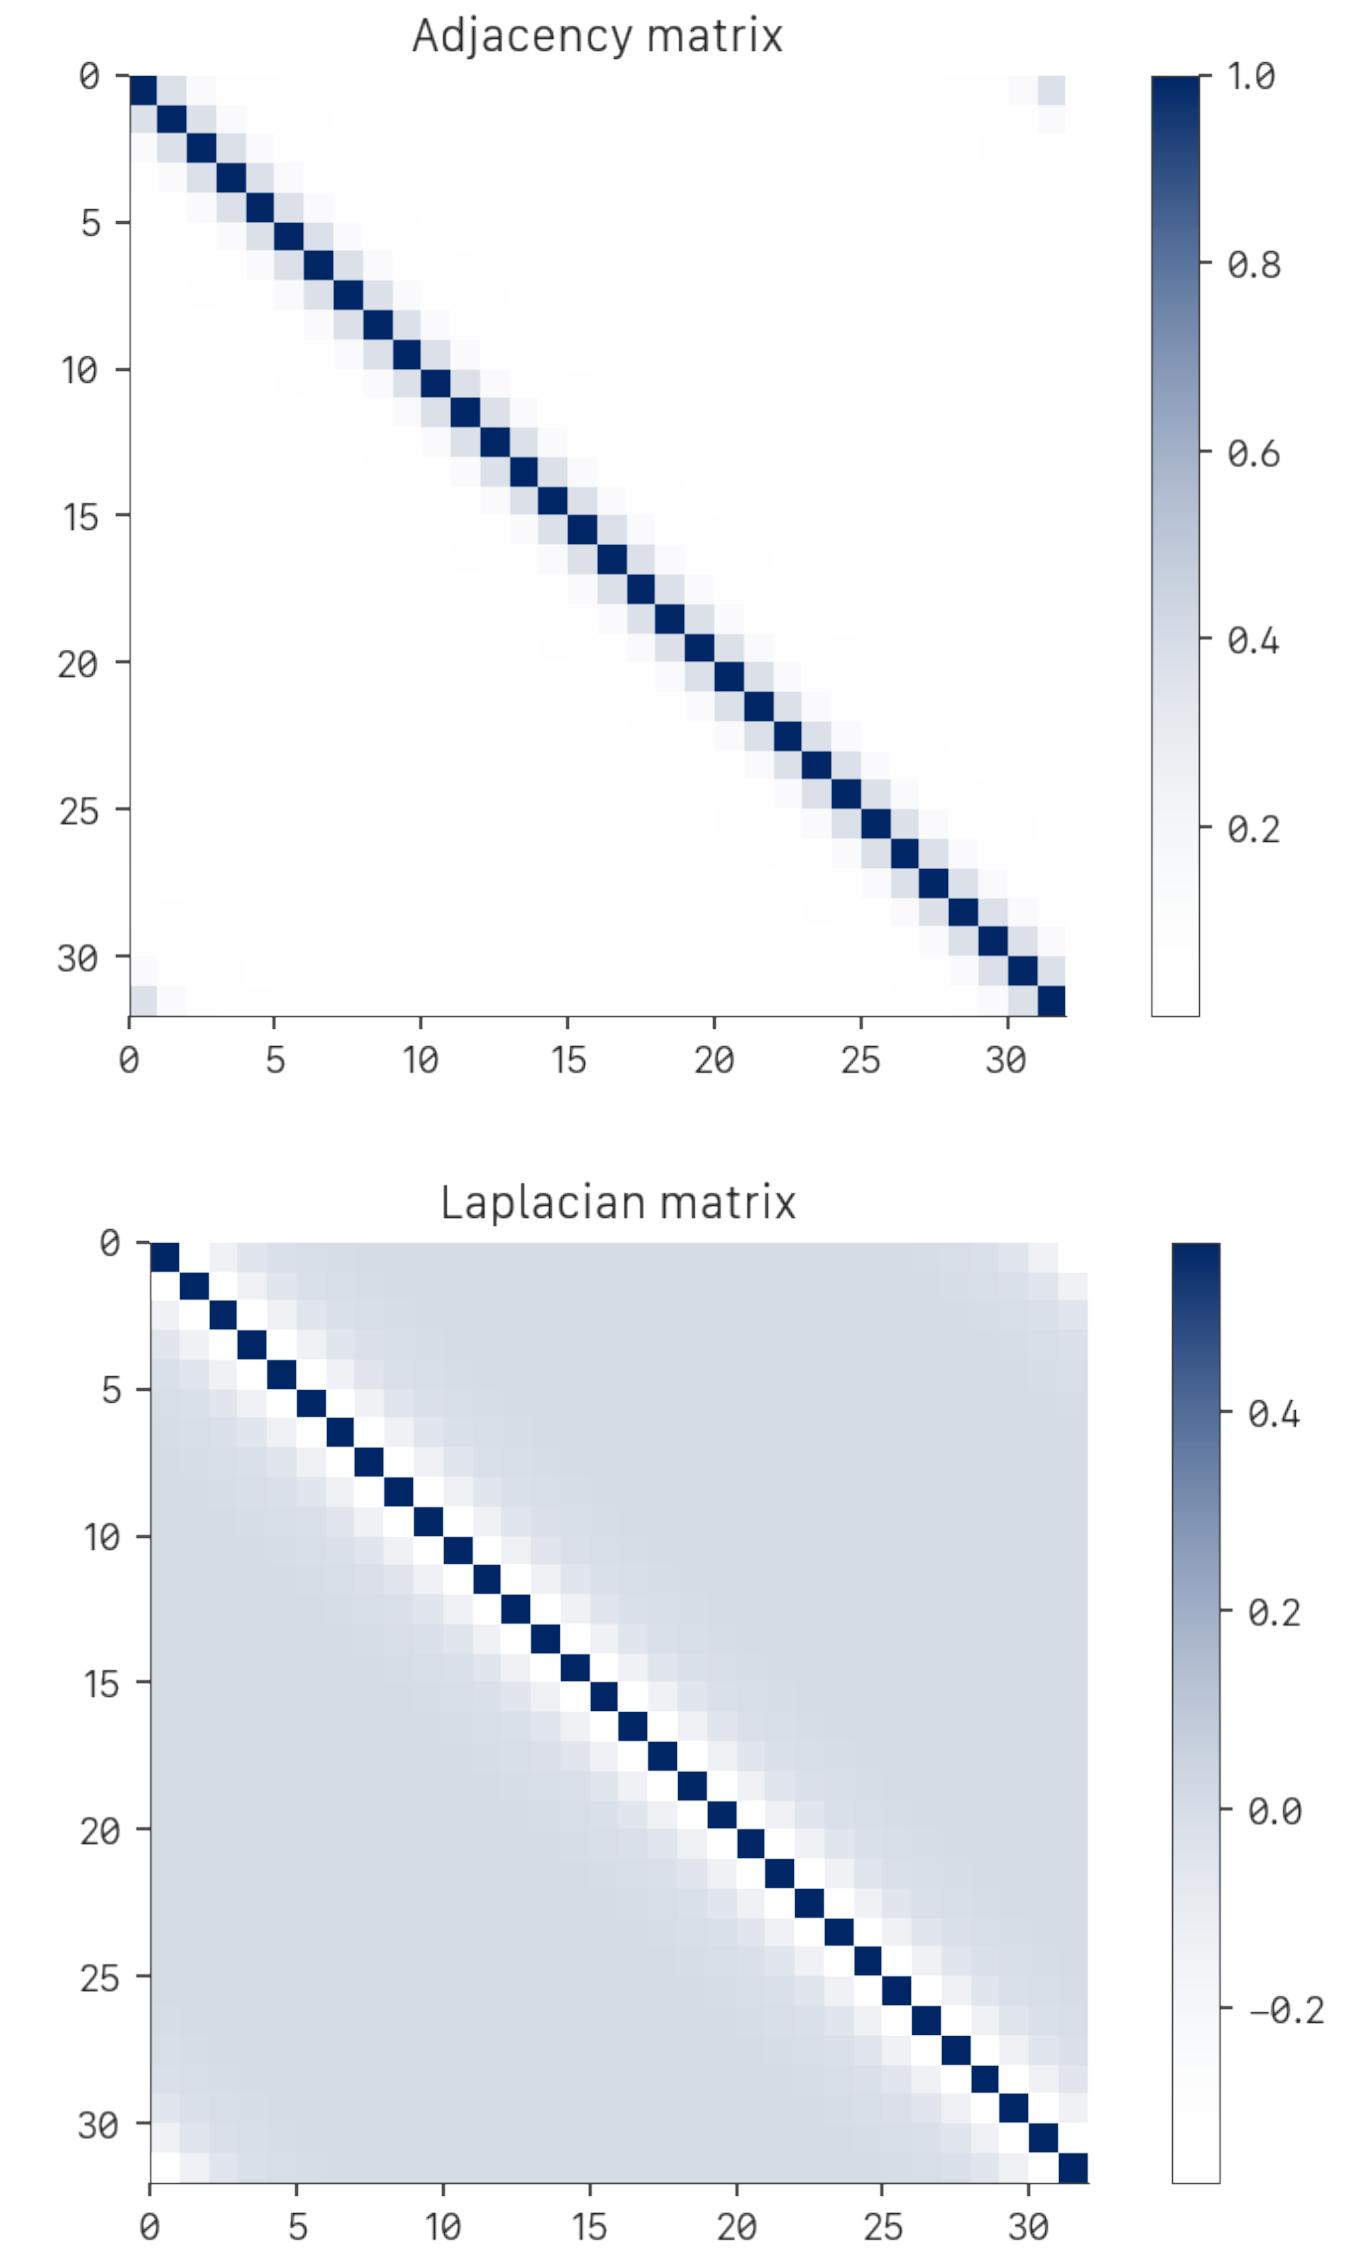
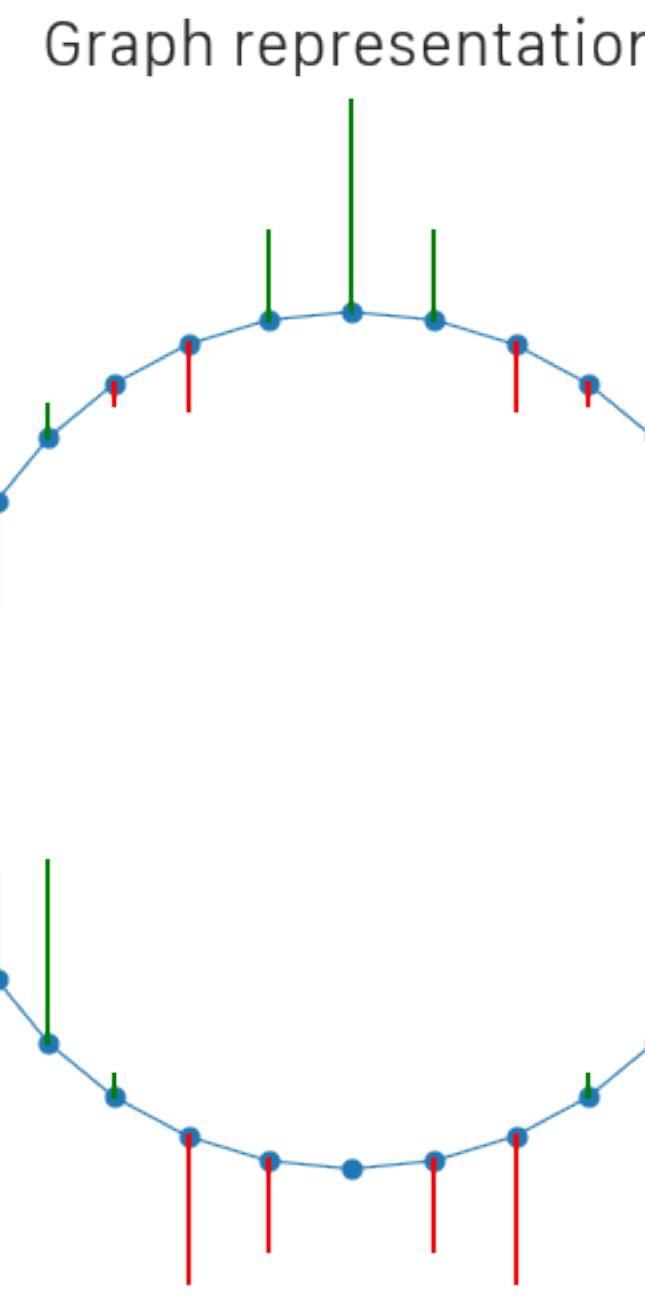
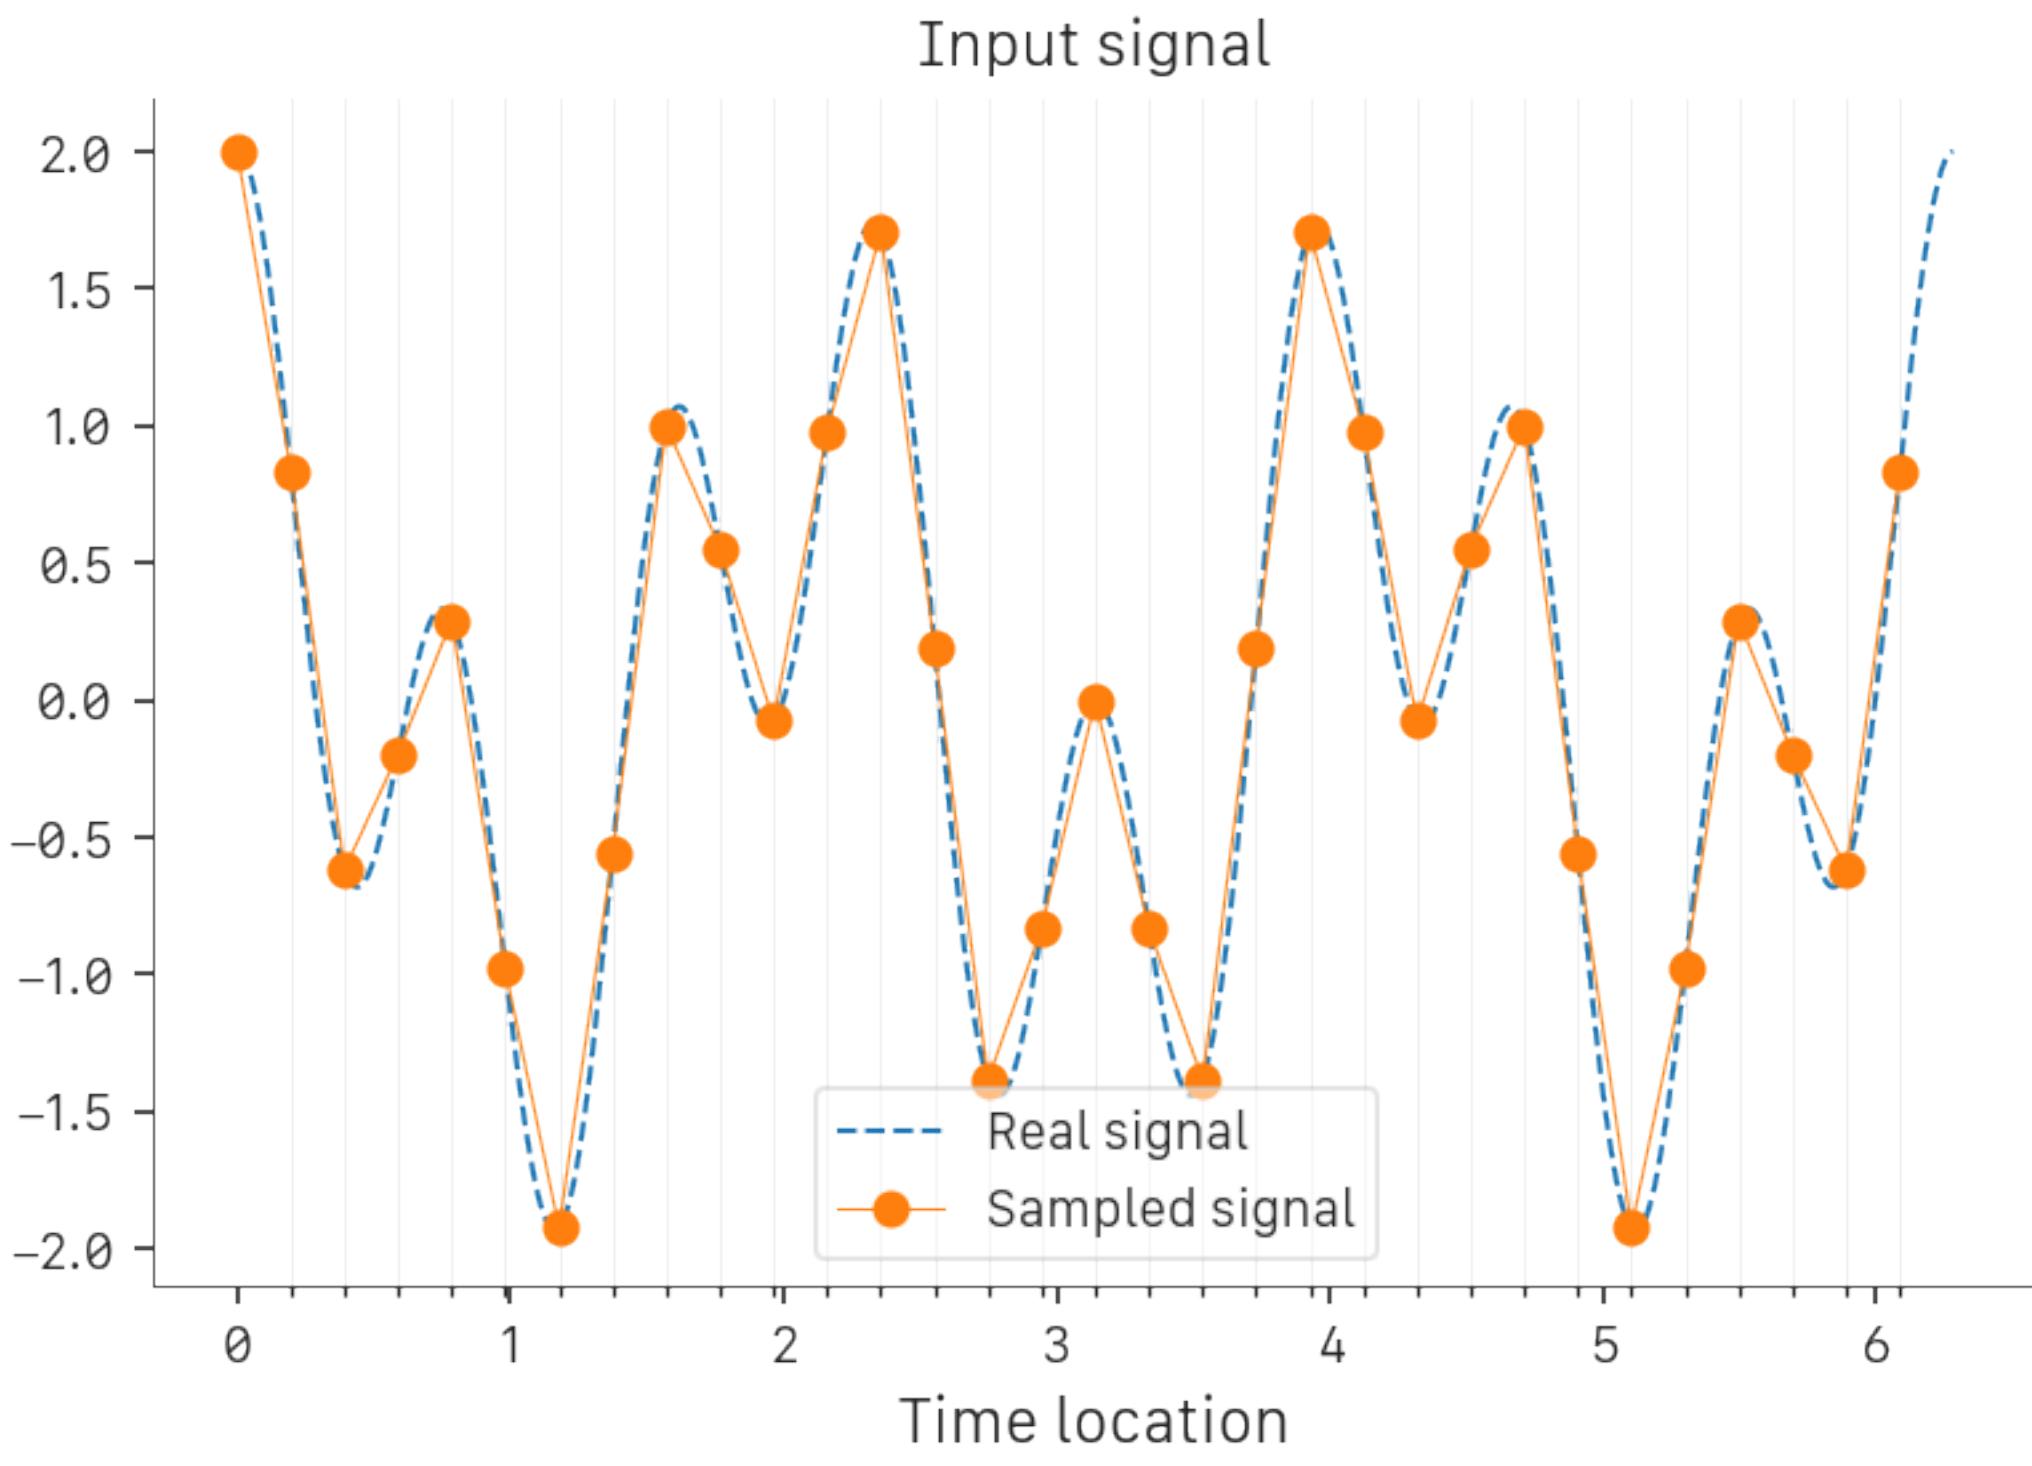
*Transform signal into frequency domain
Apply lowpass filter*

Invert to time-domain



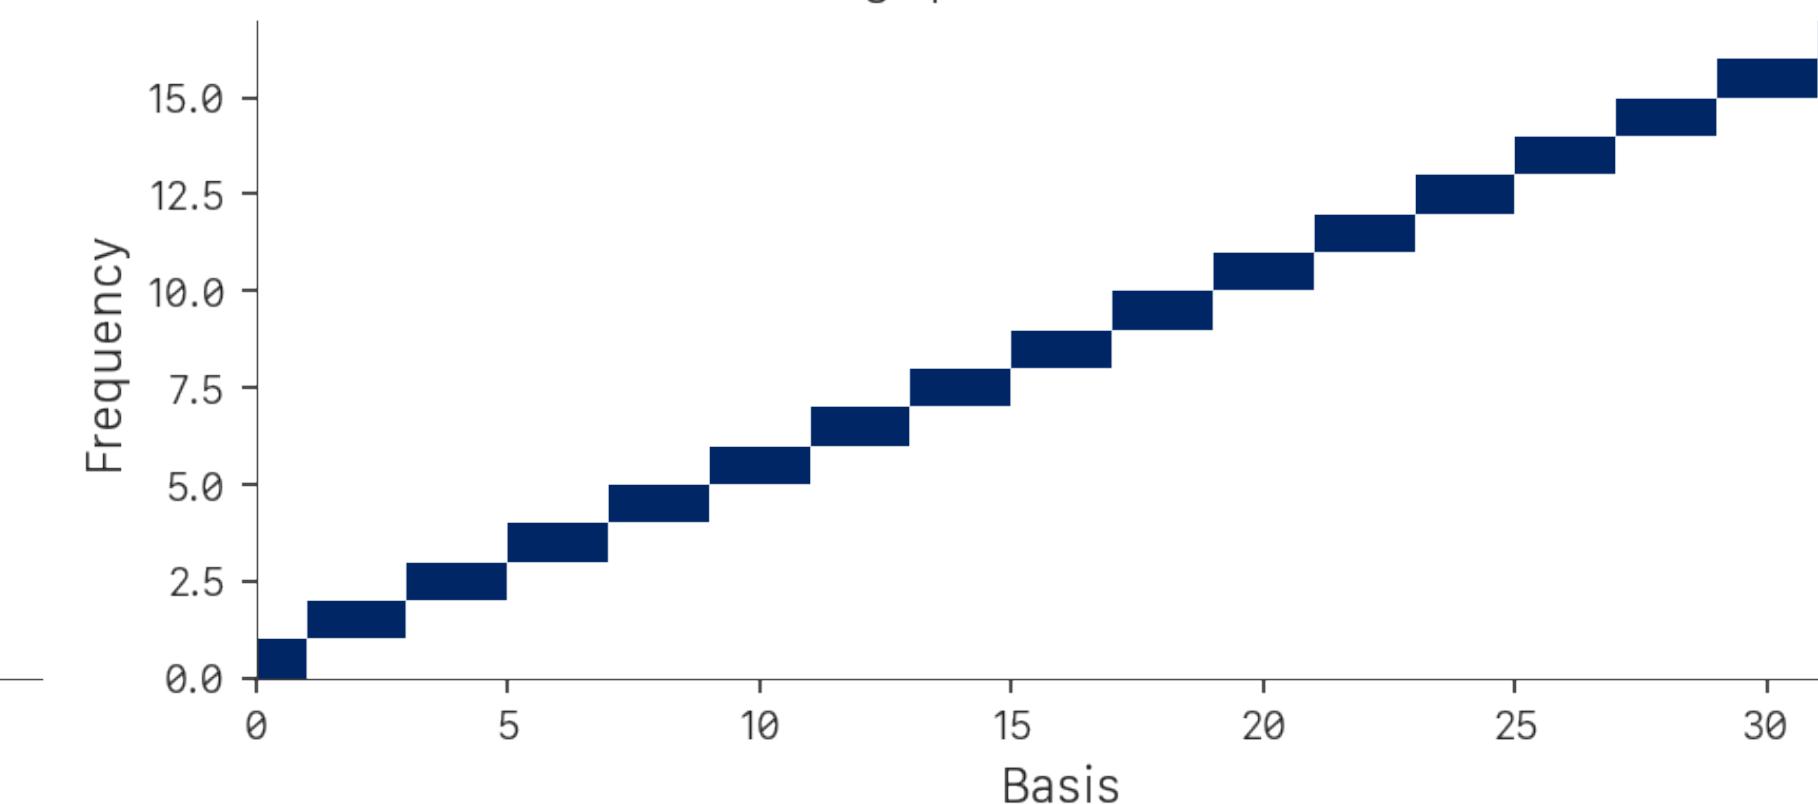
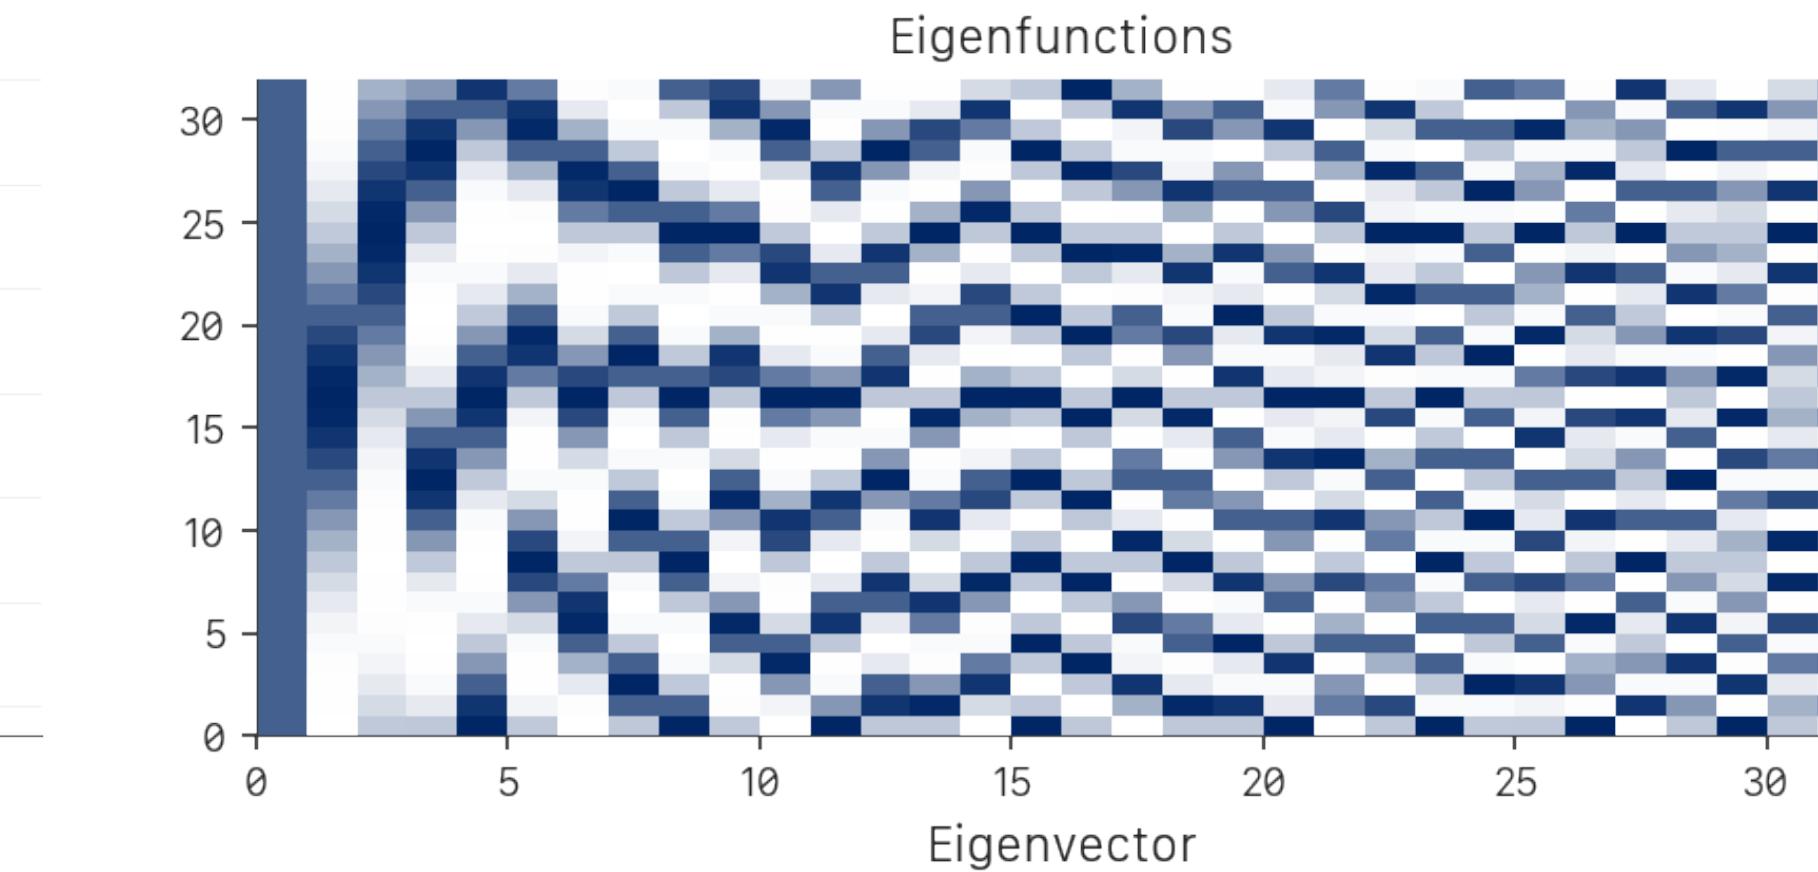
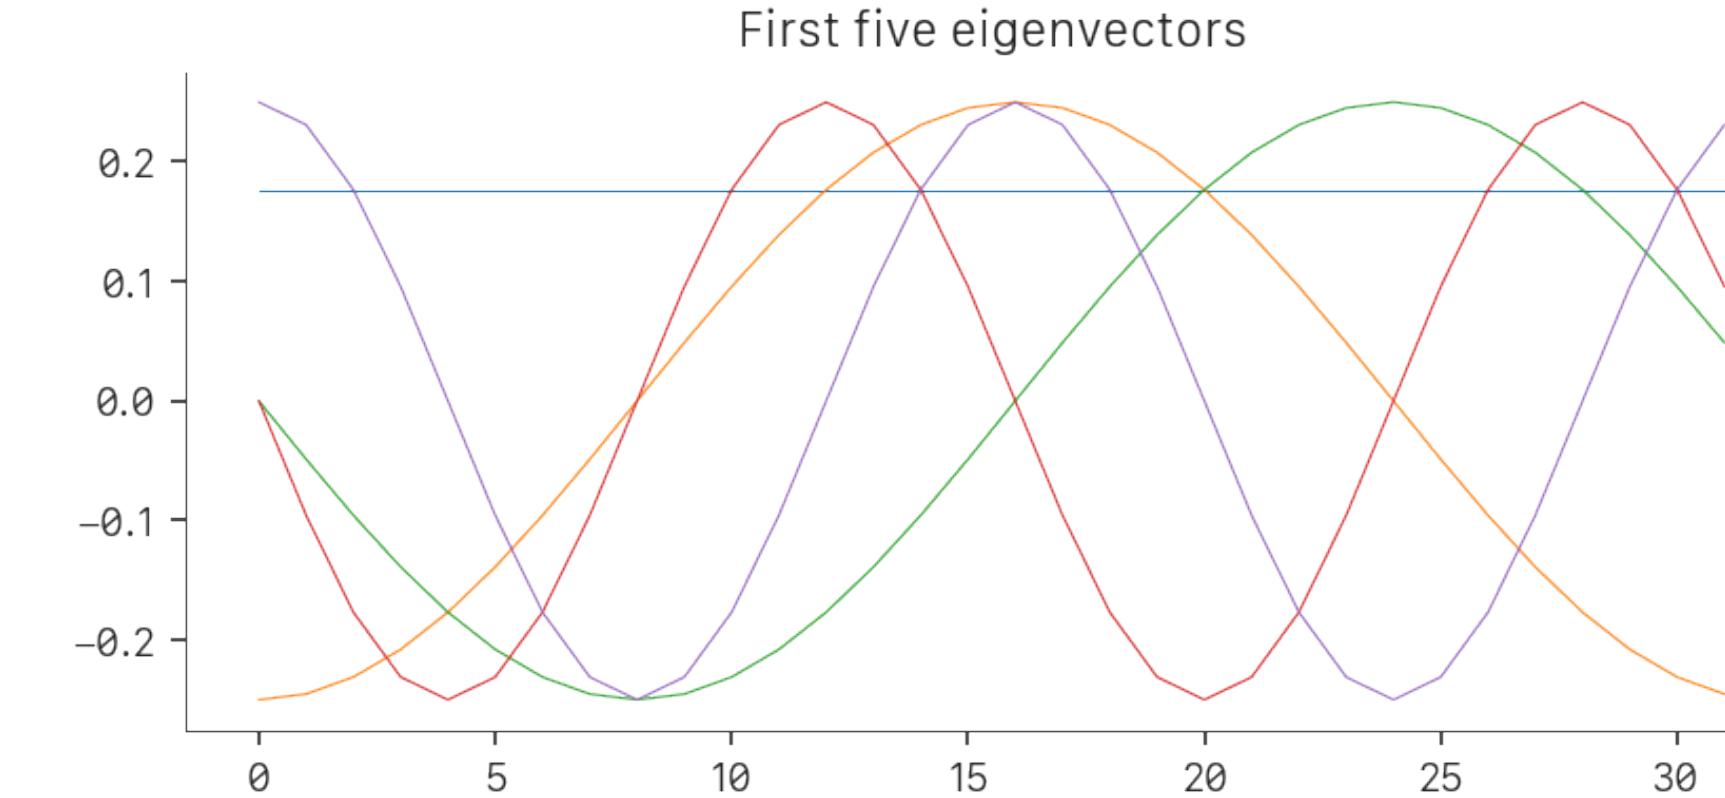
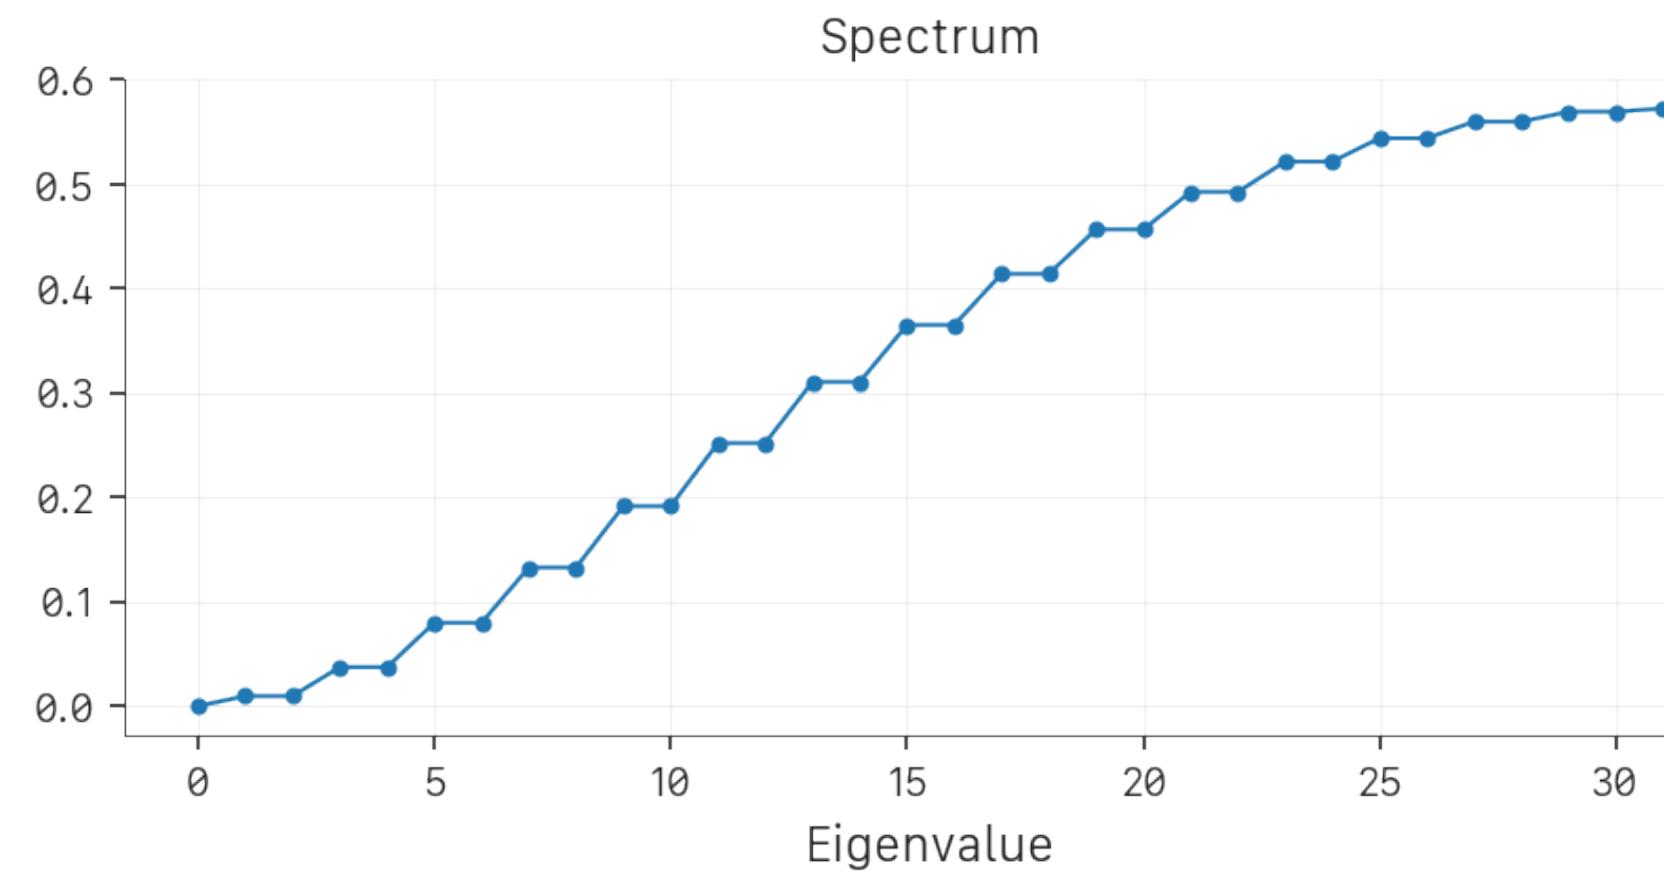
Filtering on Graphs

- Define a simple circular graph
 - Place sampled values in vertices



Step 1: Get frequency bases

- Eigendecomposition of the Graph Laplacian
 - Learn a suitable Graph Fourier transform $\mathbf{L} = \mathbf{V} \cdot \Lambda \cdot \mathbf{V}^\top$



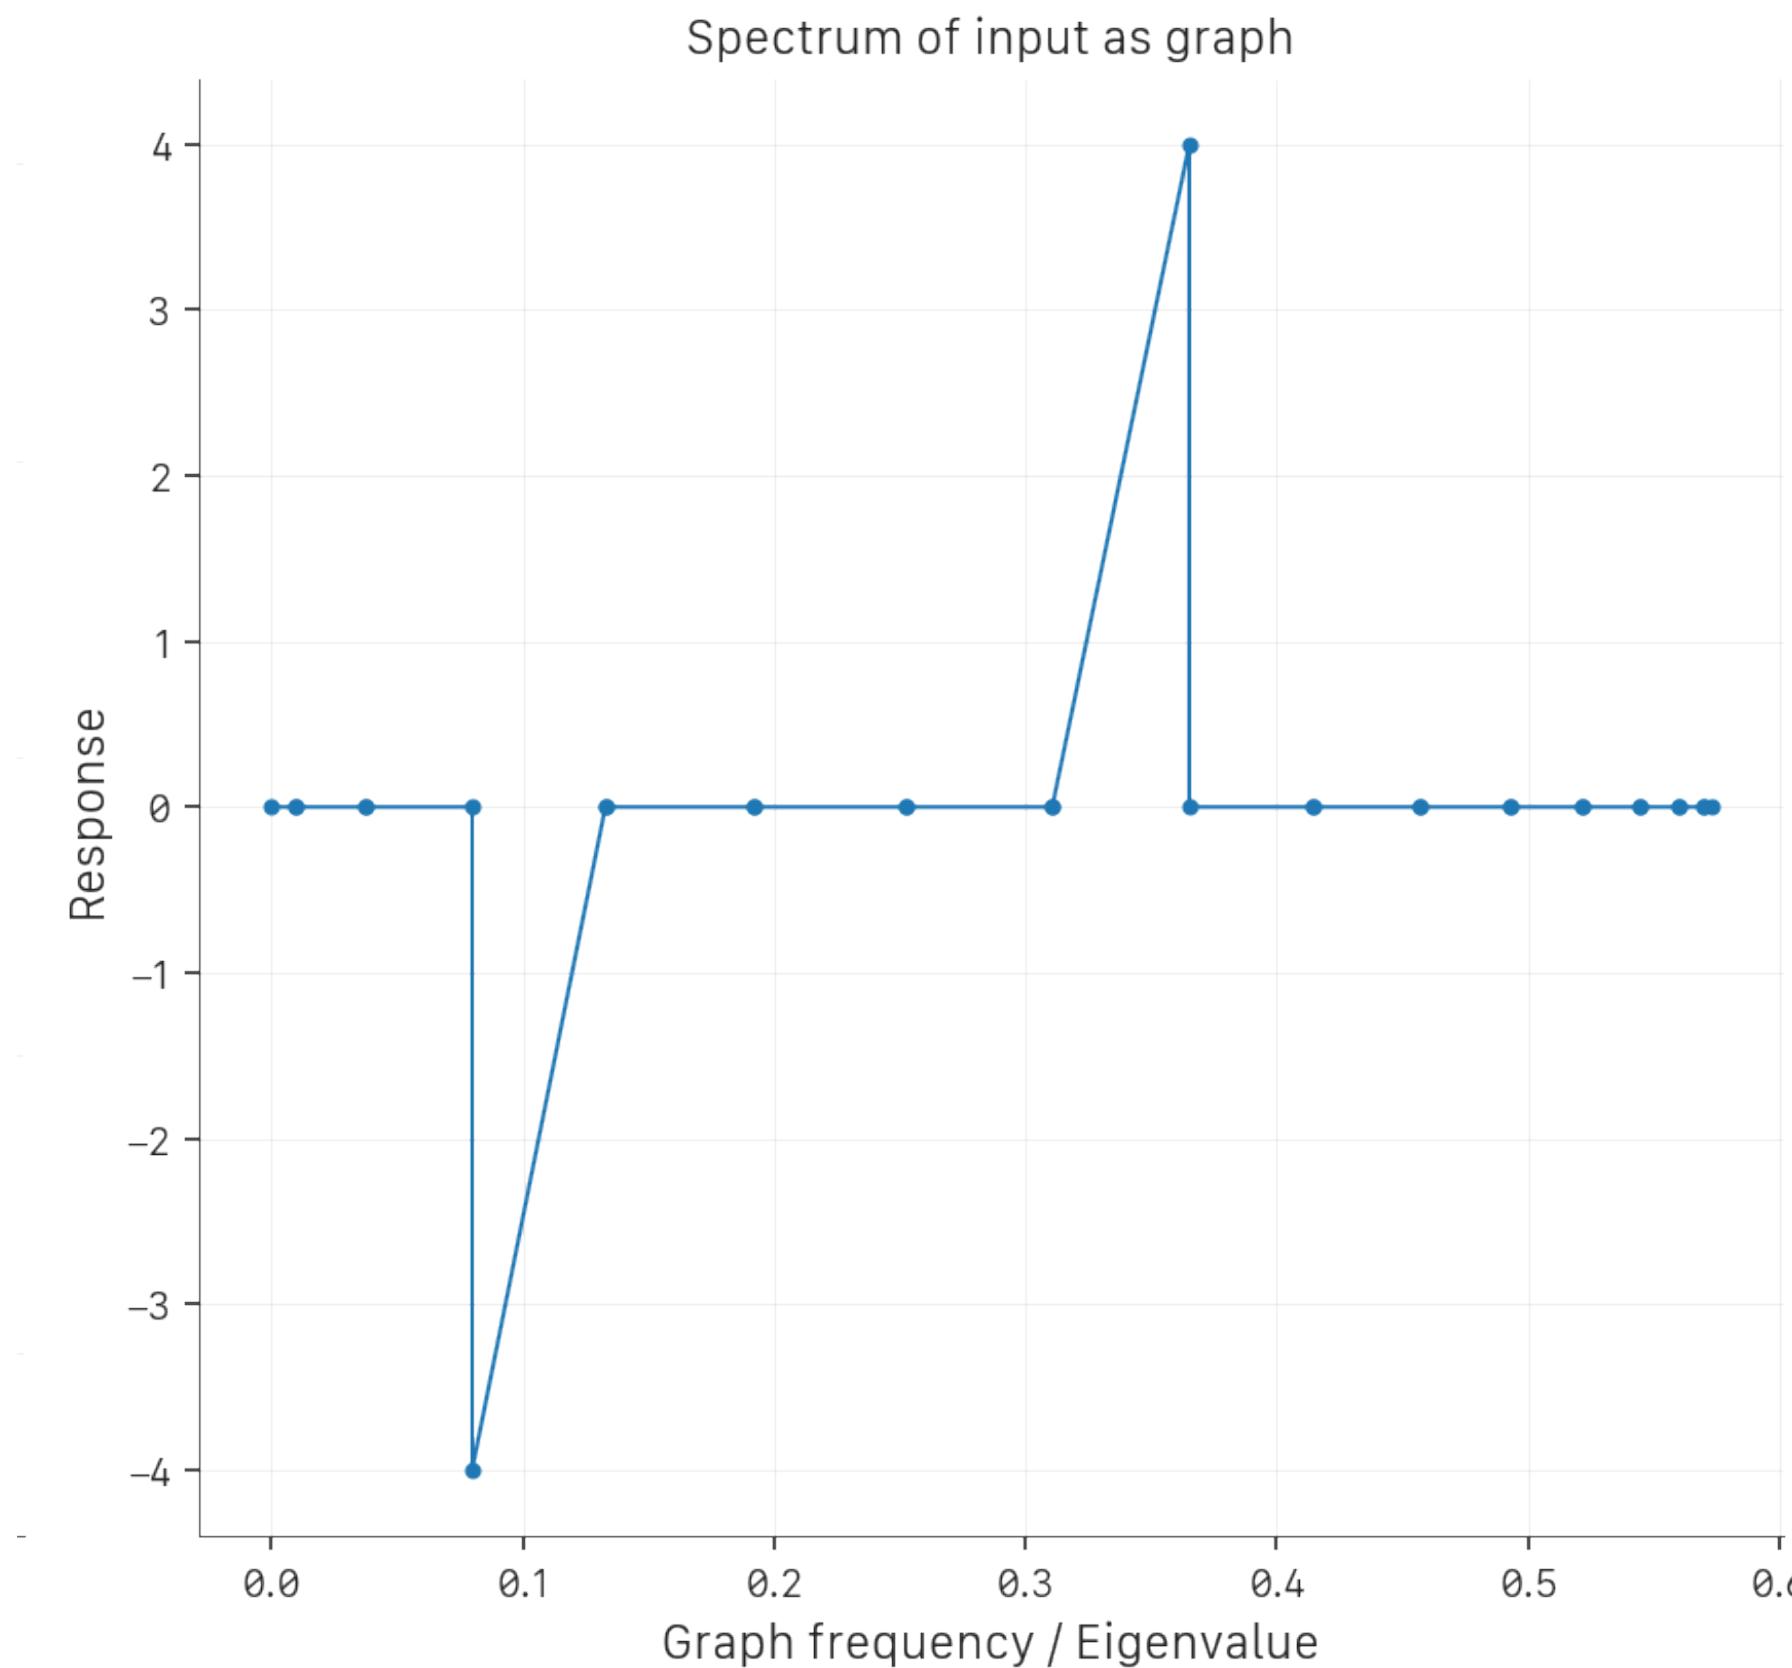
Step 2: Transform Data

- Transform data into Fourier Domain
 - Use the eigenvectors learned in previous step

$$\mathbf{X} = \mathbf{V}^\top \cdot \mathbf{x}$$

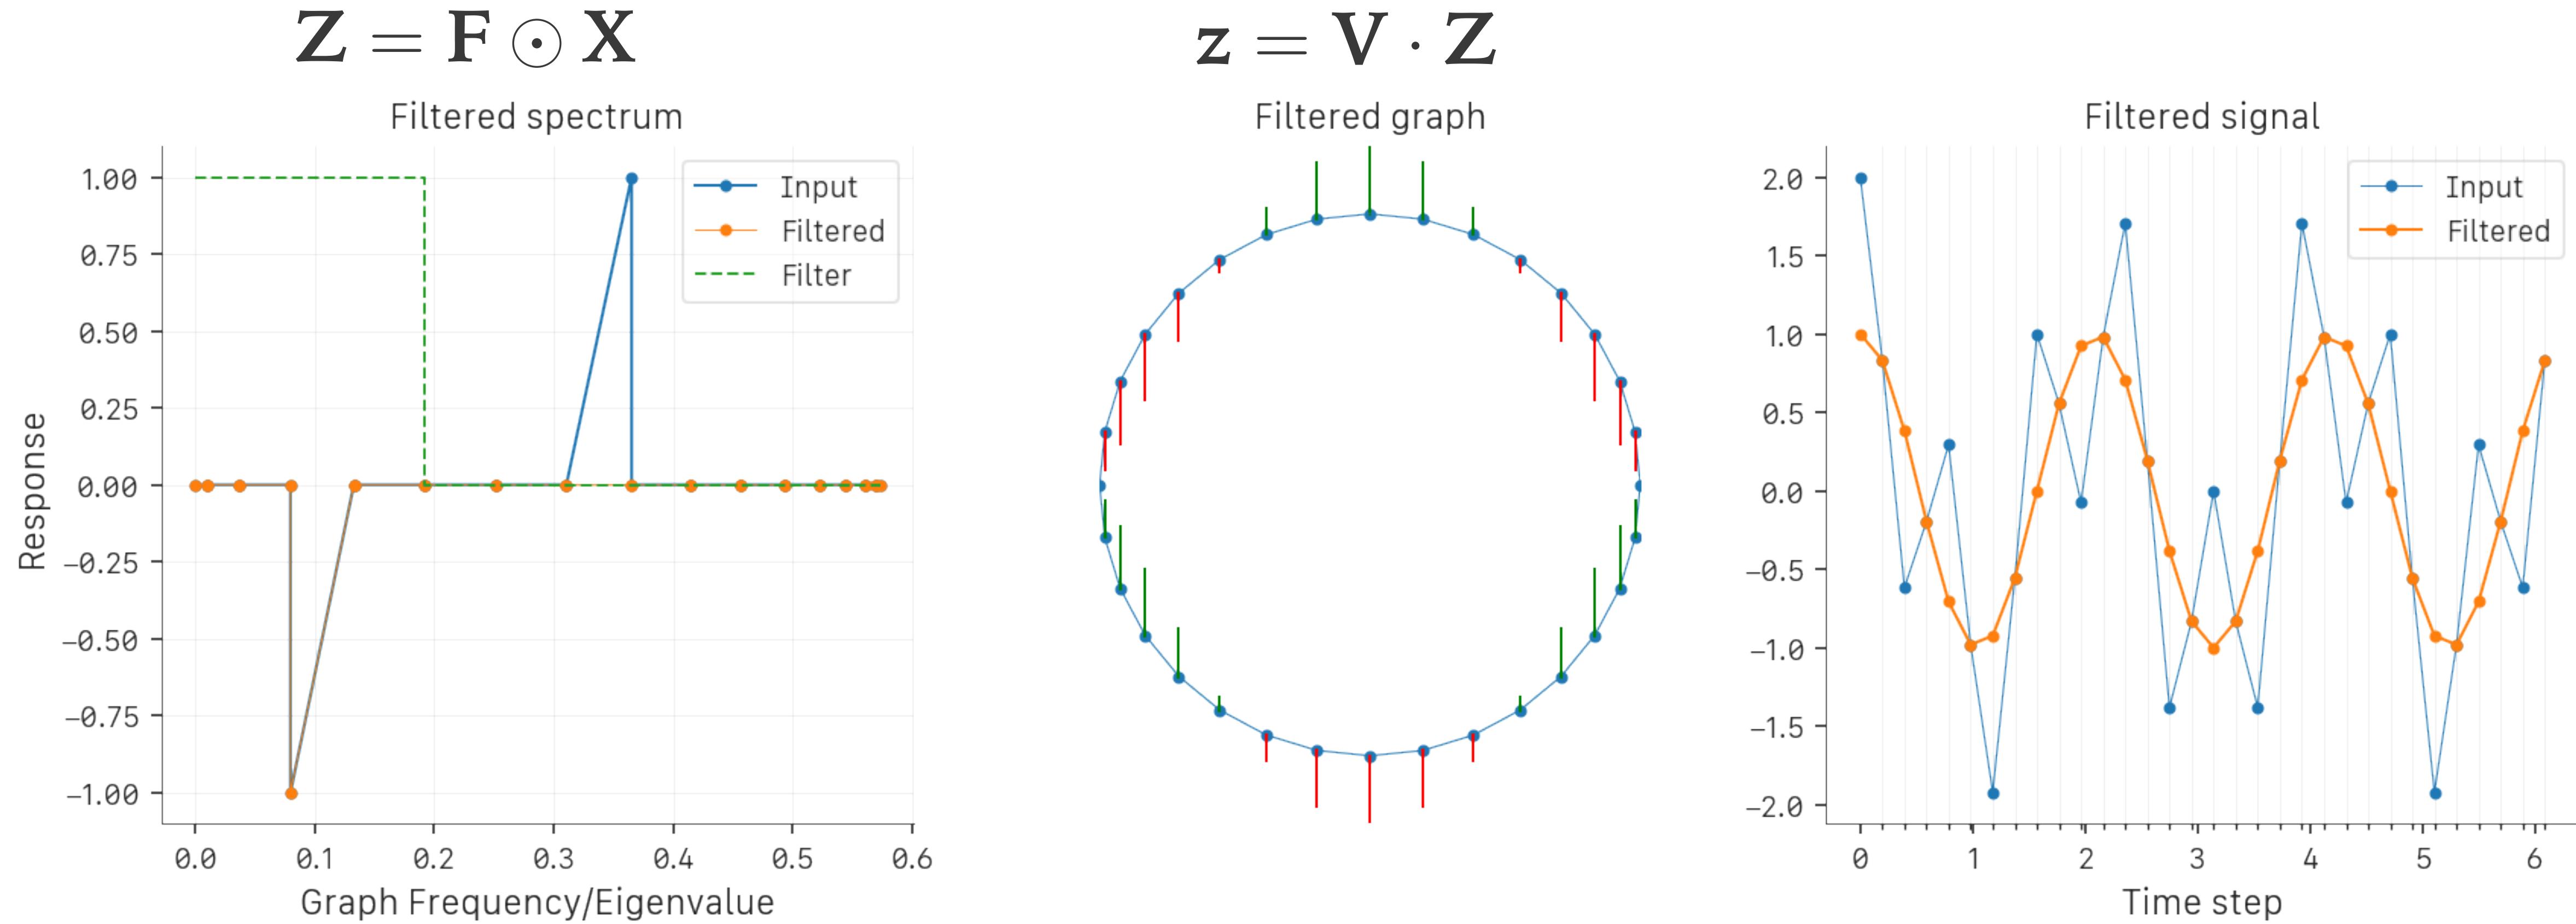
Diagram illustrating the transformation from a Graph Signal to Signal Fourier coefficients:

- Graph Fourier bases*: Points to the matrix \mathbf{V}^\top .
- Graph Signal*: Points to the vector \mathbf{x} .
- Signal Fourier coefficients*: Points to the resulting vector \mathbf{X} .



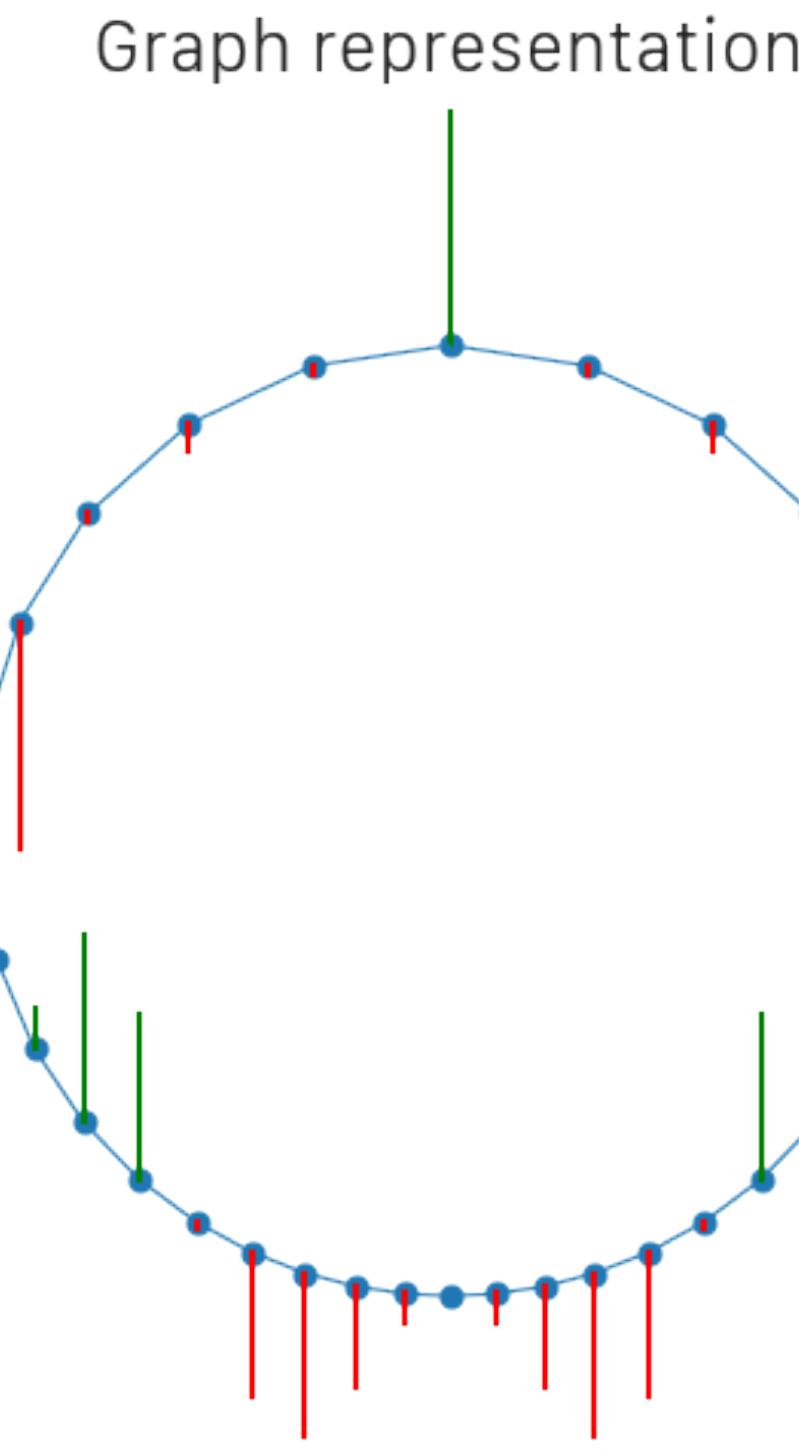
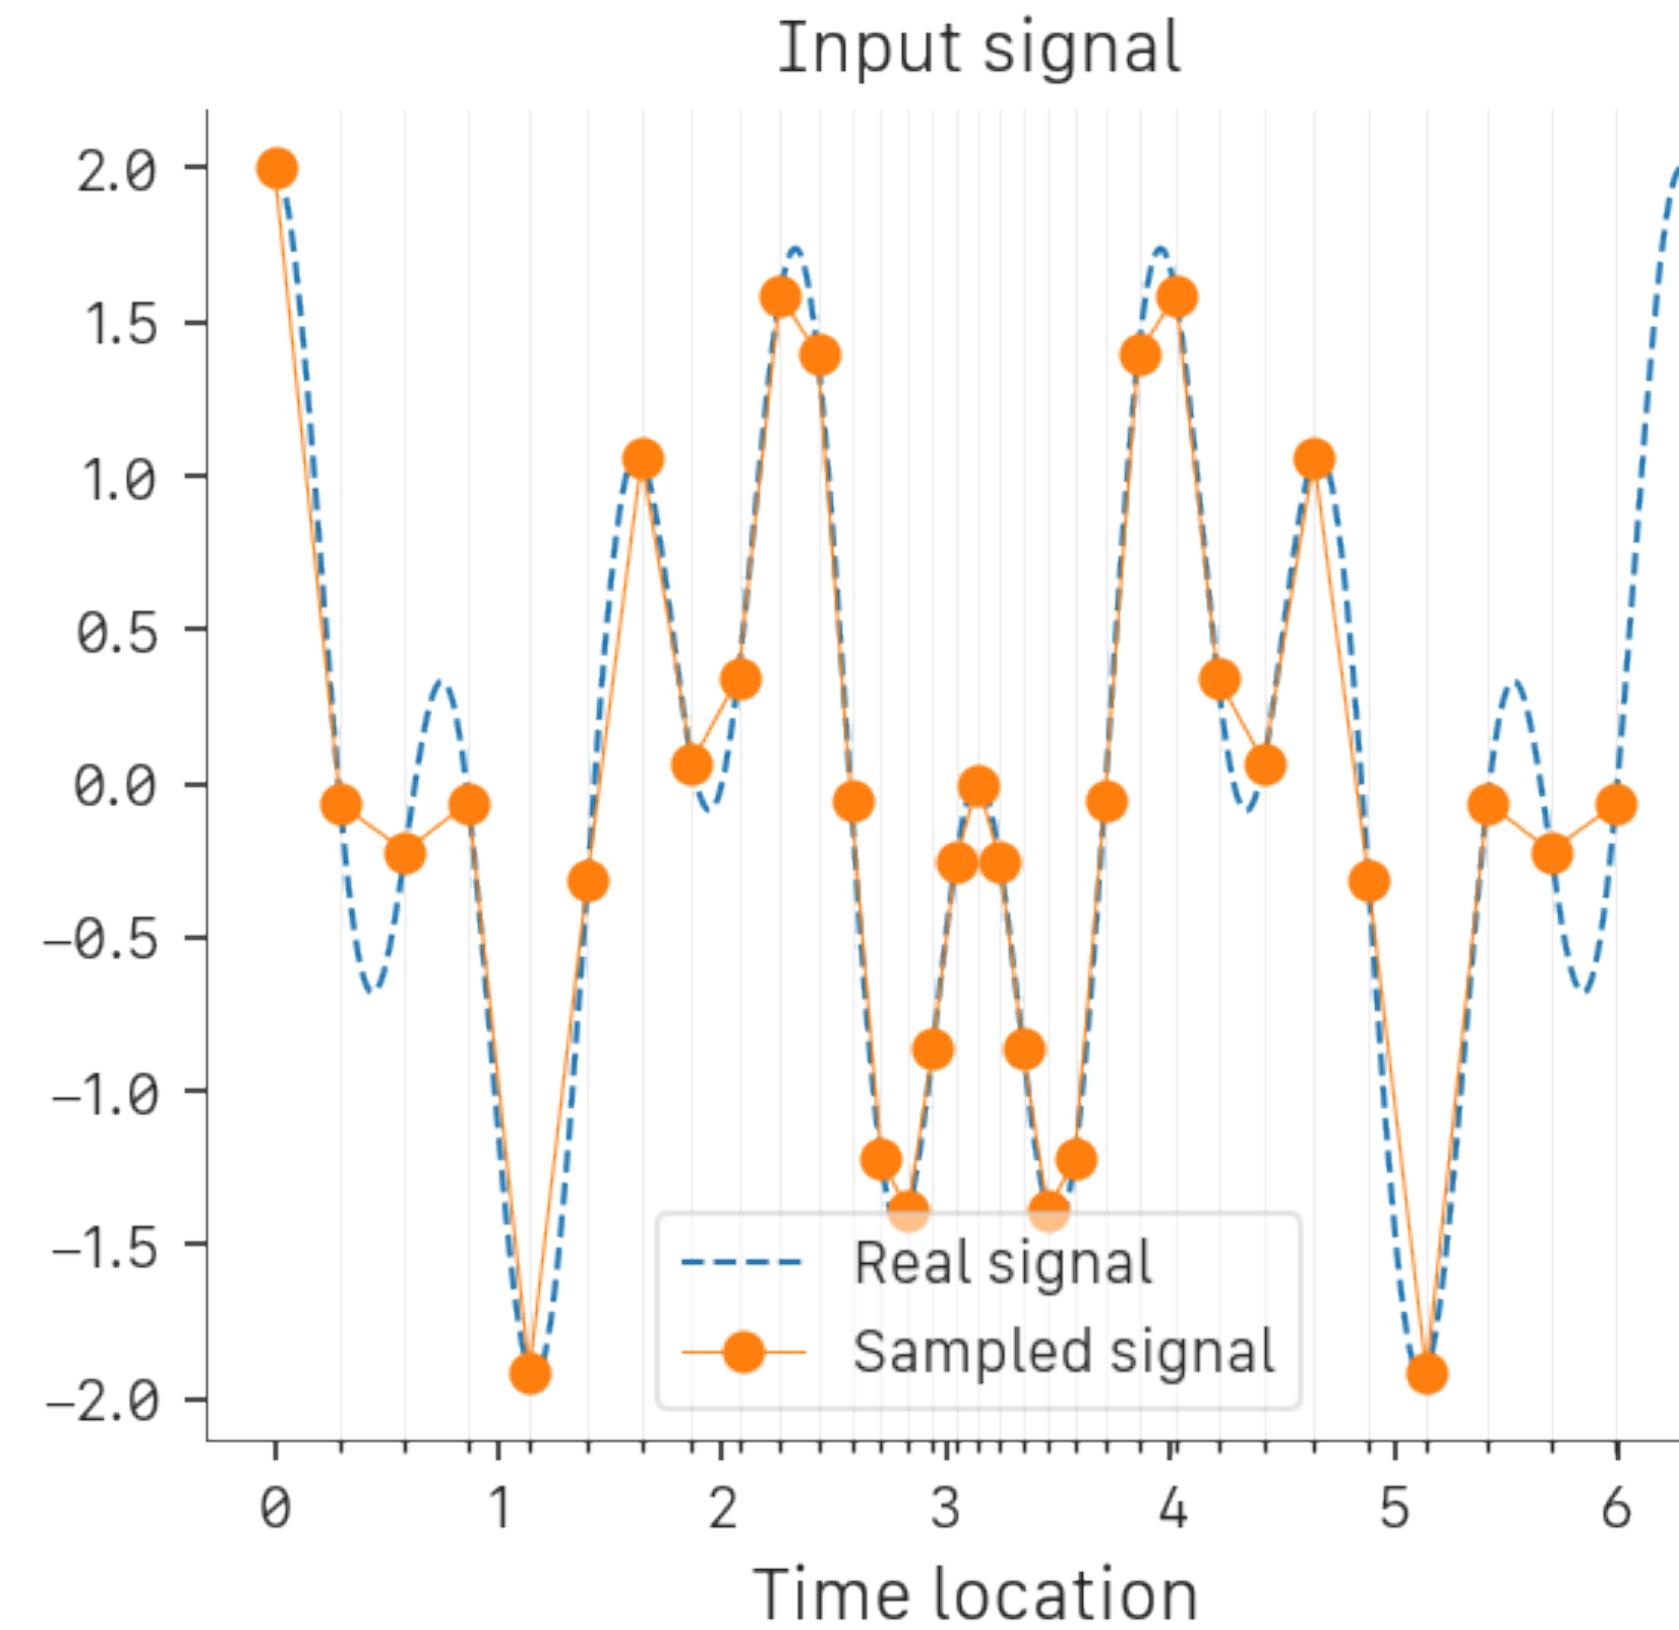
Step 3: Filter and invert

- Design and apply a lowpass filter
 - Use learned bases to invert back to time



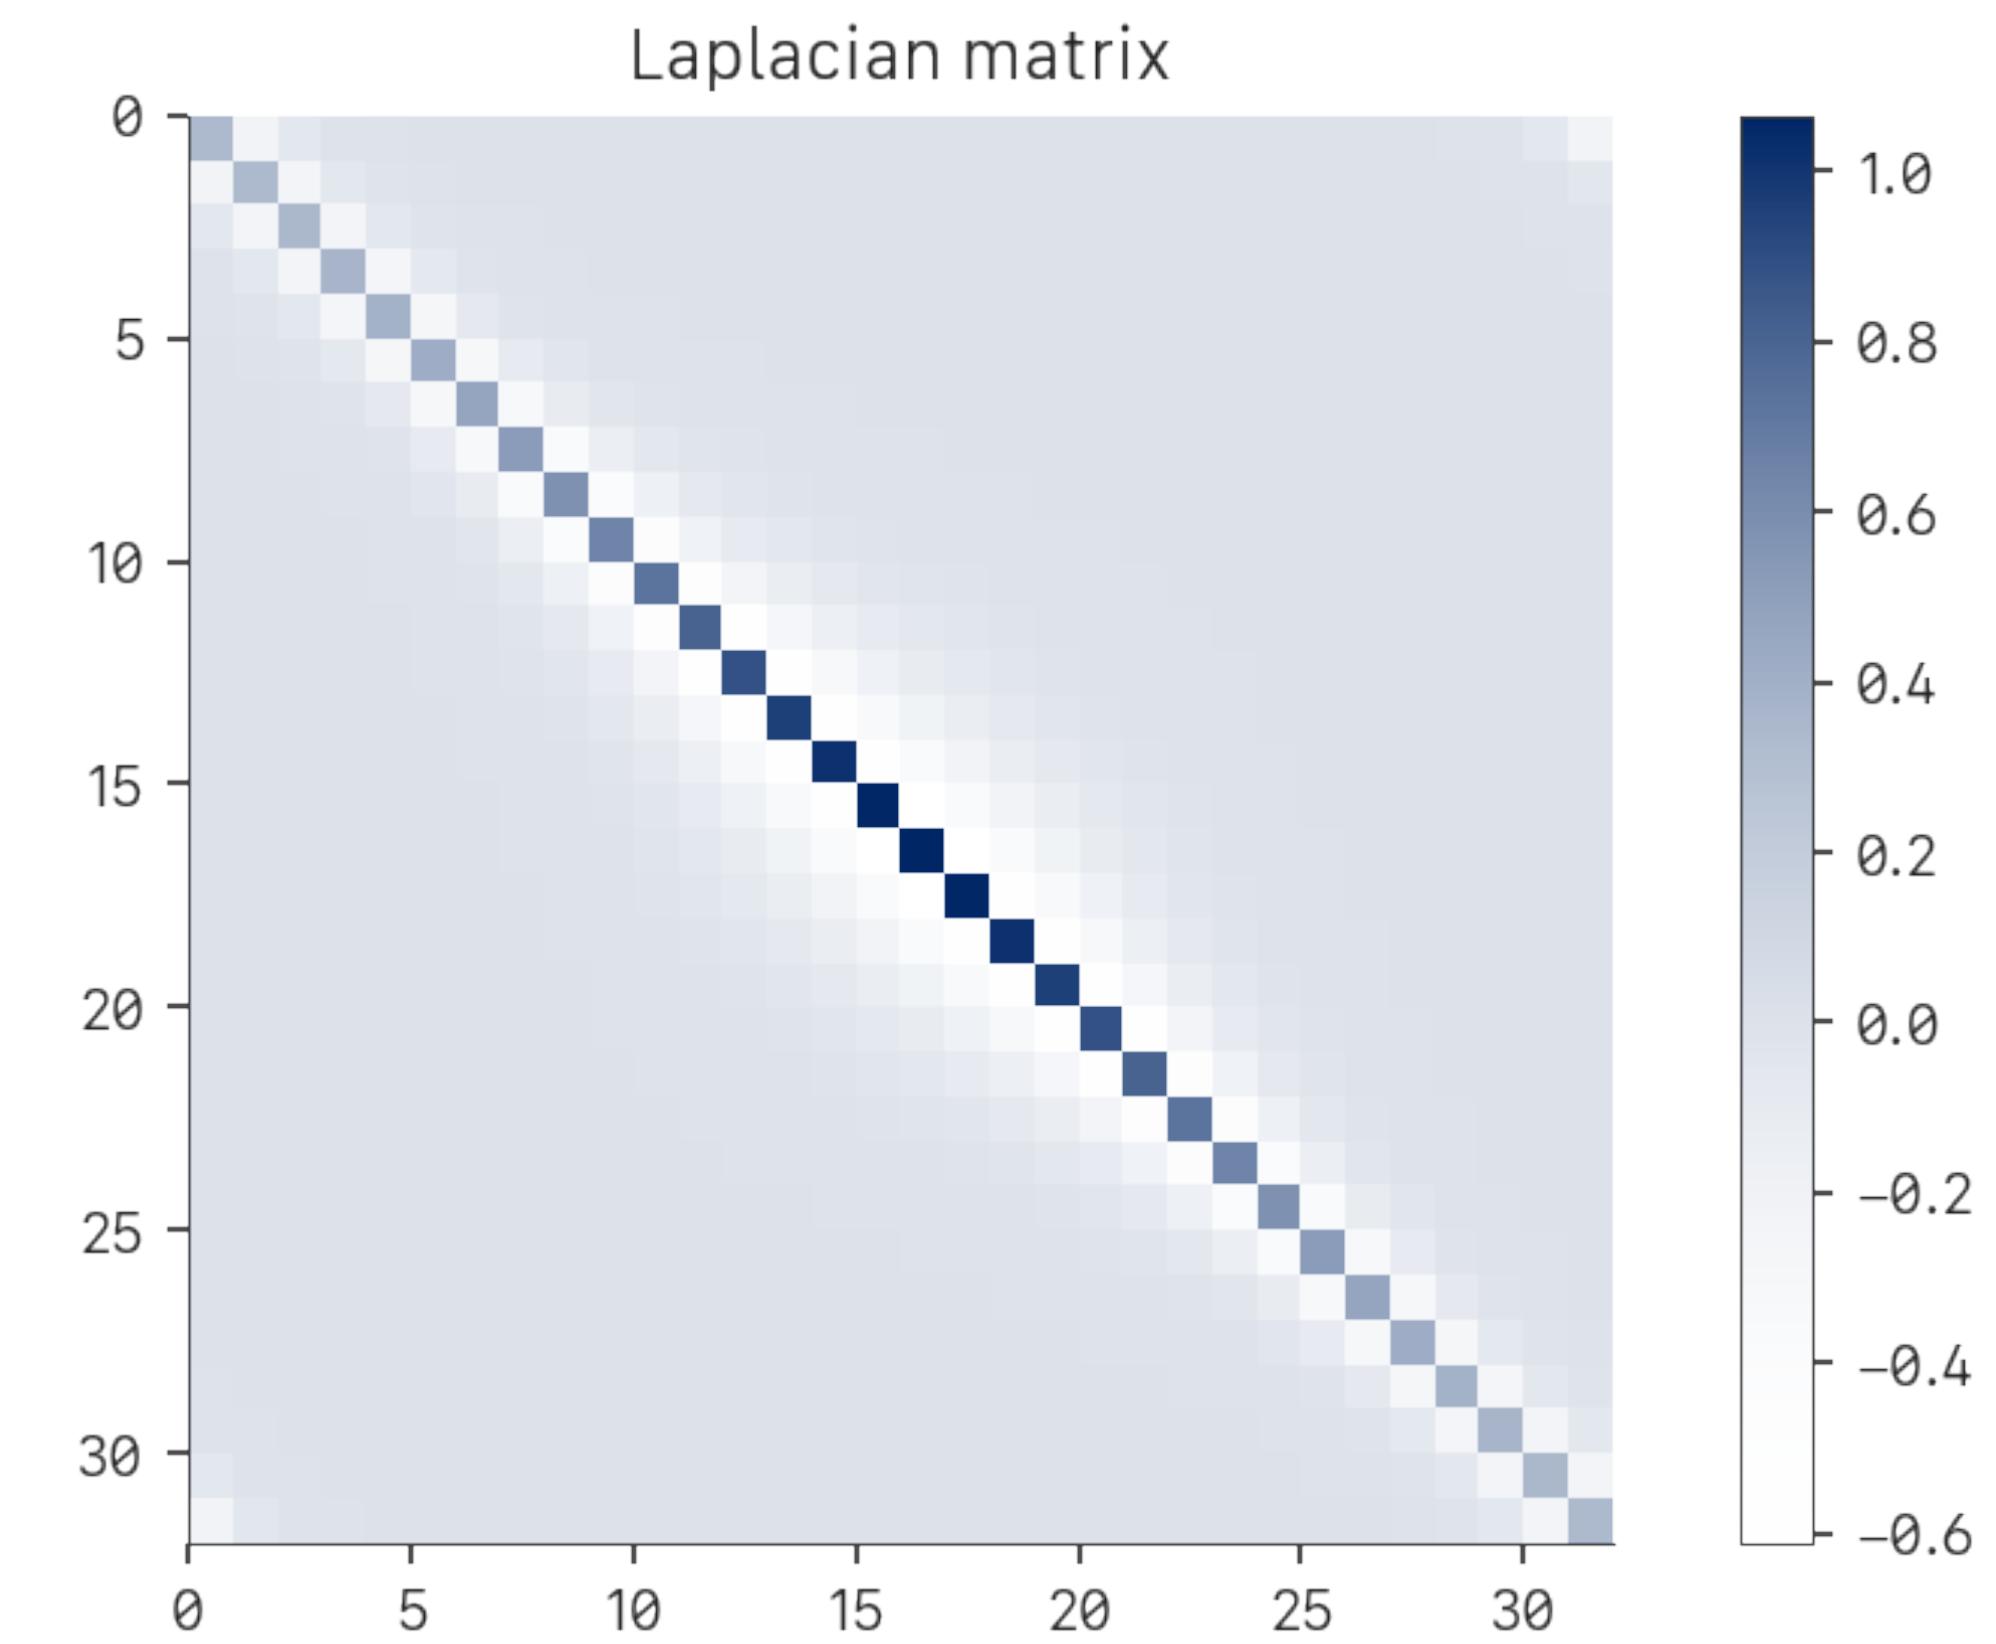
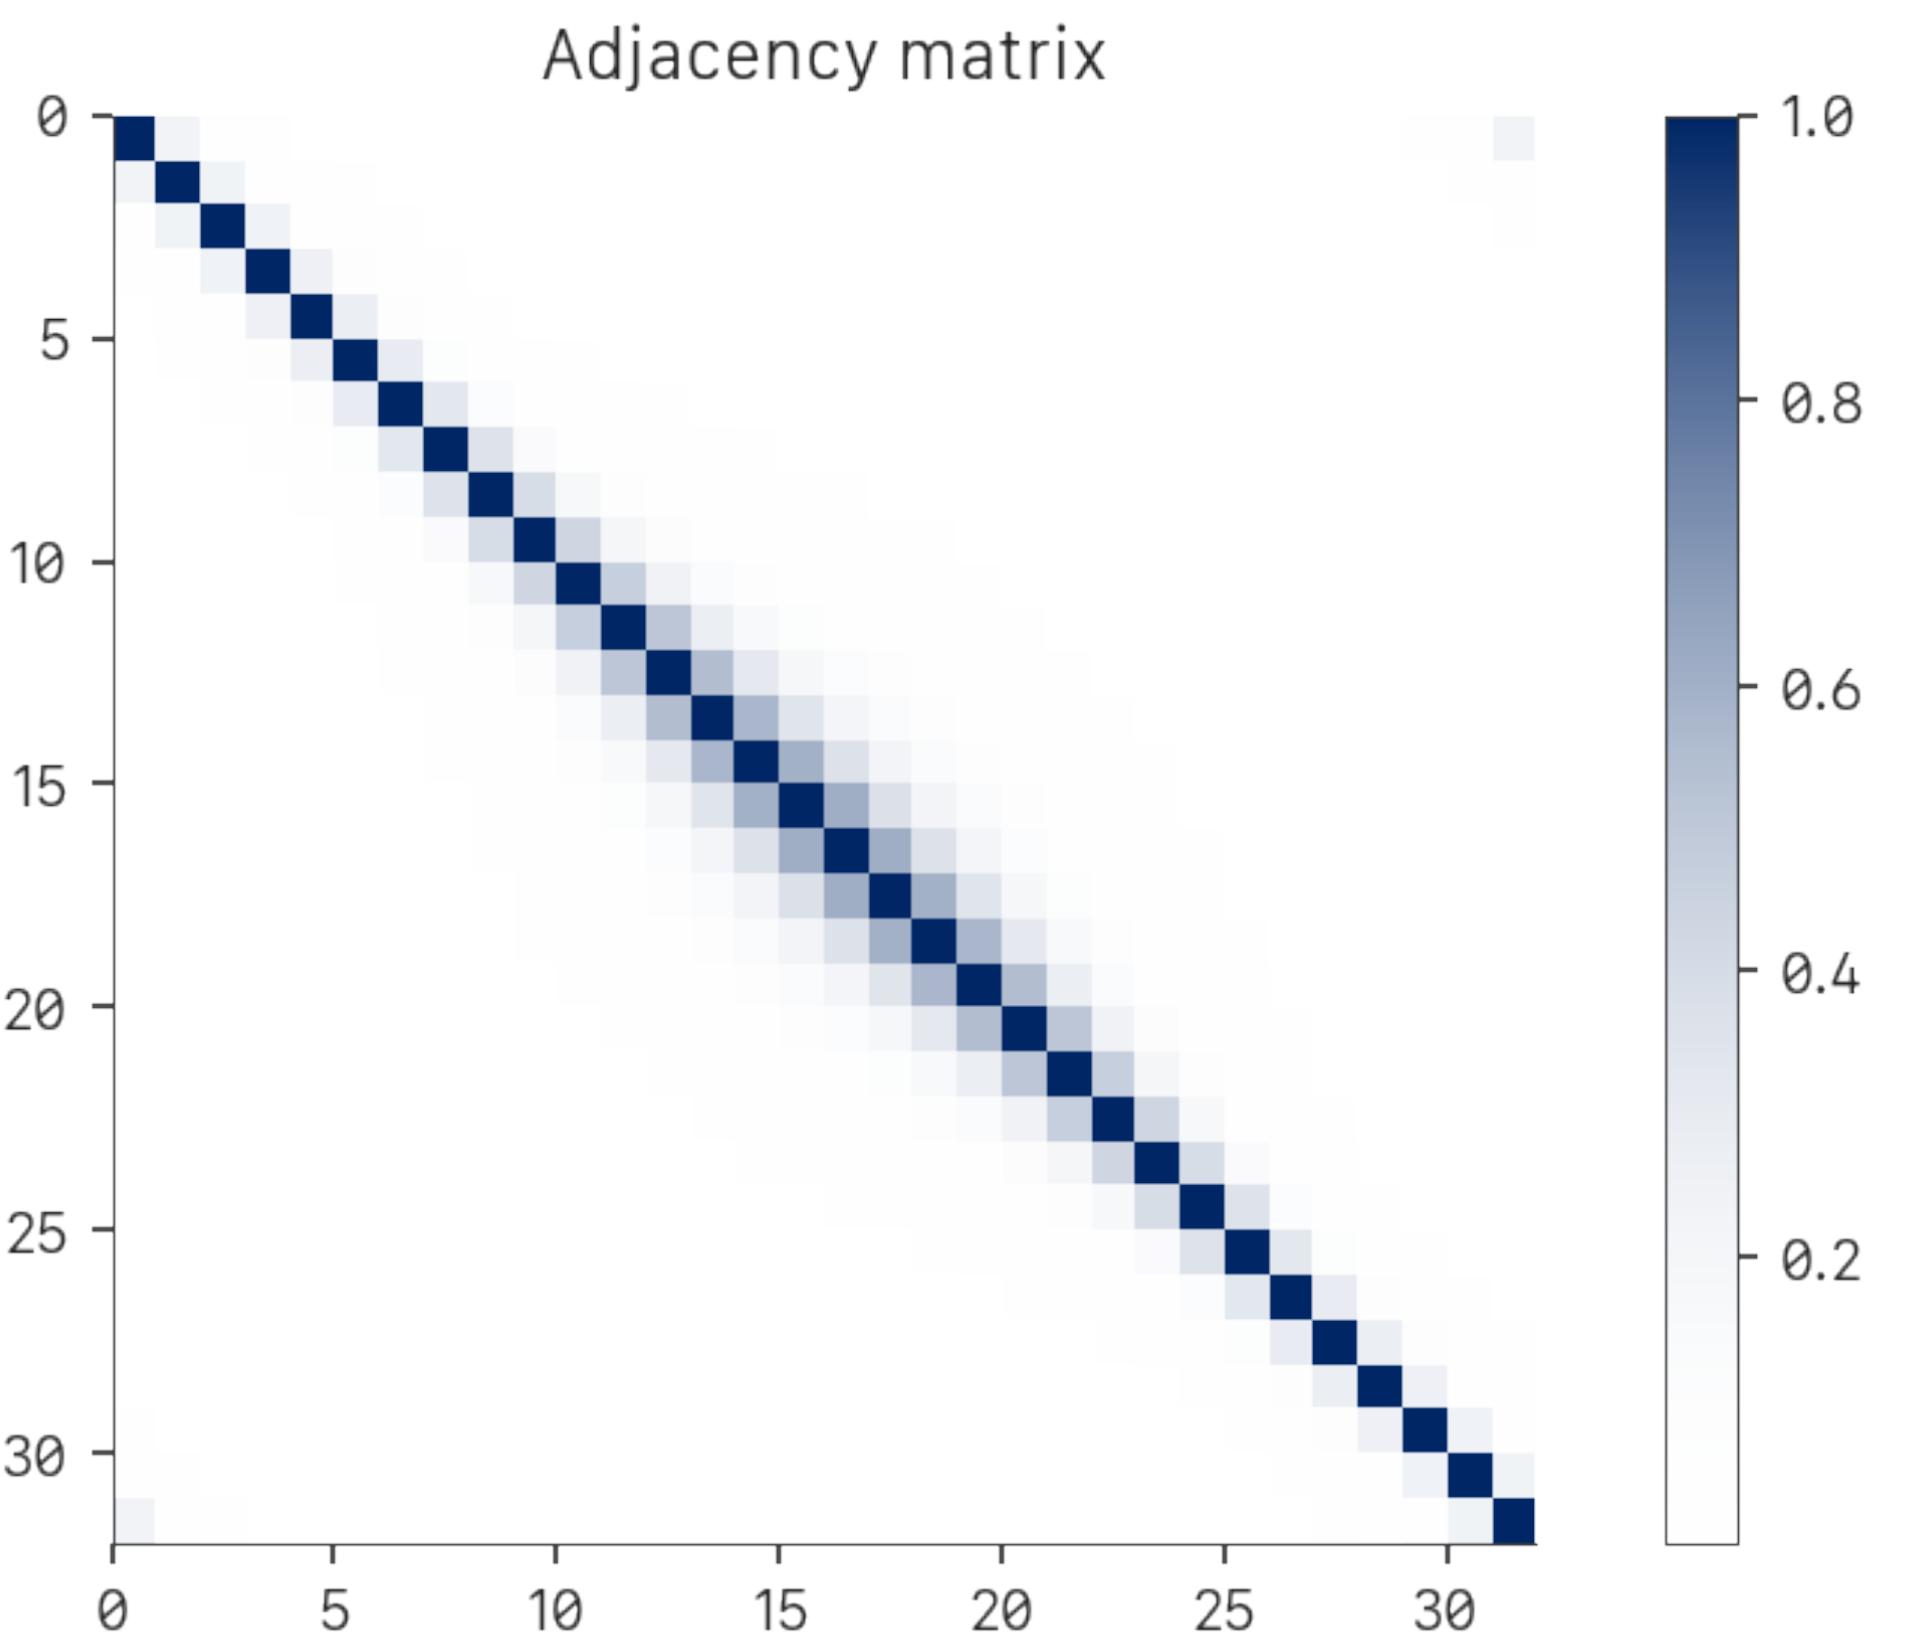
Nothing new so far, but ...

- What if my graph is not as regular?
 - What if the sample rate is variable?



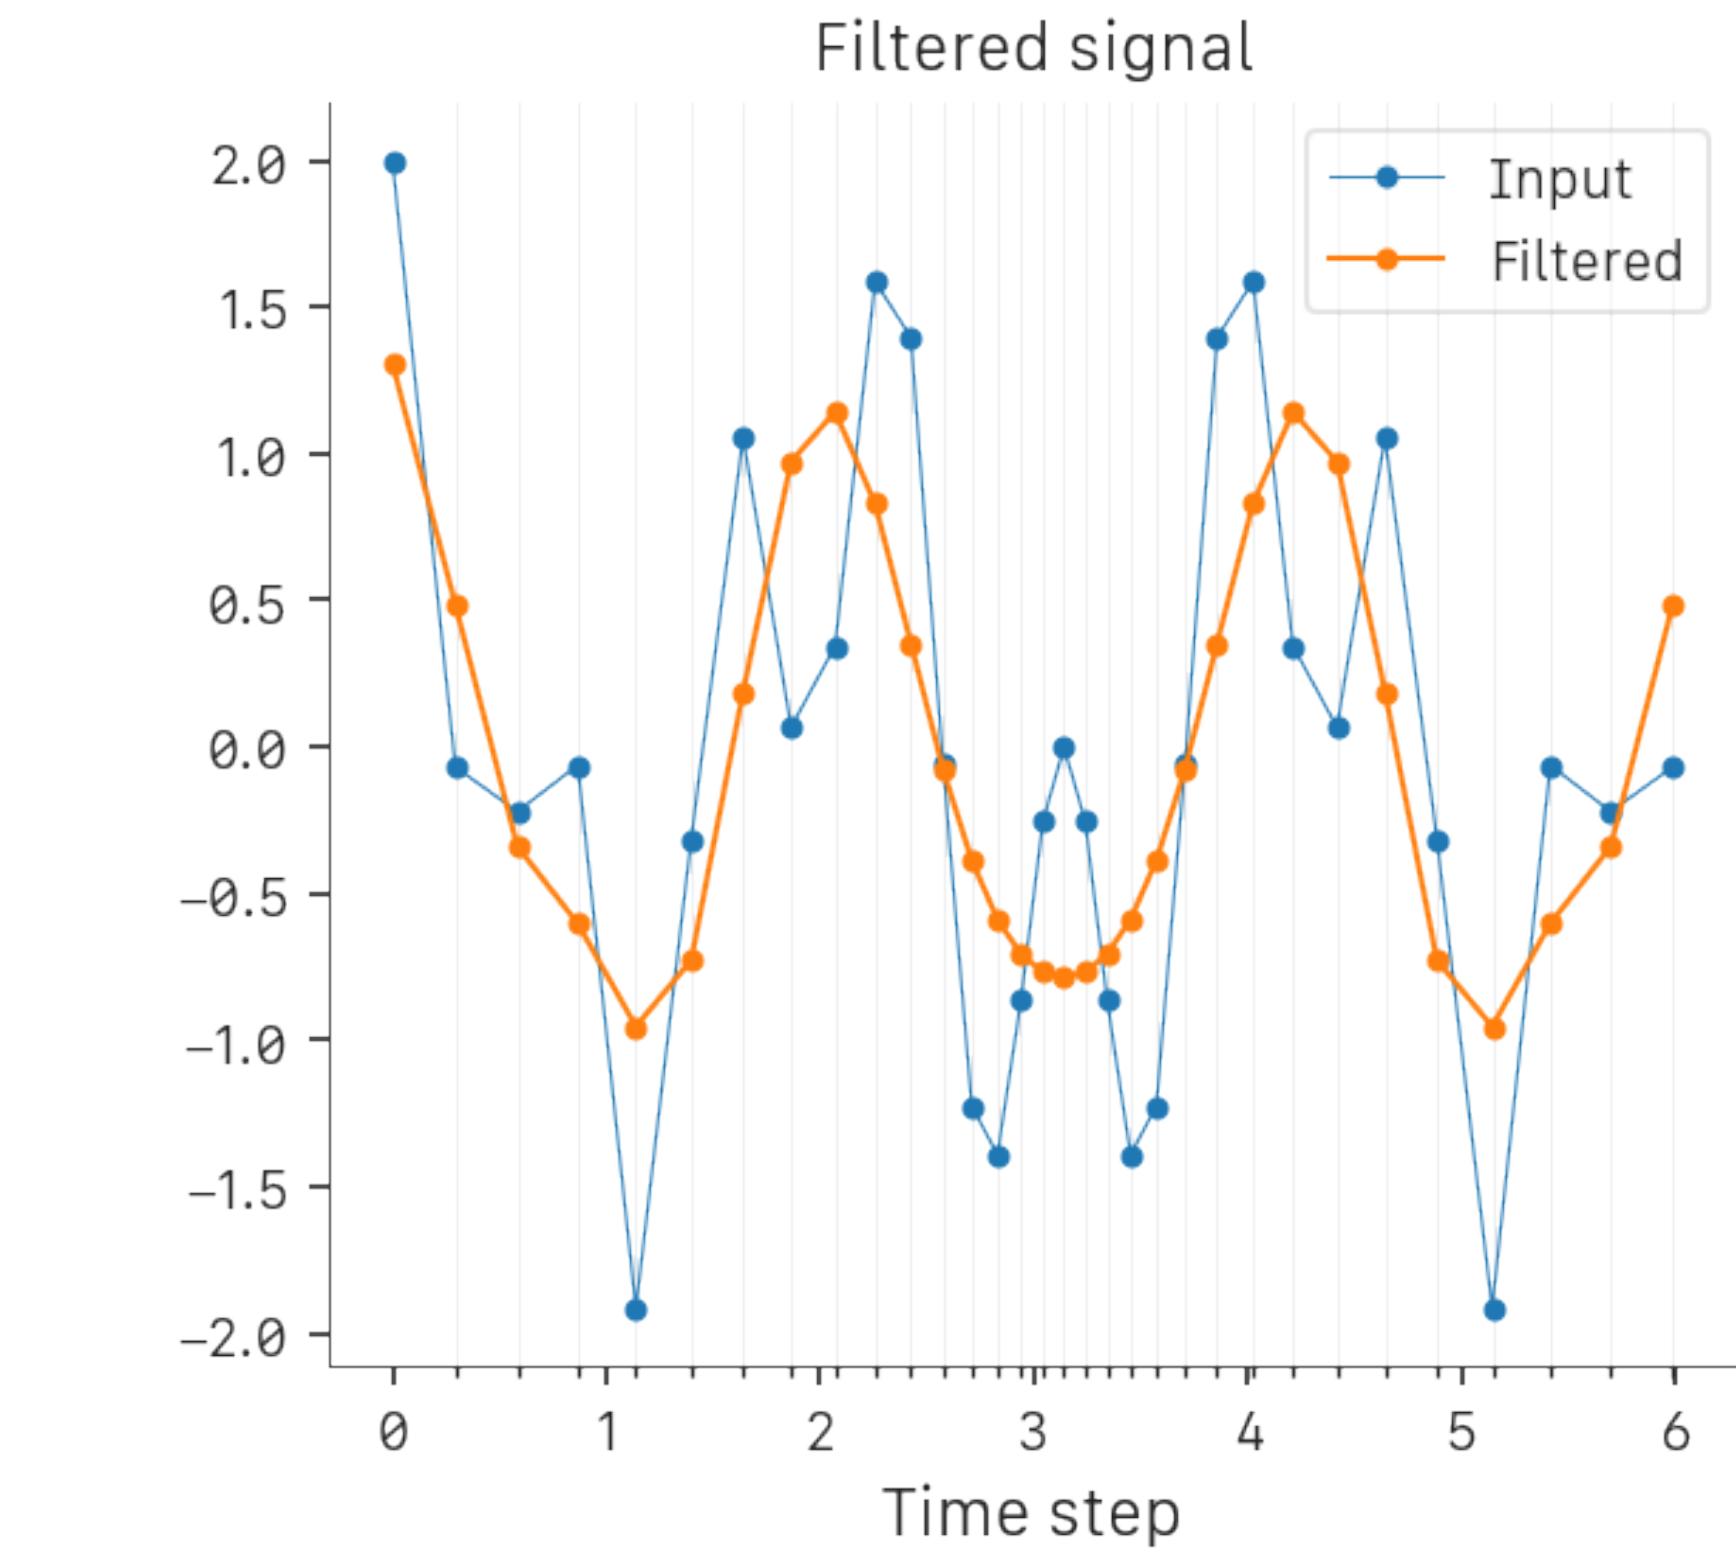
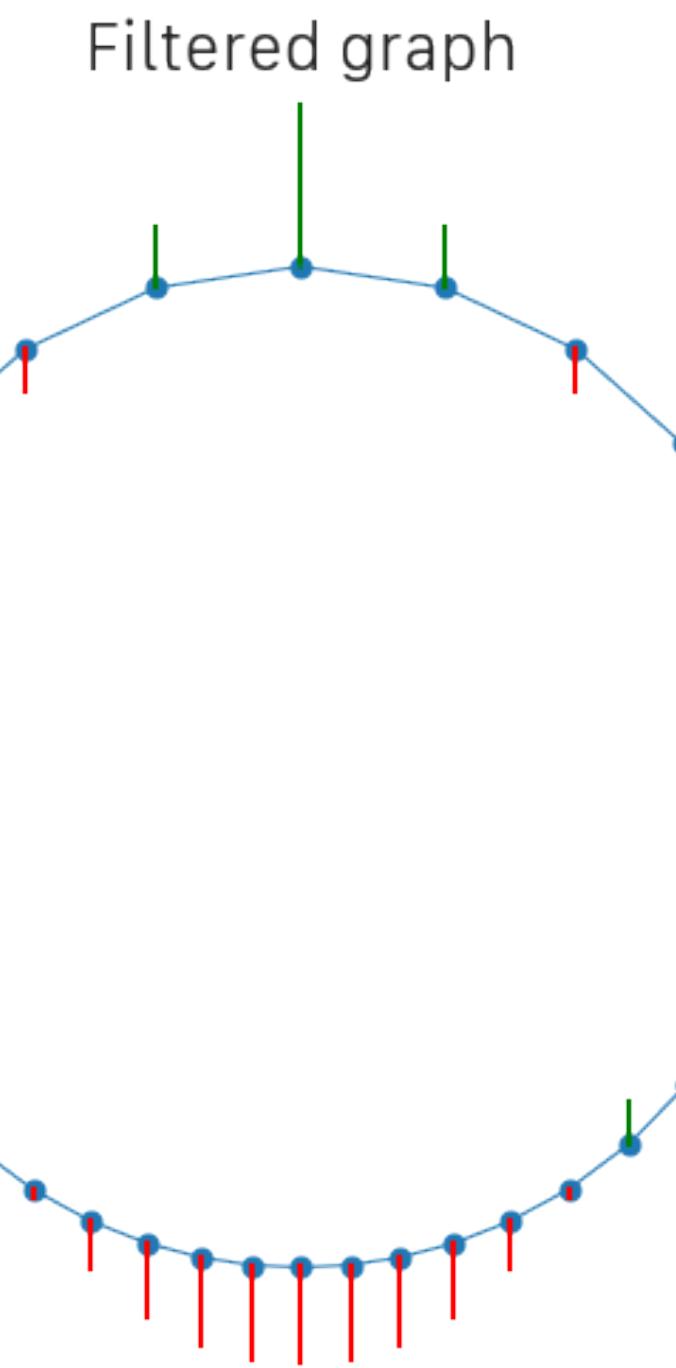
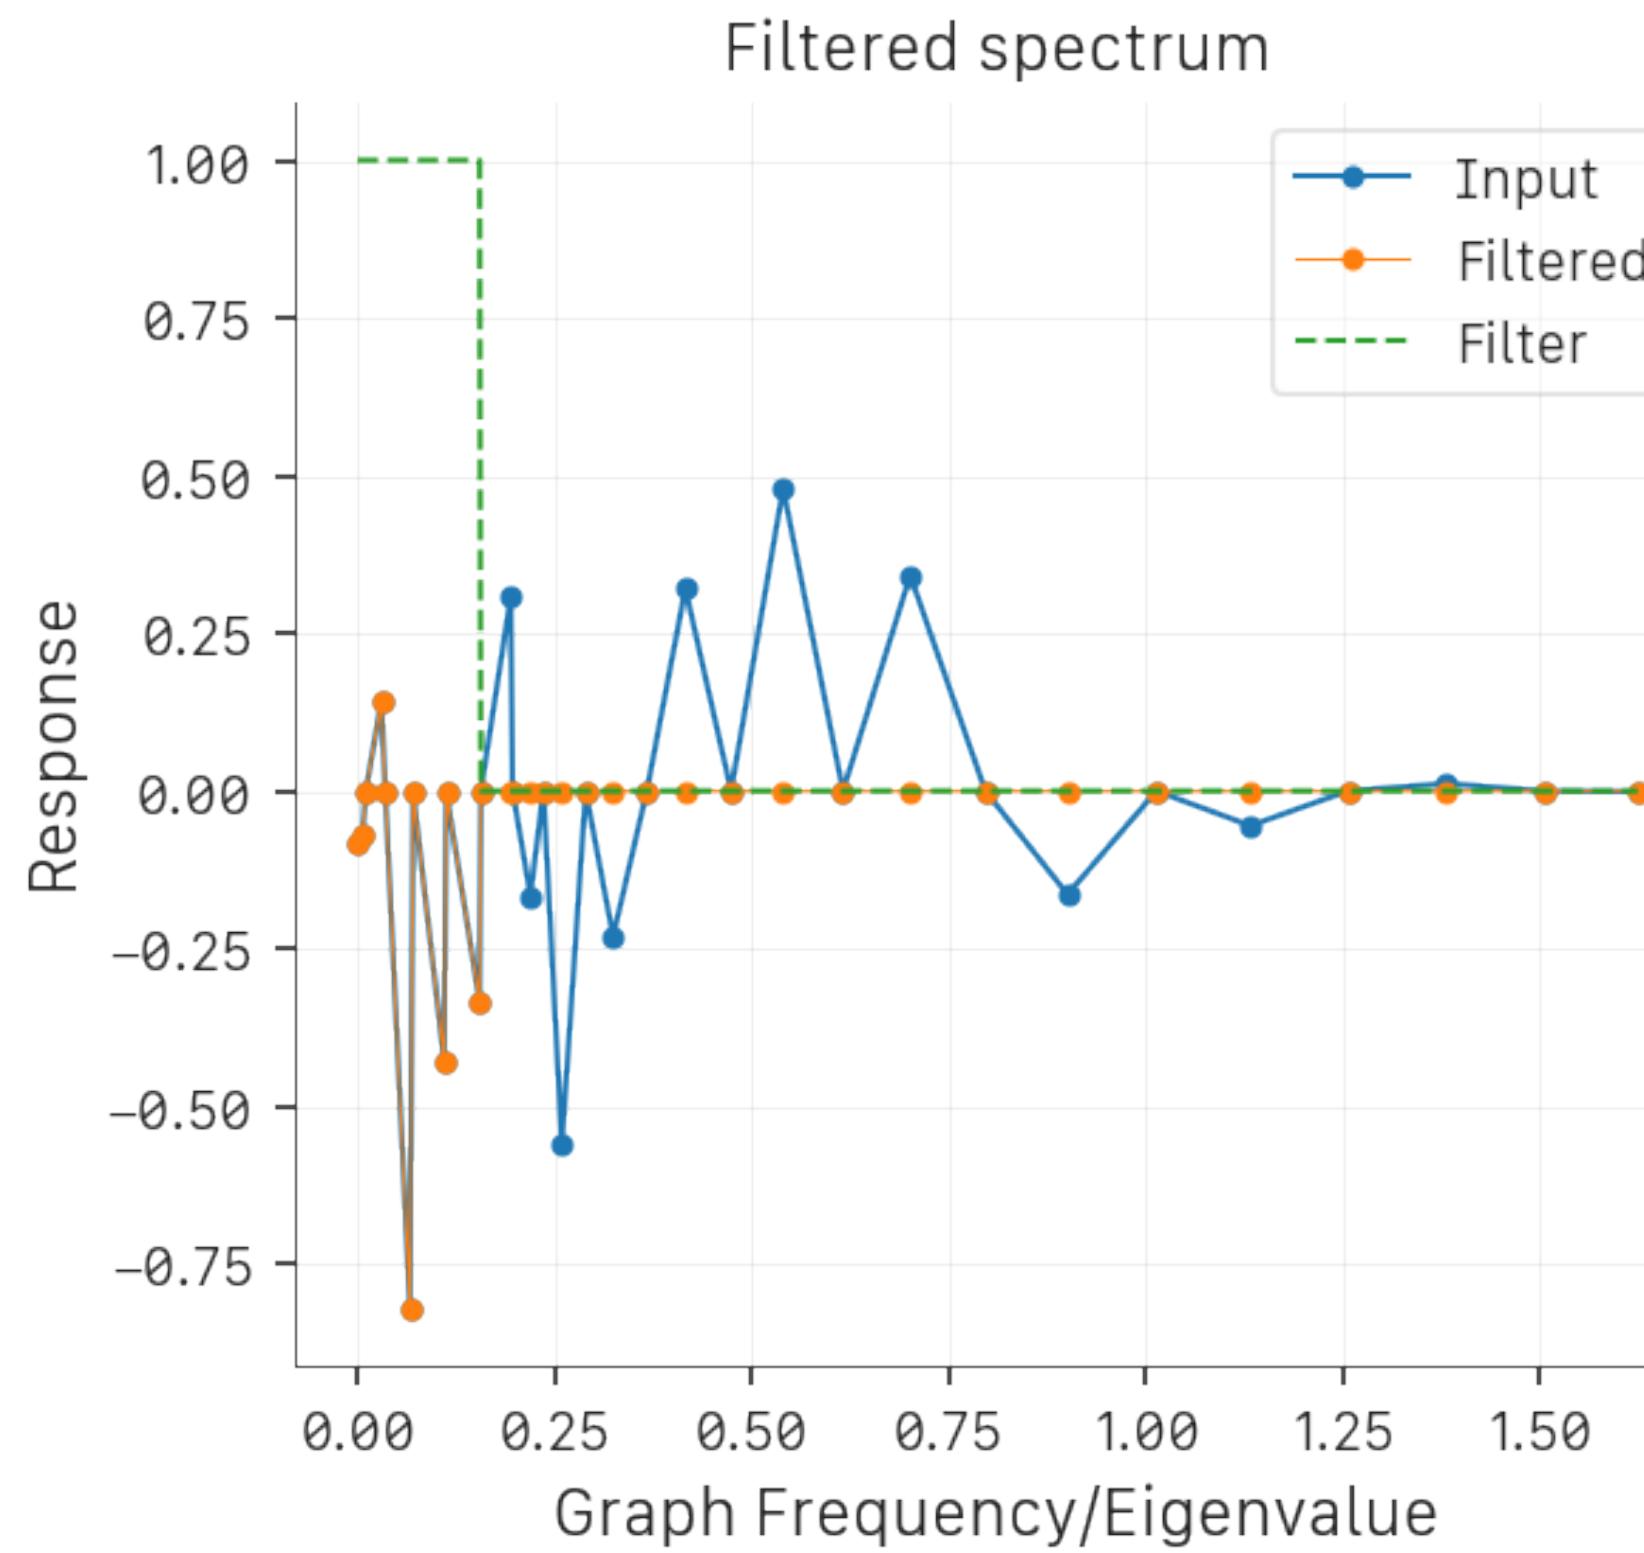
New time structure gets encoded in graph

- Nodes closer in time have a stronger edge weight



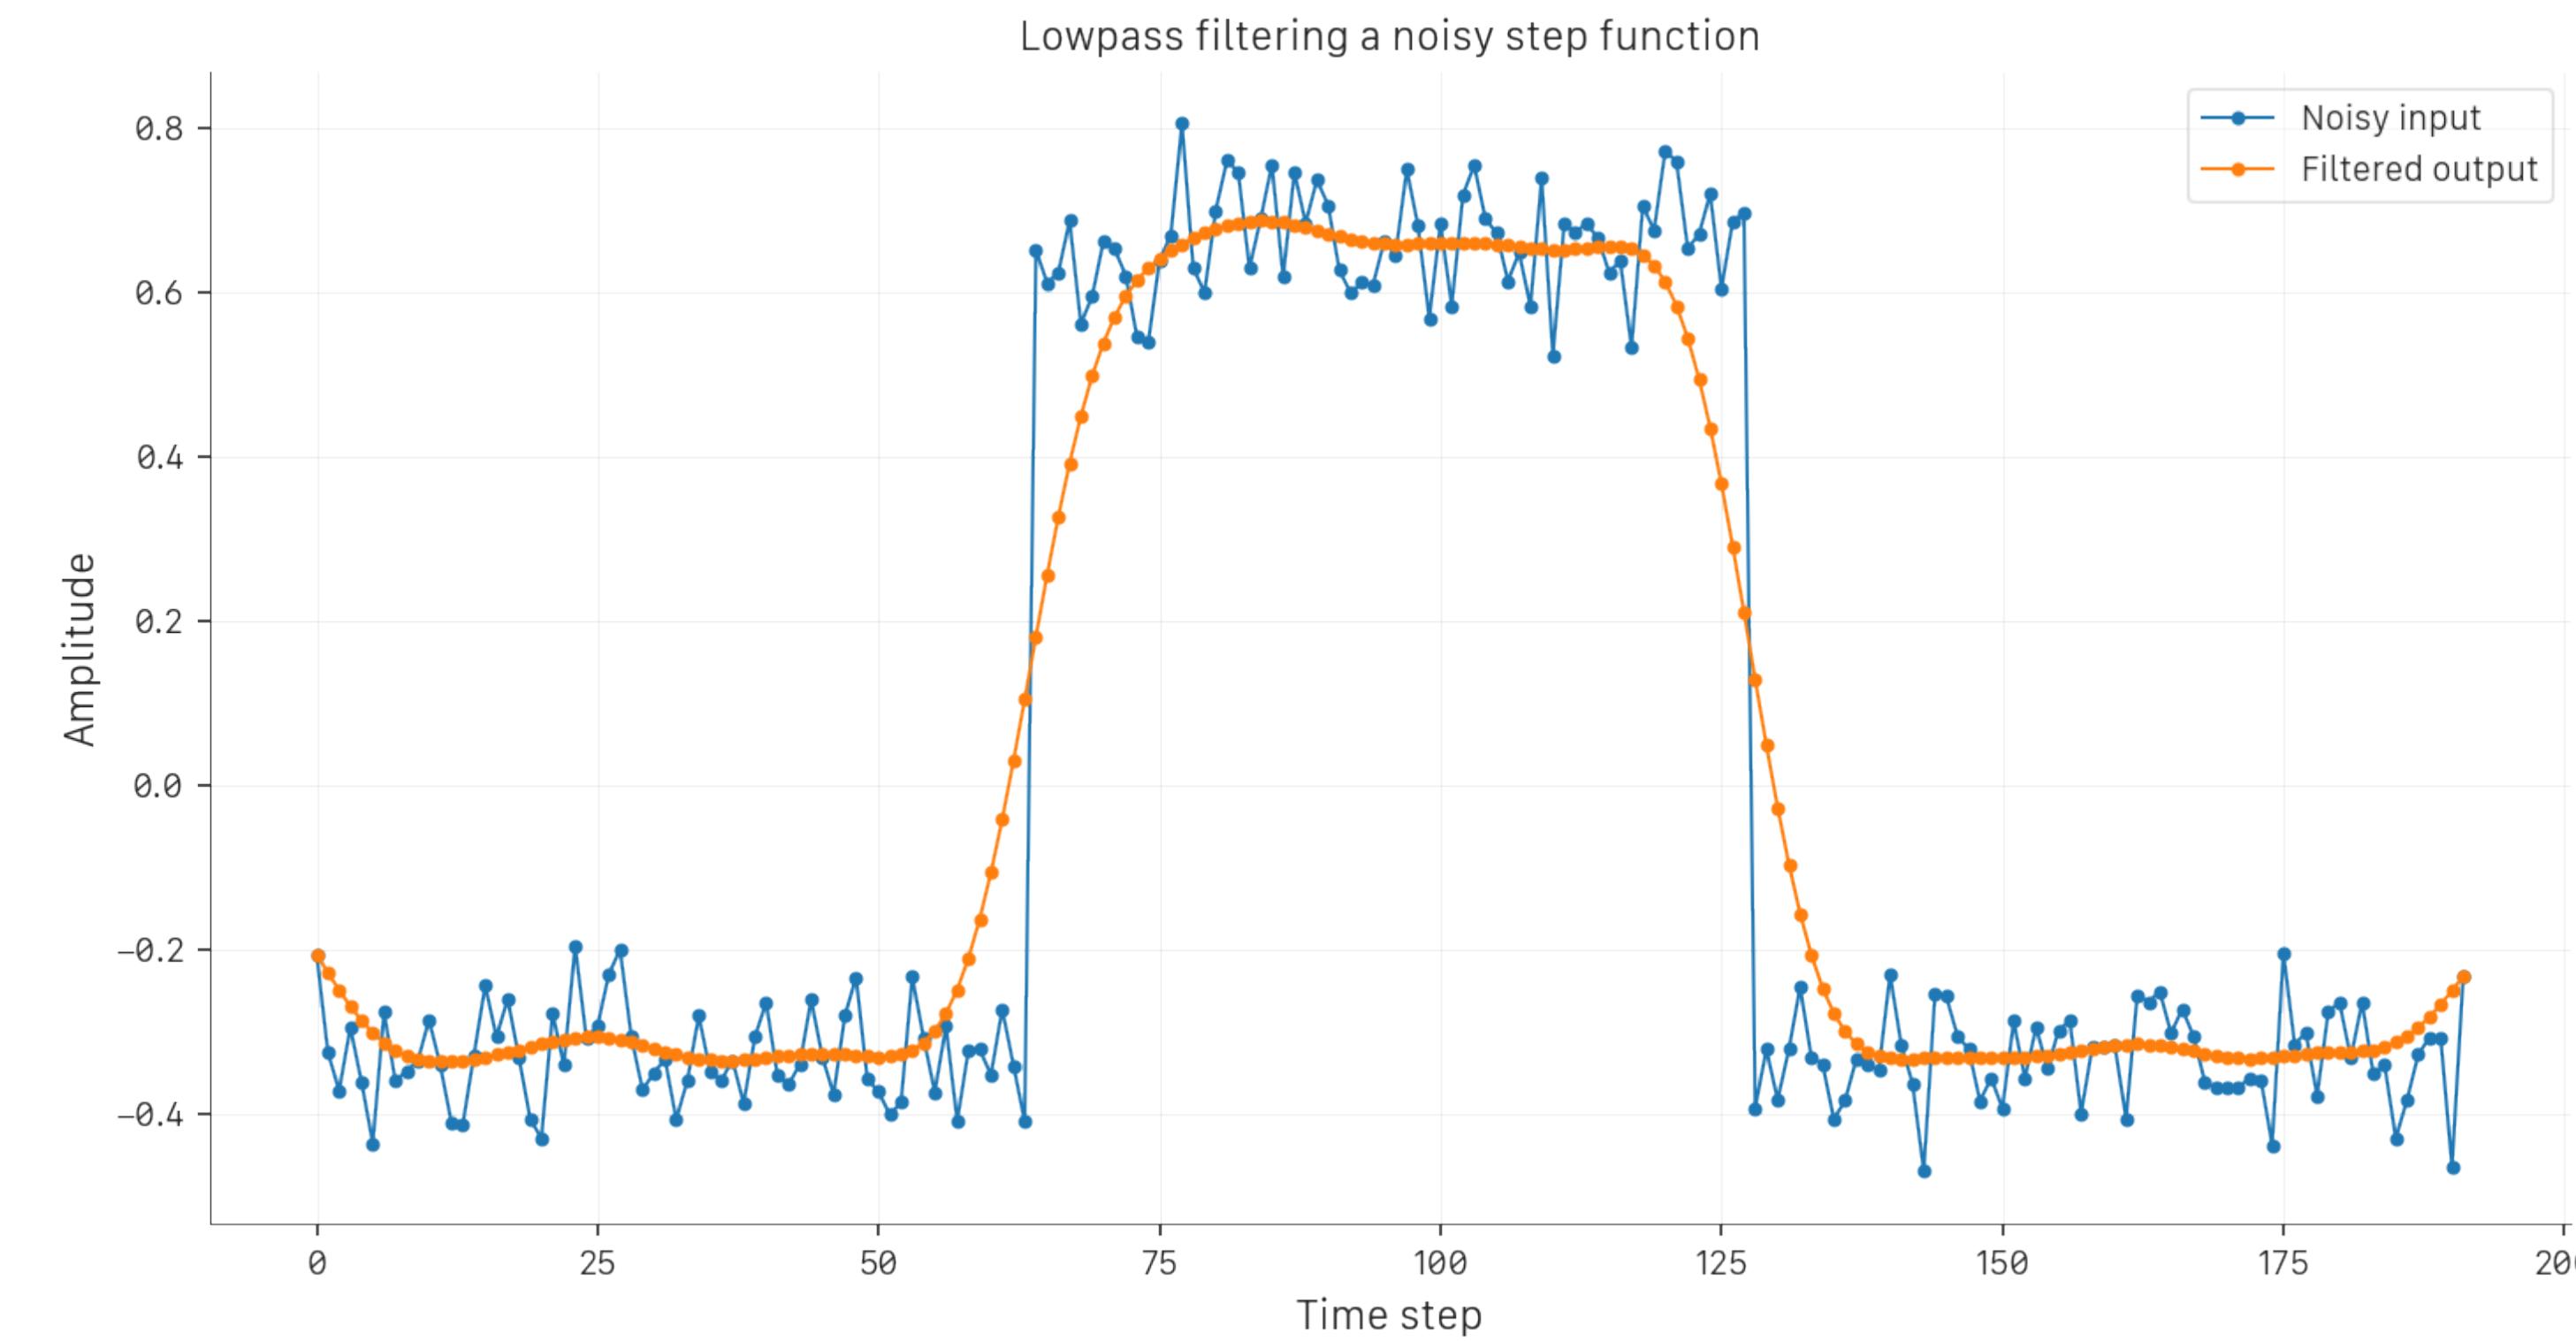
We can still apply a lowpass filter

- Using the same process as before
 - But now we can deal with the strange sample rate variation



Encoding context from a signal

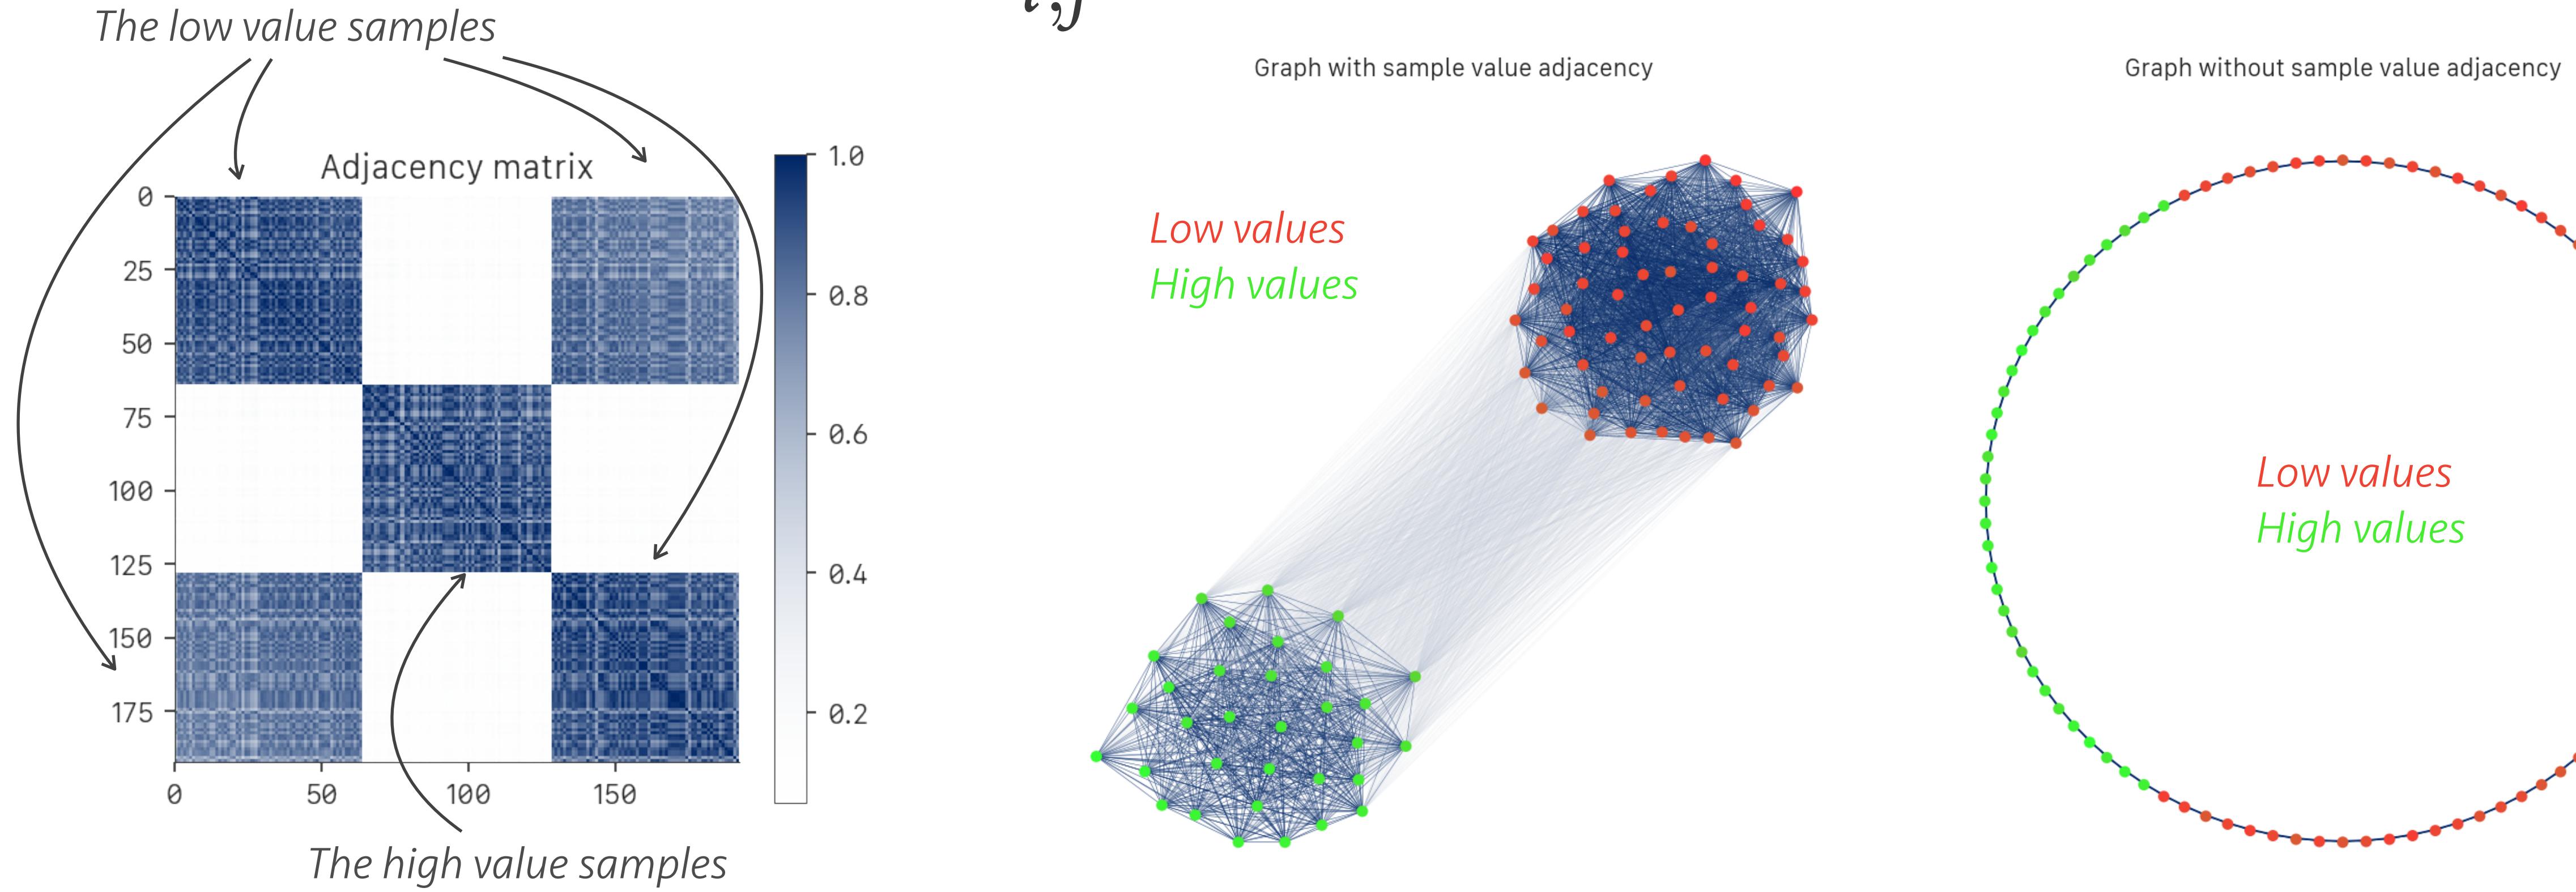
- Consider a noisy input with steep edges
 - The jumps are treated as high-frequency components
 - Filtering smooths out the jumps, altering the original signal



Using the signal in the adjacency

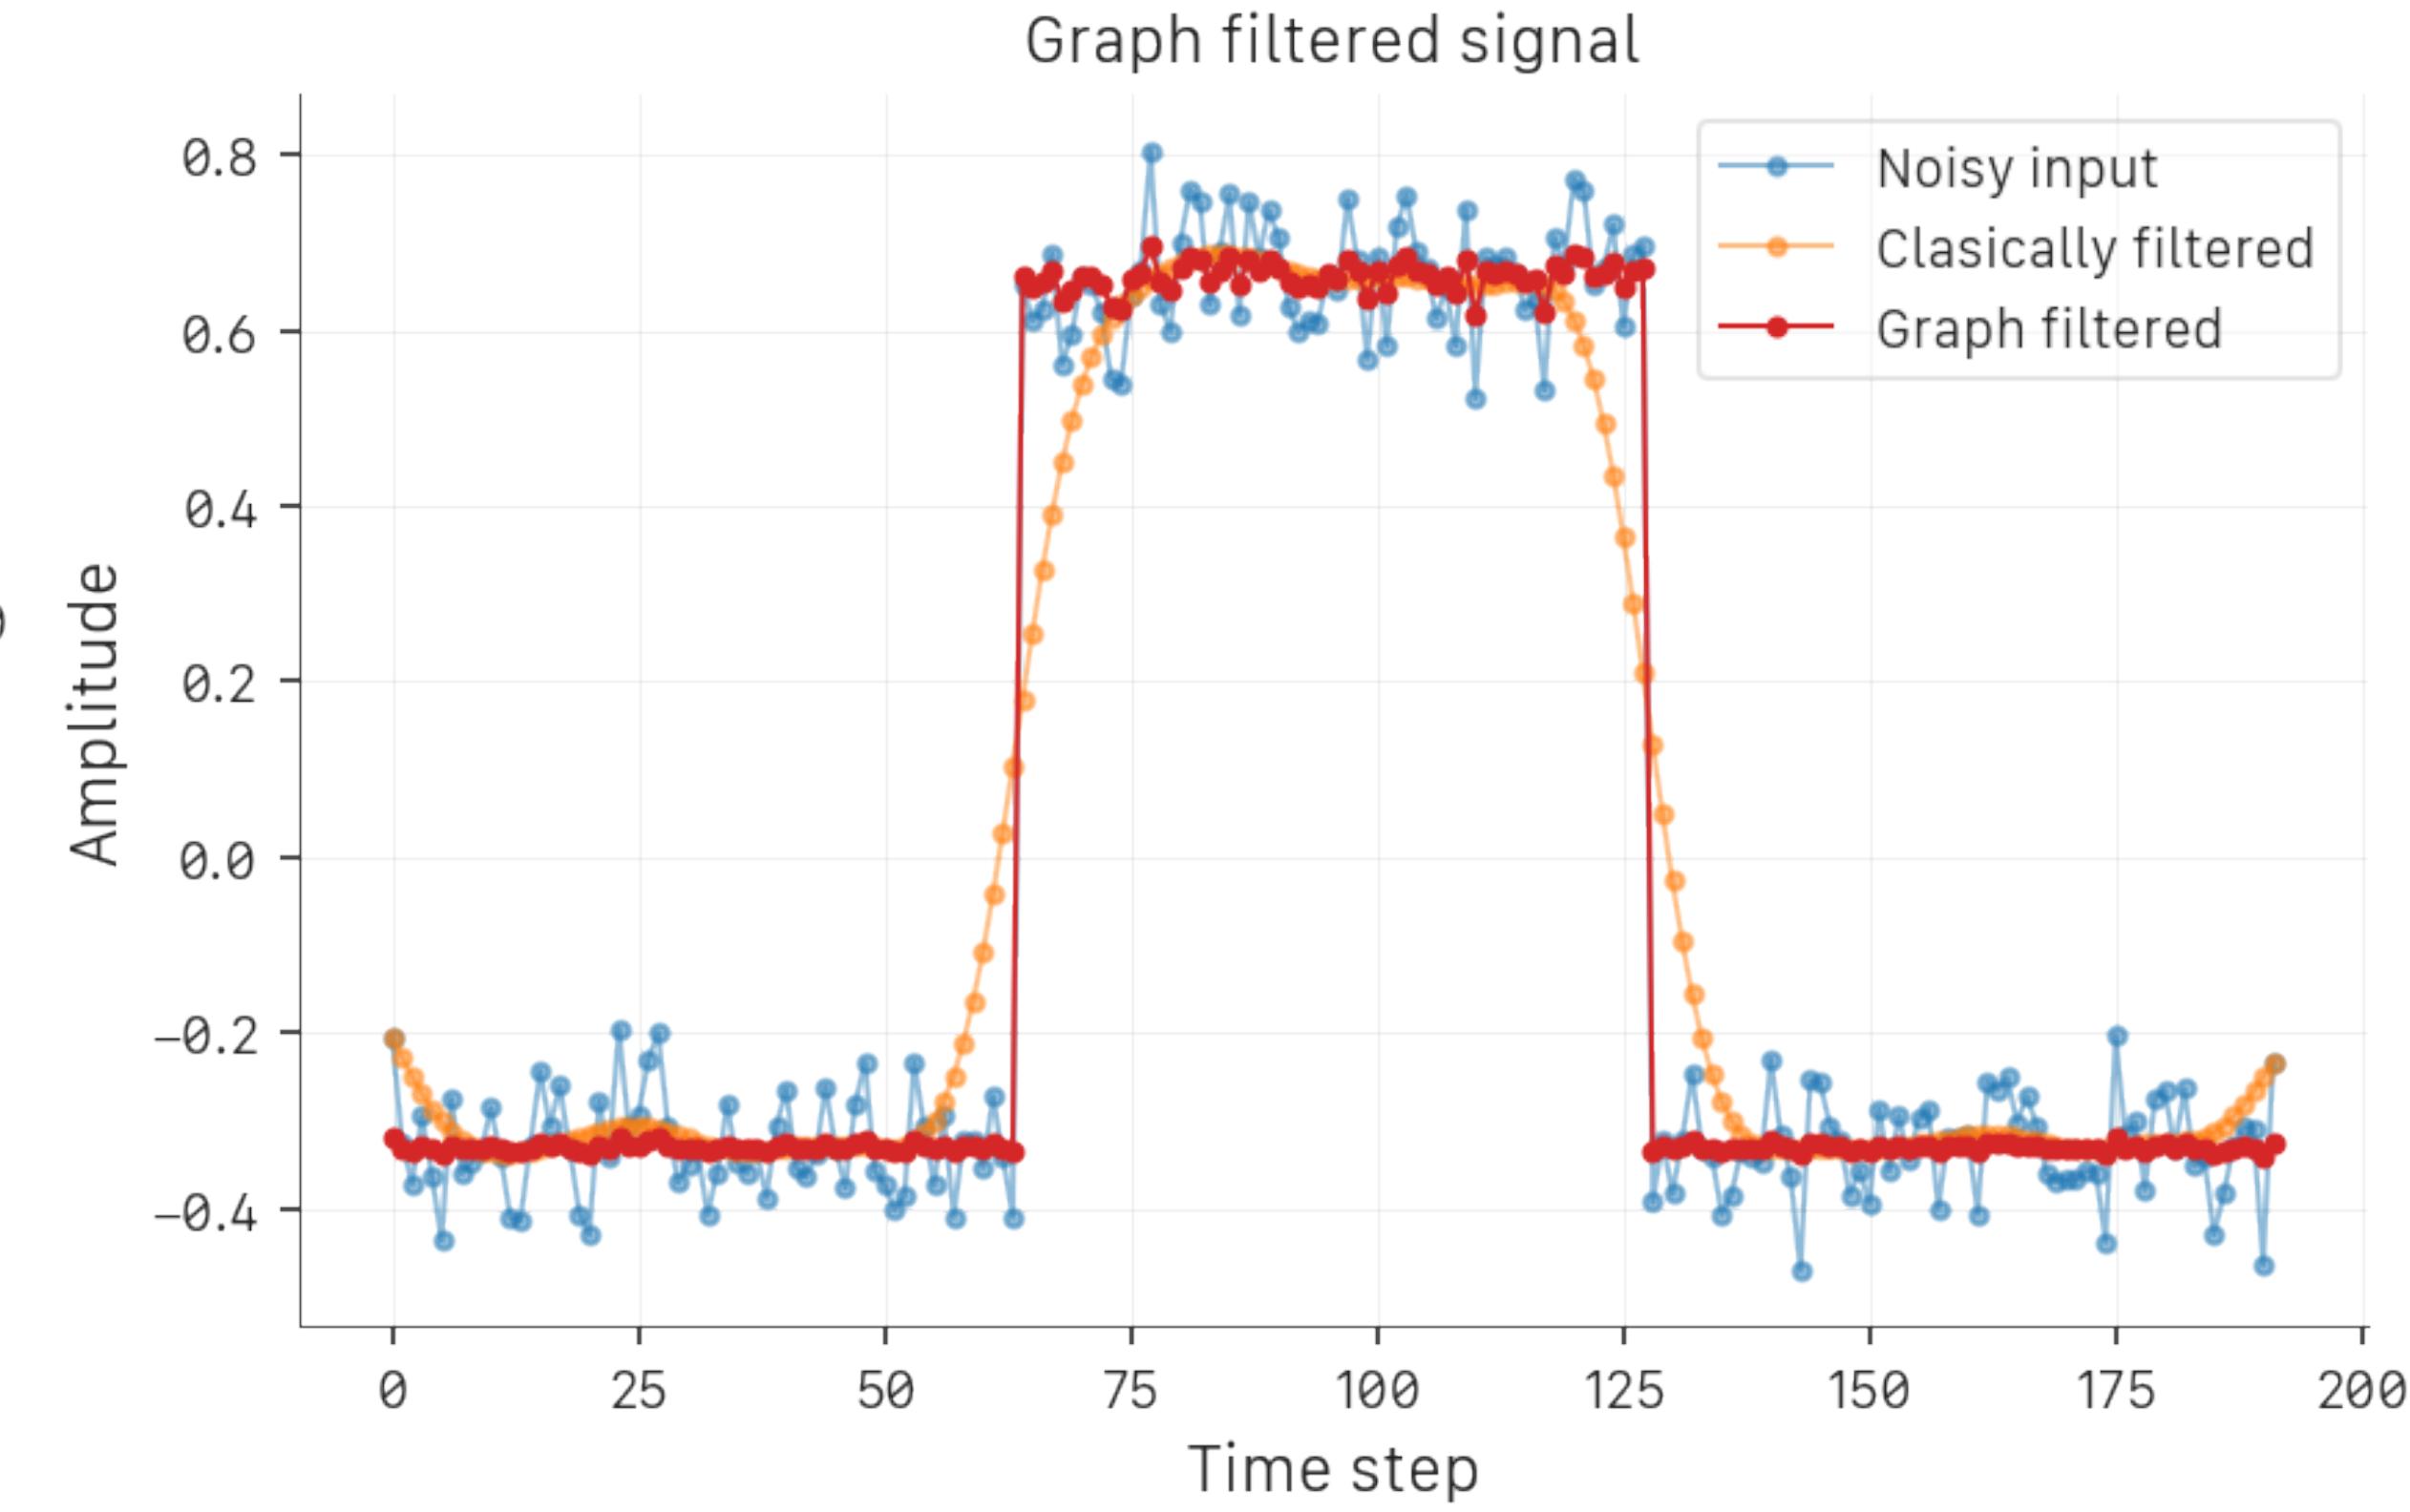
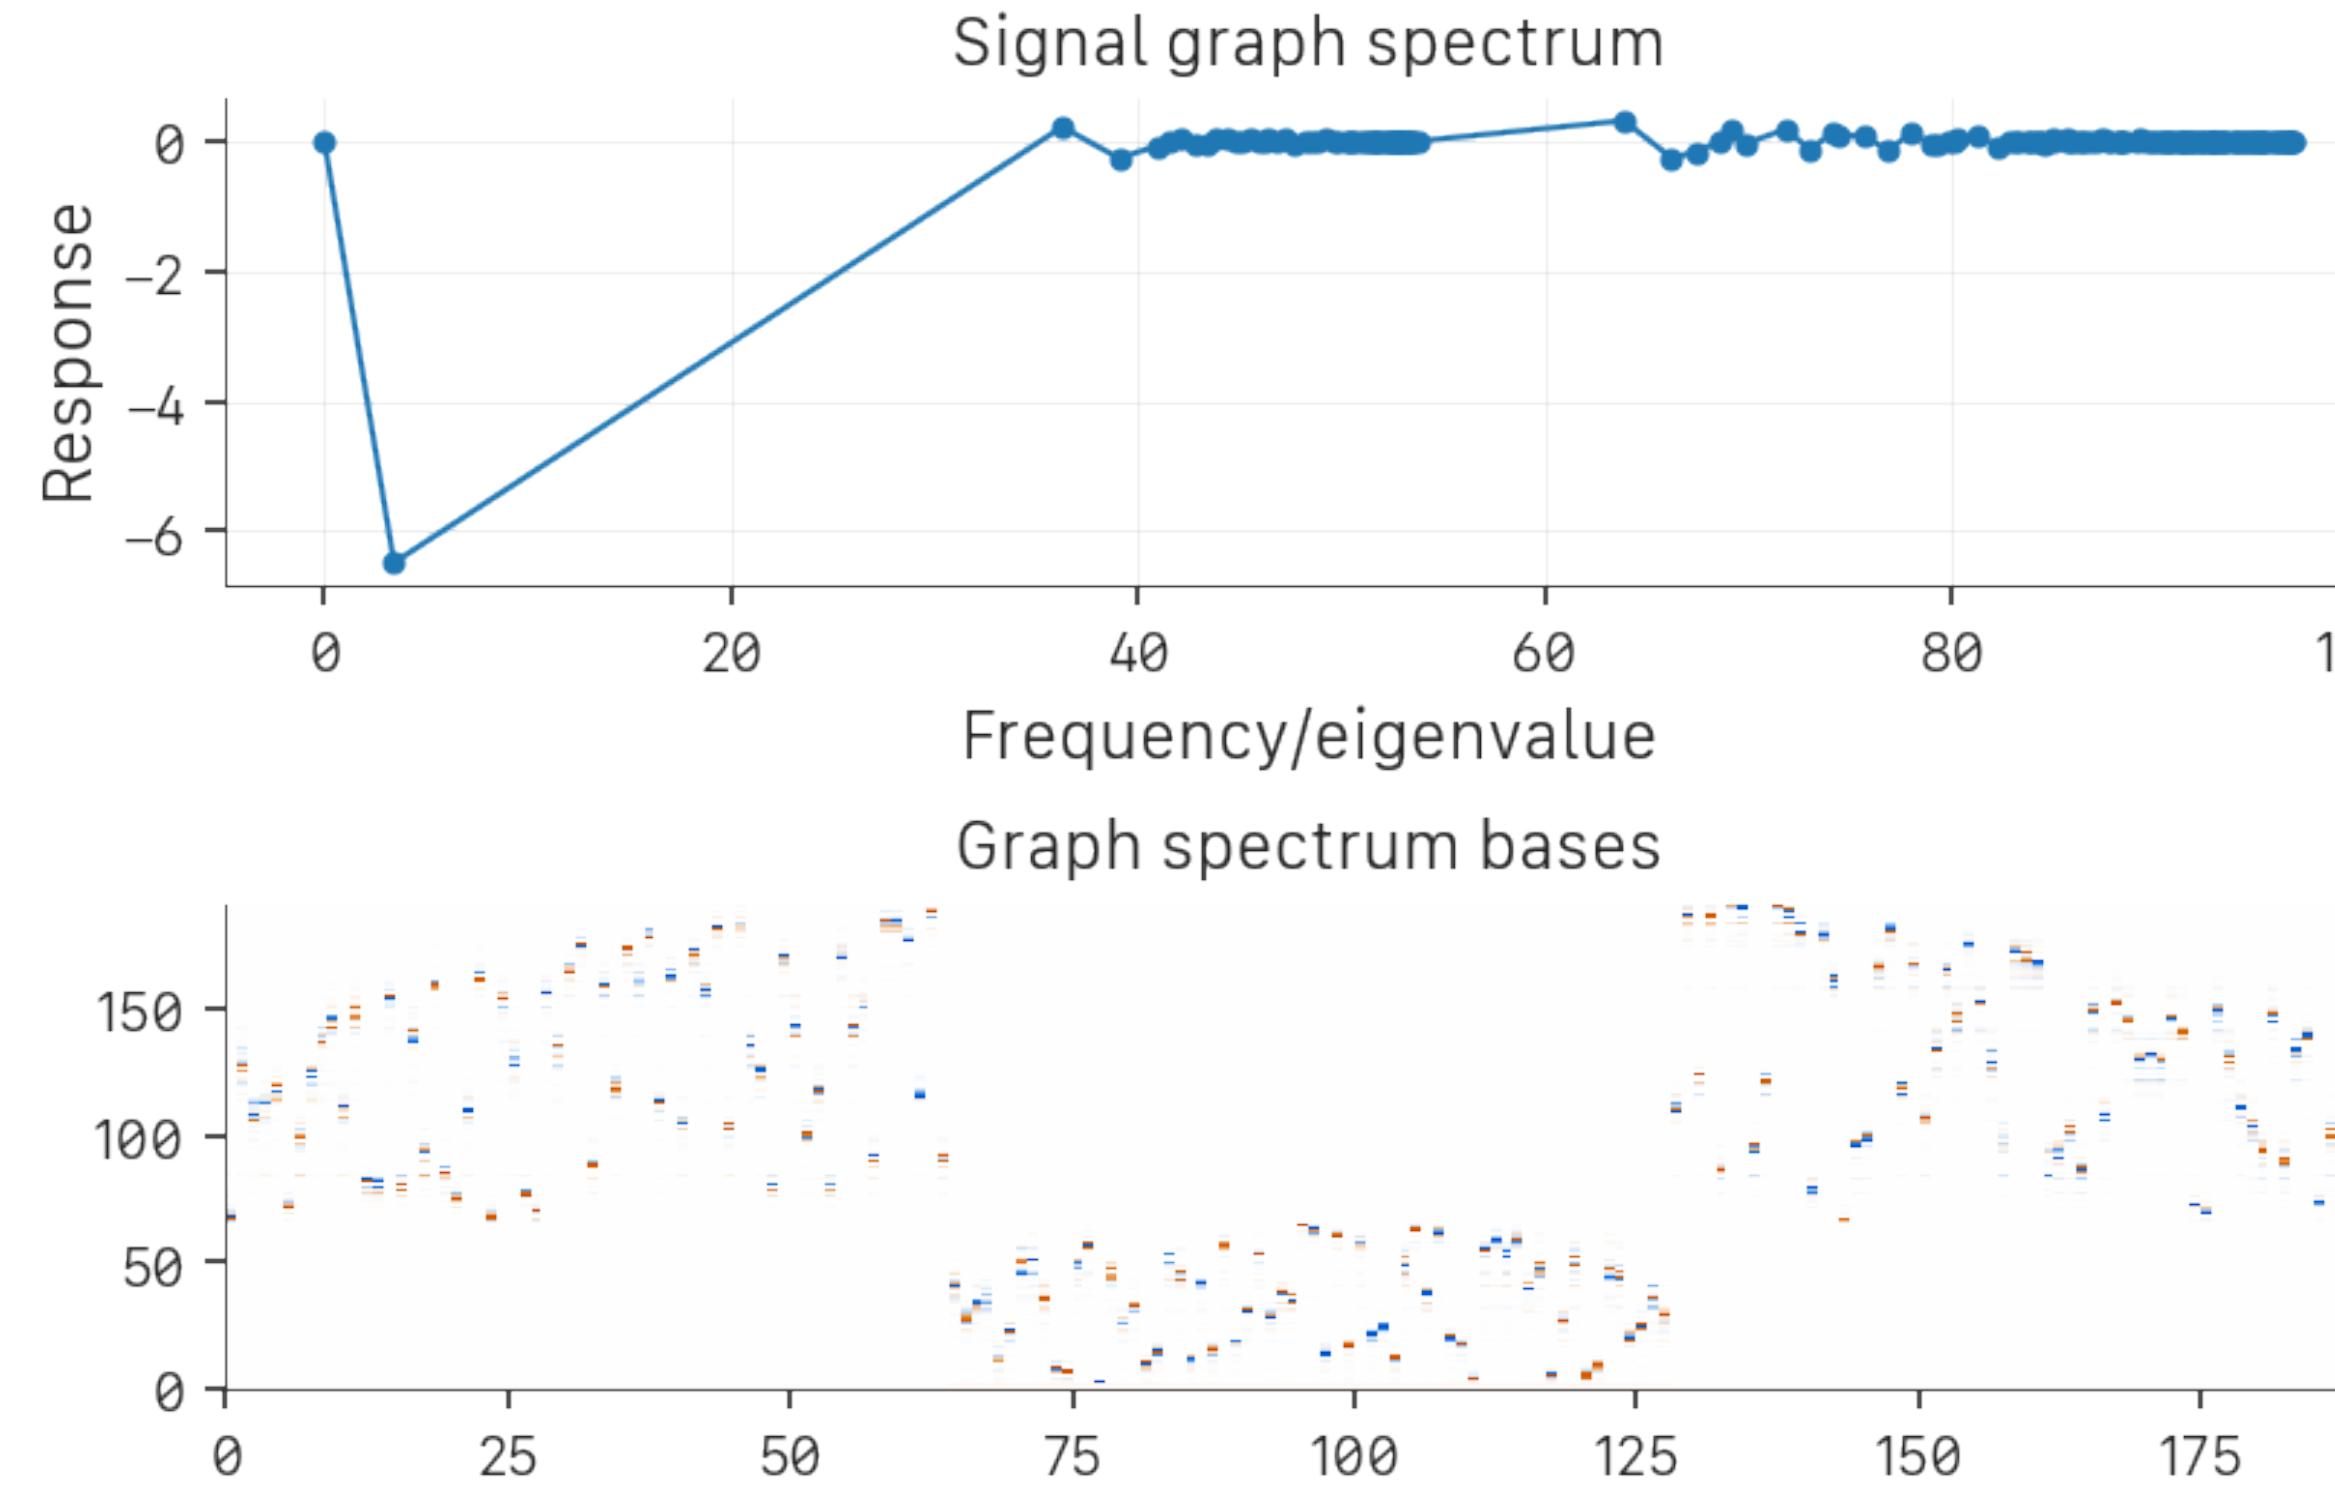
- We wish to convey that similar values are “closer”
 - Make an adjacency matrix that also uses the signal values

$$w_{i,j} = e^{-a|i-j| - \beta|x_i - x_j|}$$



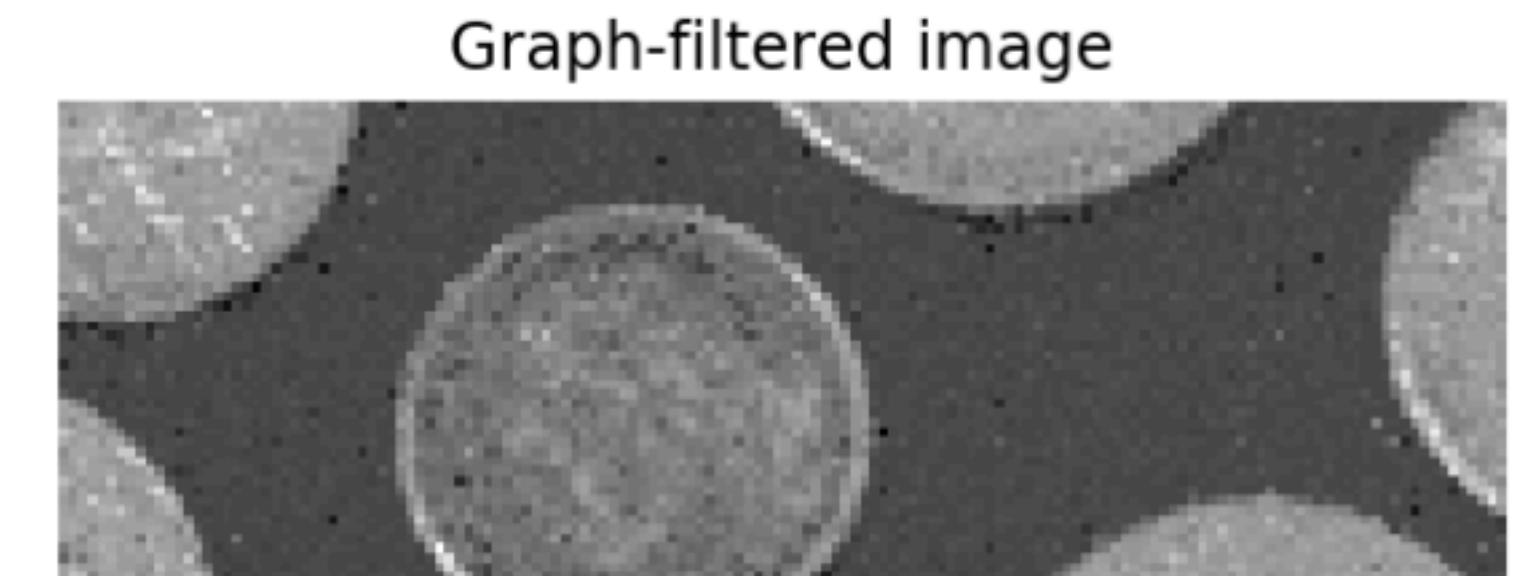
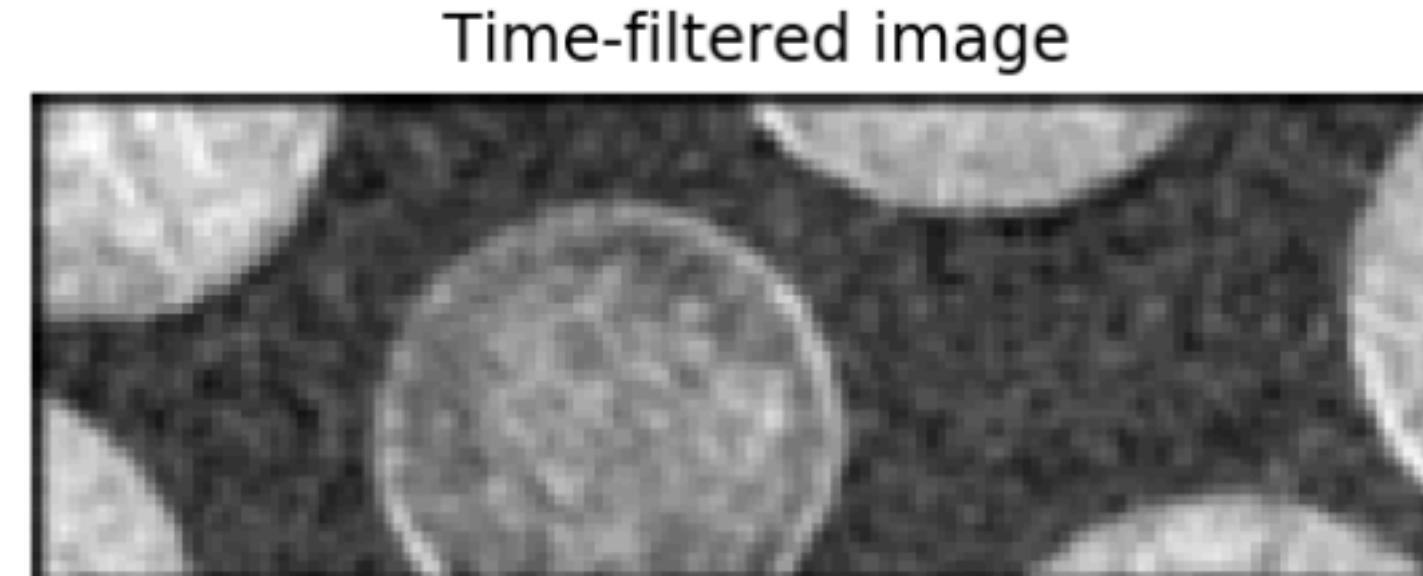
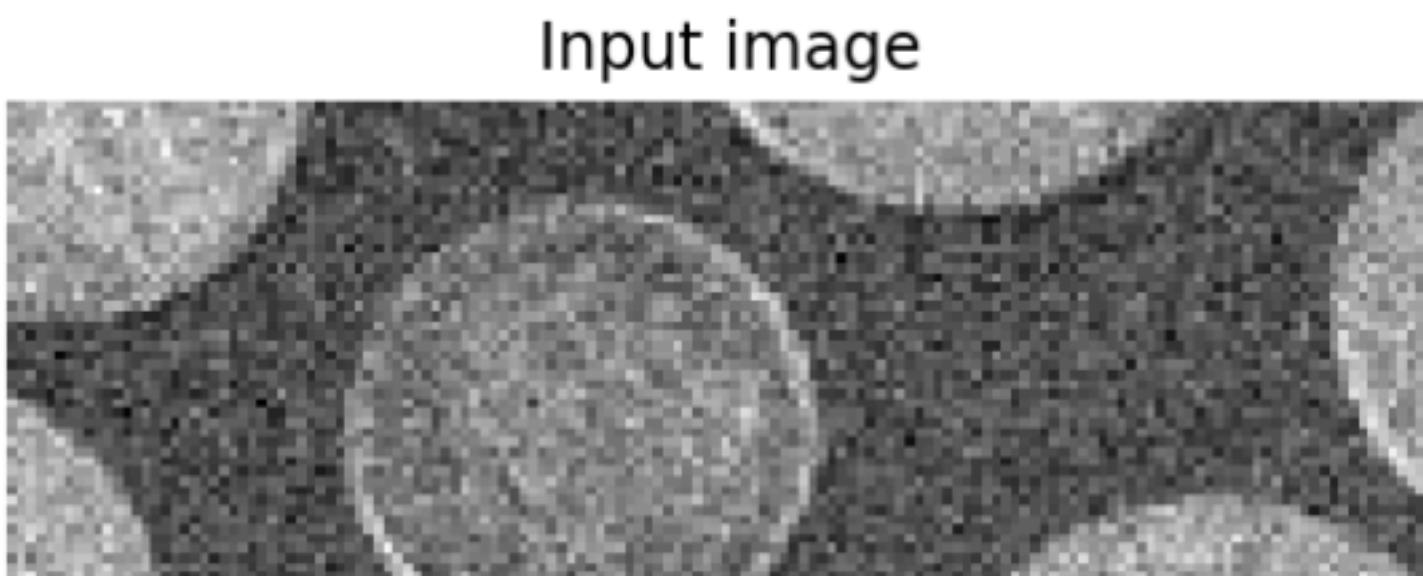
Edge-preserving filter

- Resulting filter doesn't smooth the large step
 - Effectively preserves edges in the signal



Edge preserving filters: 2D Example

- Consider a noisy image
 - Sample value is given by pixel intensity
 - Spatial location is given by pixel co-ordinates
- Vectorize the samples and spatial locations
 - Everything else stays the same as before

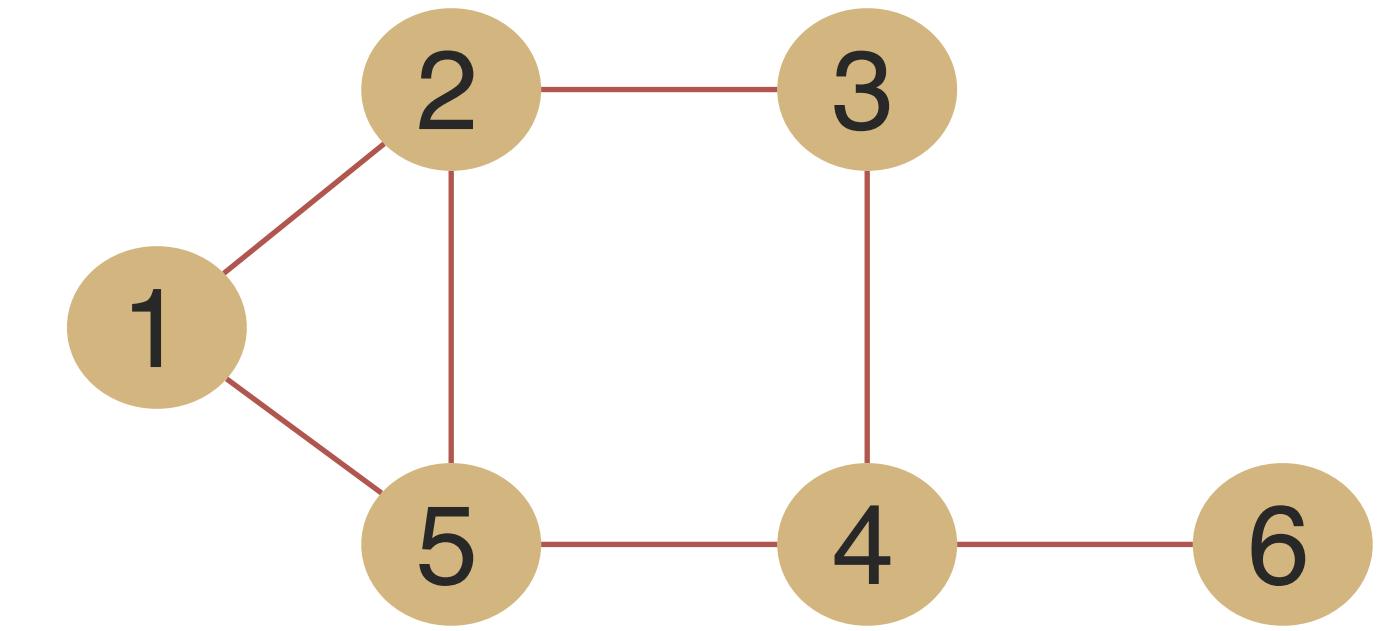


Defining graphs from data

- K-Nearest Neighbor graphs
 - Alternative way to impose a graph on data
 - Every vertex is connected to only K of its closest points
- Fourier Transforms on KNN graphs
 - Eigendecomposition on KNN adjacency matrix
 - We did that for manifold learning!

Filtering in the spatial domain

- Classical DSP: filtering in time vs frequency
 - Frequency: Multiplication \leftrightarrow Time: Convolution
 - Can we have such interpretations for graphs?
- Filter outputs are simply weighted sums of time-shifted inputs
 - We use a time shift operator to define filtering
 - Shifts a sequence by one time step
- The *graph-shift matrix*
 - Non-zero for connected edges, e.g. the Laplacian
 - Can (confusingly) otherwise be arbitrary!



$$S = \begin{bmatrix} S_{11} & S_{12} & 0 & 0 & S_{15} & 0 \\ S_{21} & S_{22} & S_{23} & 0 & S_{25} & 0 \\ 0 & S_{32} & S_{33} & S_{34} & 0 & 0 \\ 0 & 0 & S_{43} & S_{44} & S_{45} & S_{46} \\ S_{51} & S_{52} & 0 & S_{54} & S_{55} & 0 \\ 0 & 0 & 0 & S_{64} & 0 & S_{66} \end{bmatrix}$$

Filtering in the spatial domain

- The graph shift can be computed locally
 - $\mathbf{S} \cdot \mathbf{x}$ can be computed from current and 1-hop nodes
 - $\mathbf{S}^2 \cdot \mathbf{x}$ uses up to 2-hop nodes
 - Used to define filters
- Example:

$$\mathbf{x} = [-1, 2, 0, 0, 0, 0]^\top$$

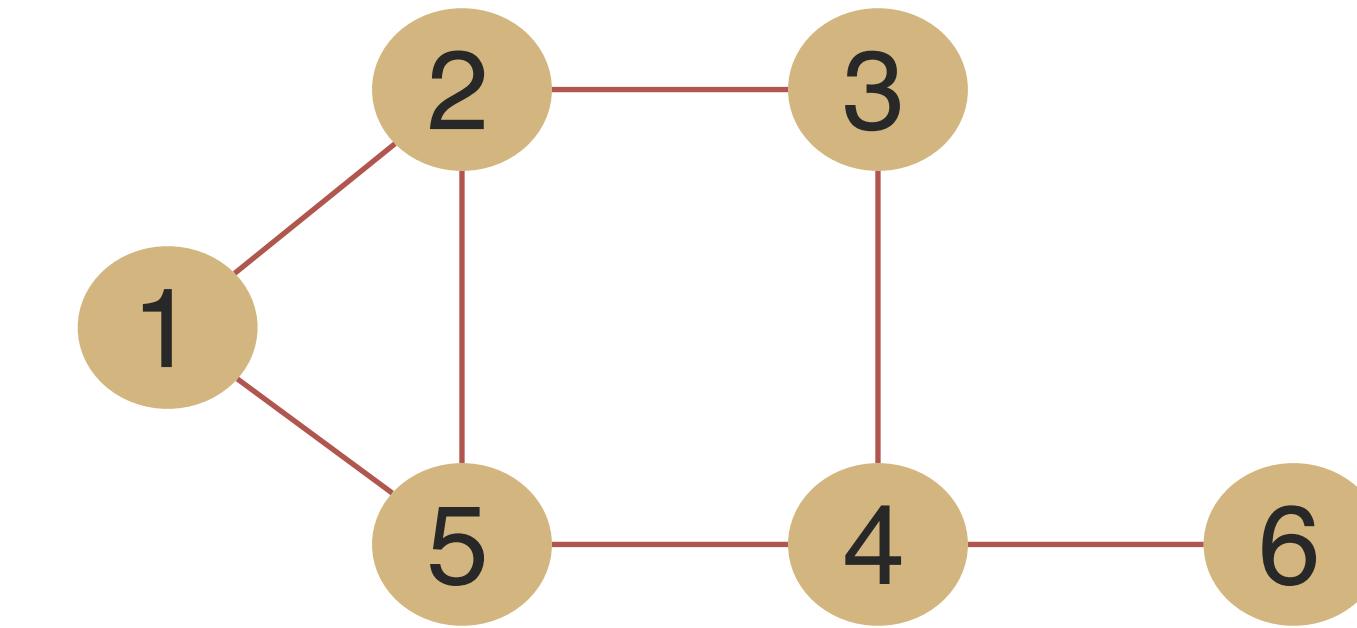
$$\mathbf{h} = [1, 1, 0.5]^\top$$

$\mathbf{S} = \mathbf{I} + \mathbf{W}$ \leftarrow Adjacency matrix

$$\mathbf{y} = h_0 \mathbf{x} + h_1 \mathbf{S} \cdot \mathbf{x} + h_2 \mathbf{S}^2 \cdot \mathbf{x}$$

$$\mathbf{S} \cdot \mathbf{x} = \begin{bmatrix} 1 \\ 1 \\ 2 \\ 0 \\ 1 \\ 0 \end{bmatrix}, \quad \mathbf{S}^2 \cdot \mathbf{x} = \begin{bmatrix} 3 \\ 5 \\ 3 \\ 3 \\ 3 \\ 0 \end{bmatrix}, \quad \mathbf{y} = h_1 \mathbf{x} + h_2 \mathbf{S} \cdot \mathbf{x} + h_3 \mathbf{S}^2 \cdot \mathbf{x} = \begin{bmatrix} 1 \\ 5 \\ 2.5 \\ 1.5 \\ 2 \\ 0 \end{bmatrix}$$

Filtering



To summarize

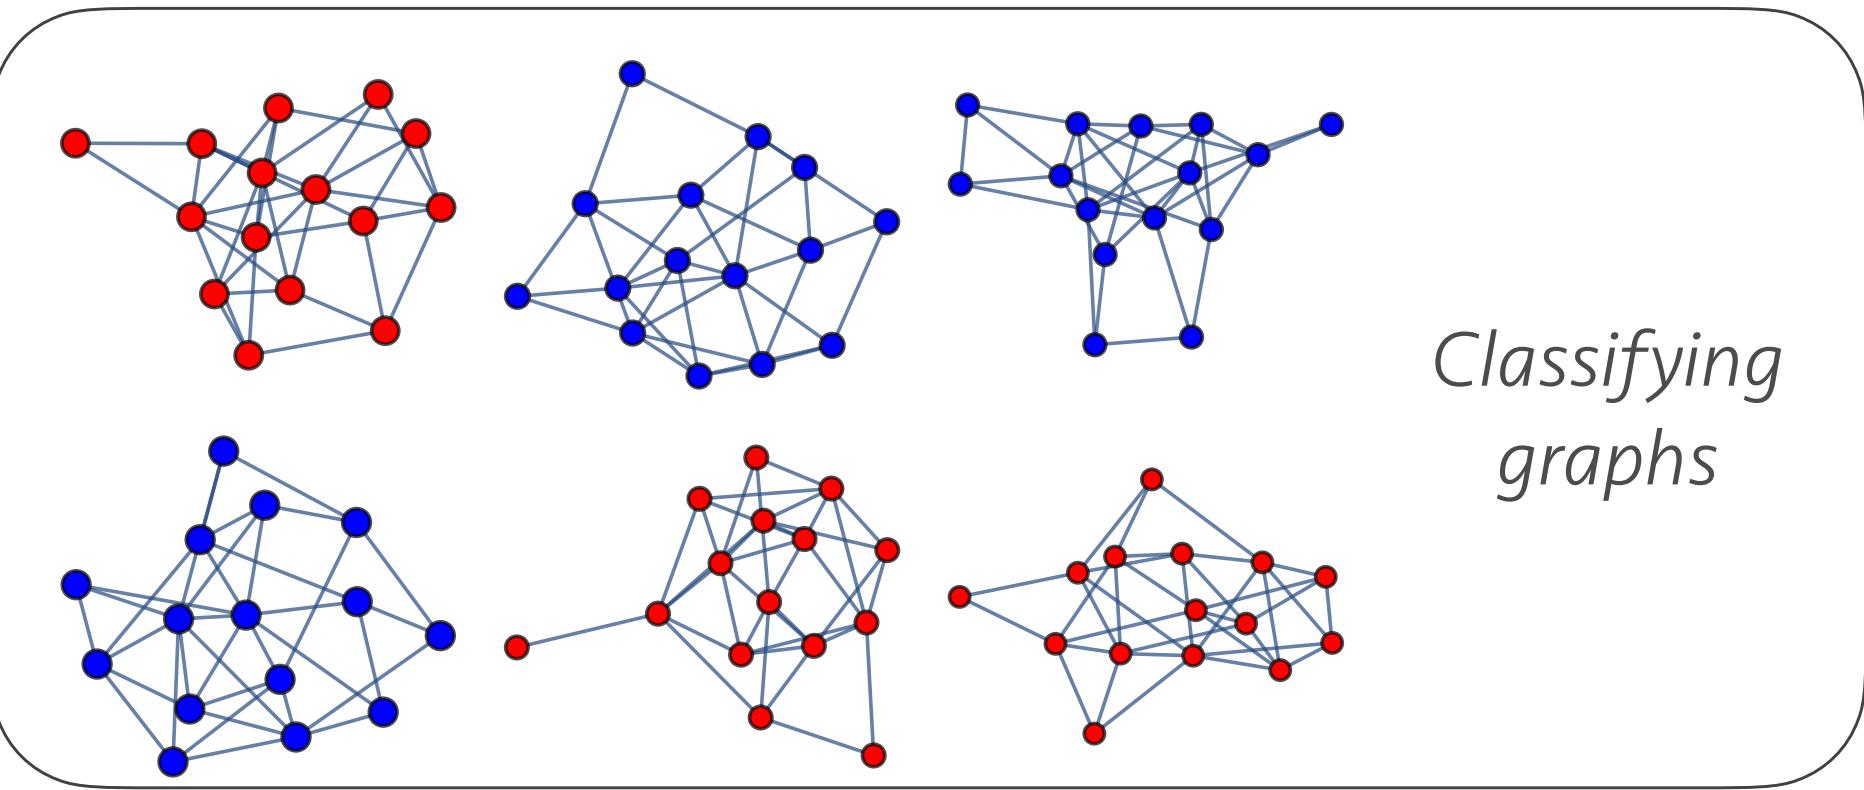
- Signals can be defined as graphs
 - Allows us to represent richer sample relationships
- We can define direct analogs to classical DSP
 - Time/frequency domain
 - Filtering operations
 - Not mentioned today: sampling, wavelets, etc ...
- So now we can use these to do machine learning!

Graph Neural Networks

- Neural nets that operate on graphs
 - Most often rely on graph convolutions (CNN equivalents)
 - Using the graph-shift matrix we can also define graph RNNs
- Useful for non-conventionally sampled data
 - Molecules, social graphs, Mesh representations, LIDAR data, ...
- One of the hot topics recently

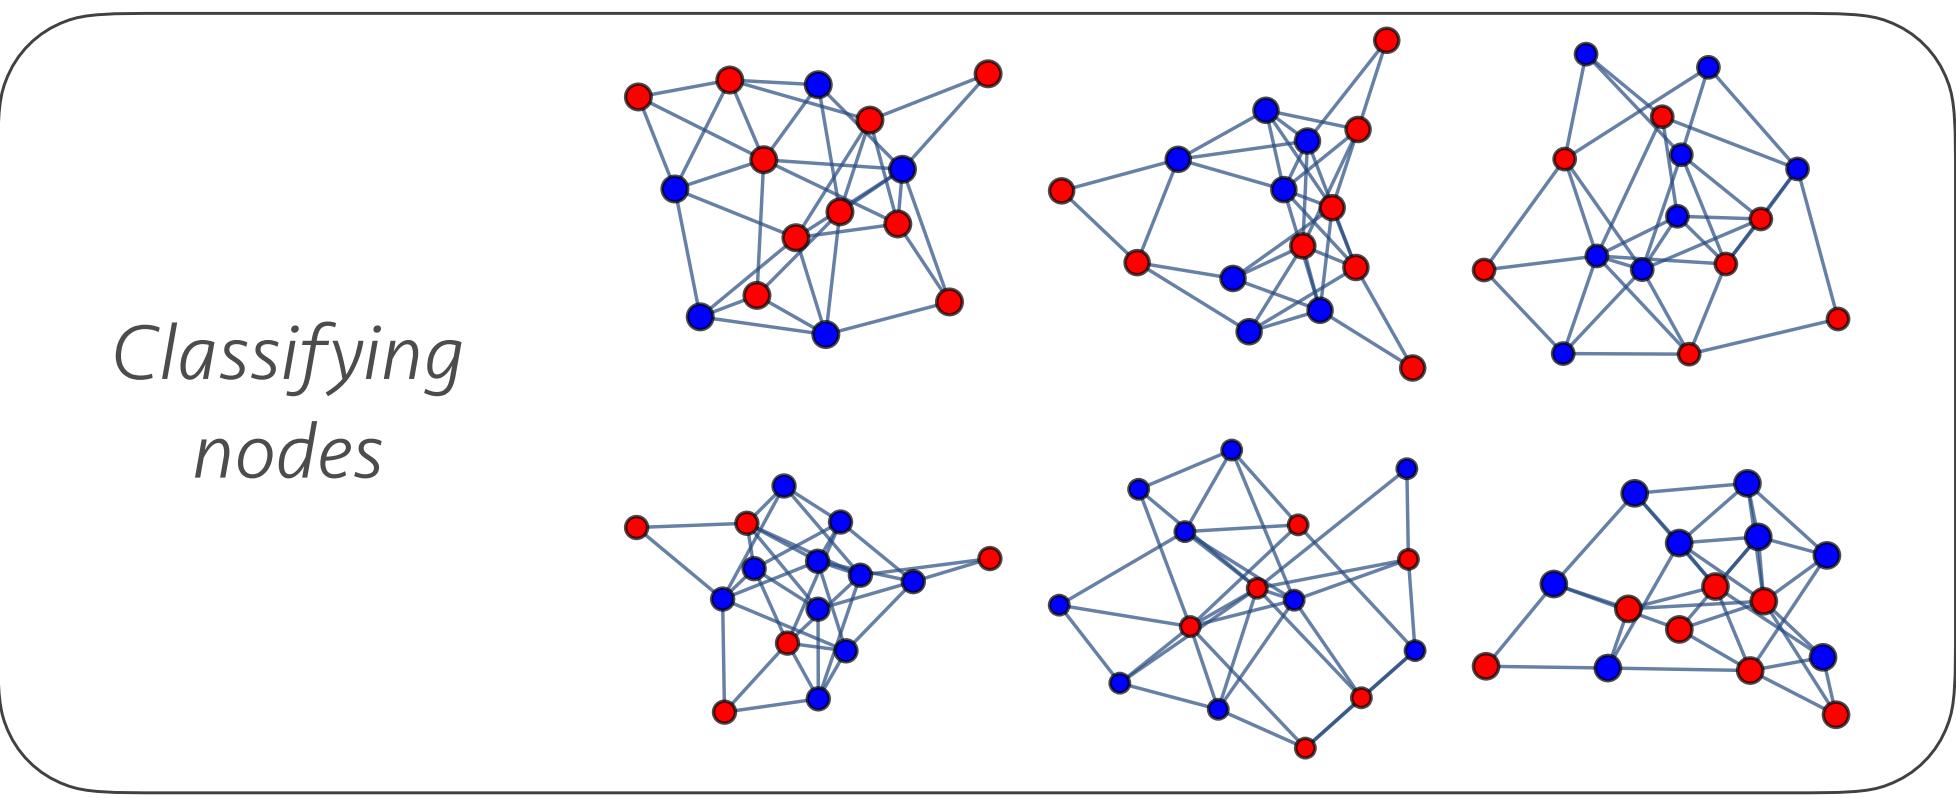
Types of graph processing

- Graph-level processing
 - Labeling an entire graph
 - E.g. classify a molecule



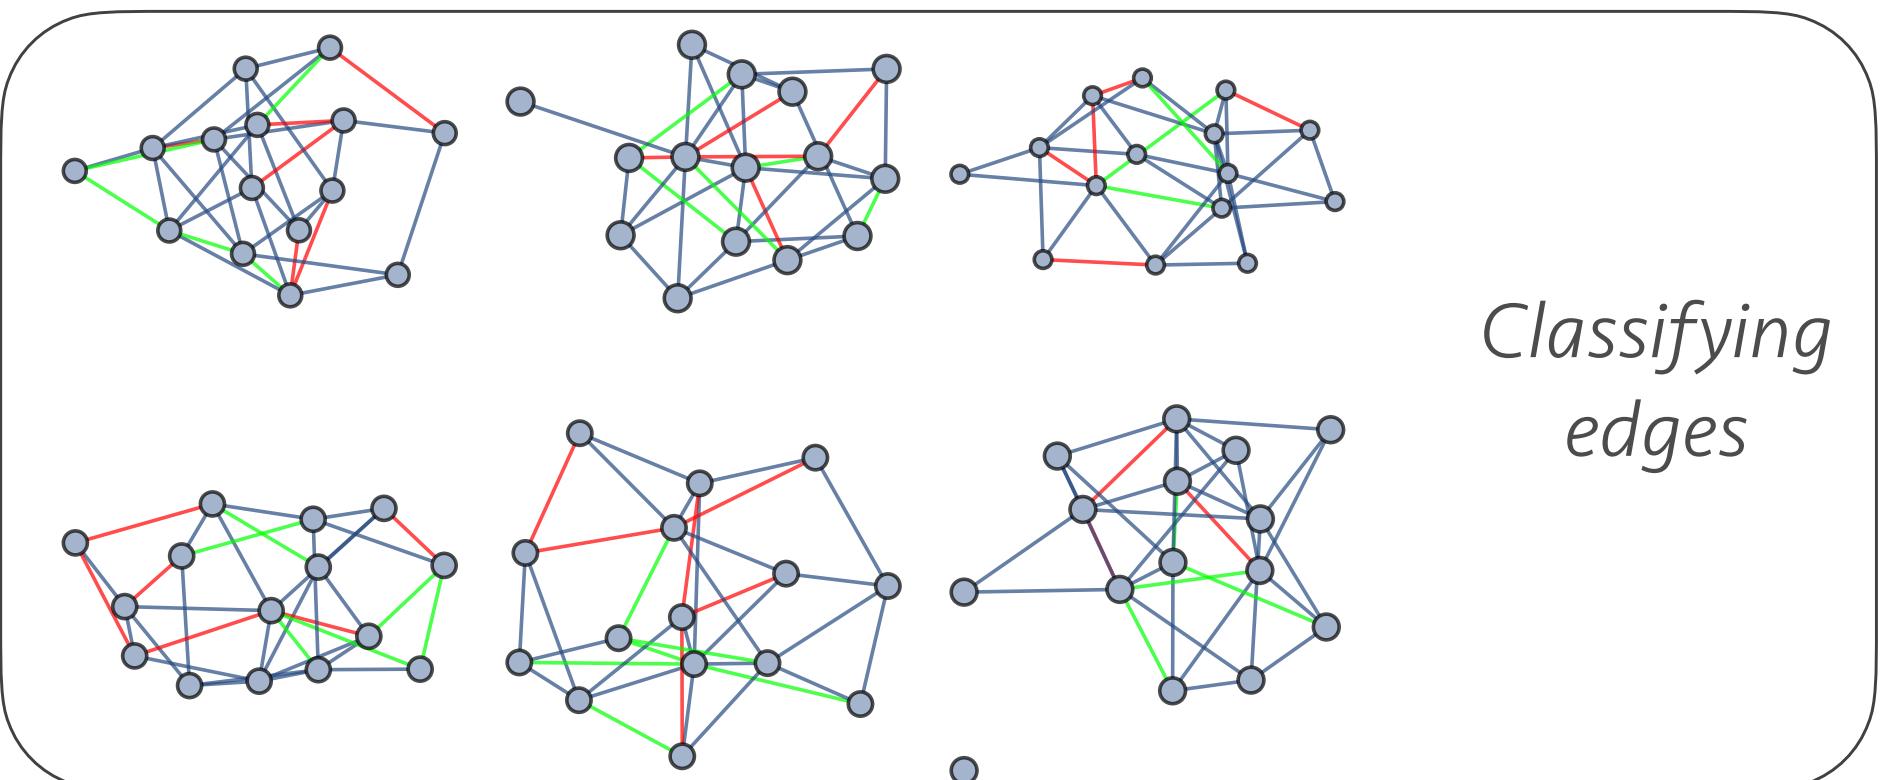
Classifying graphs

- Node-level processing
 - Labelling each node
 - E.g. LIDAR pixel classification



Classifying nodes

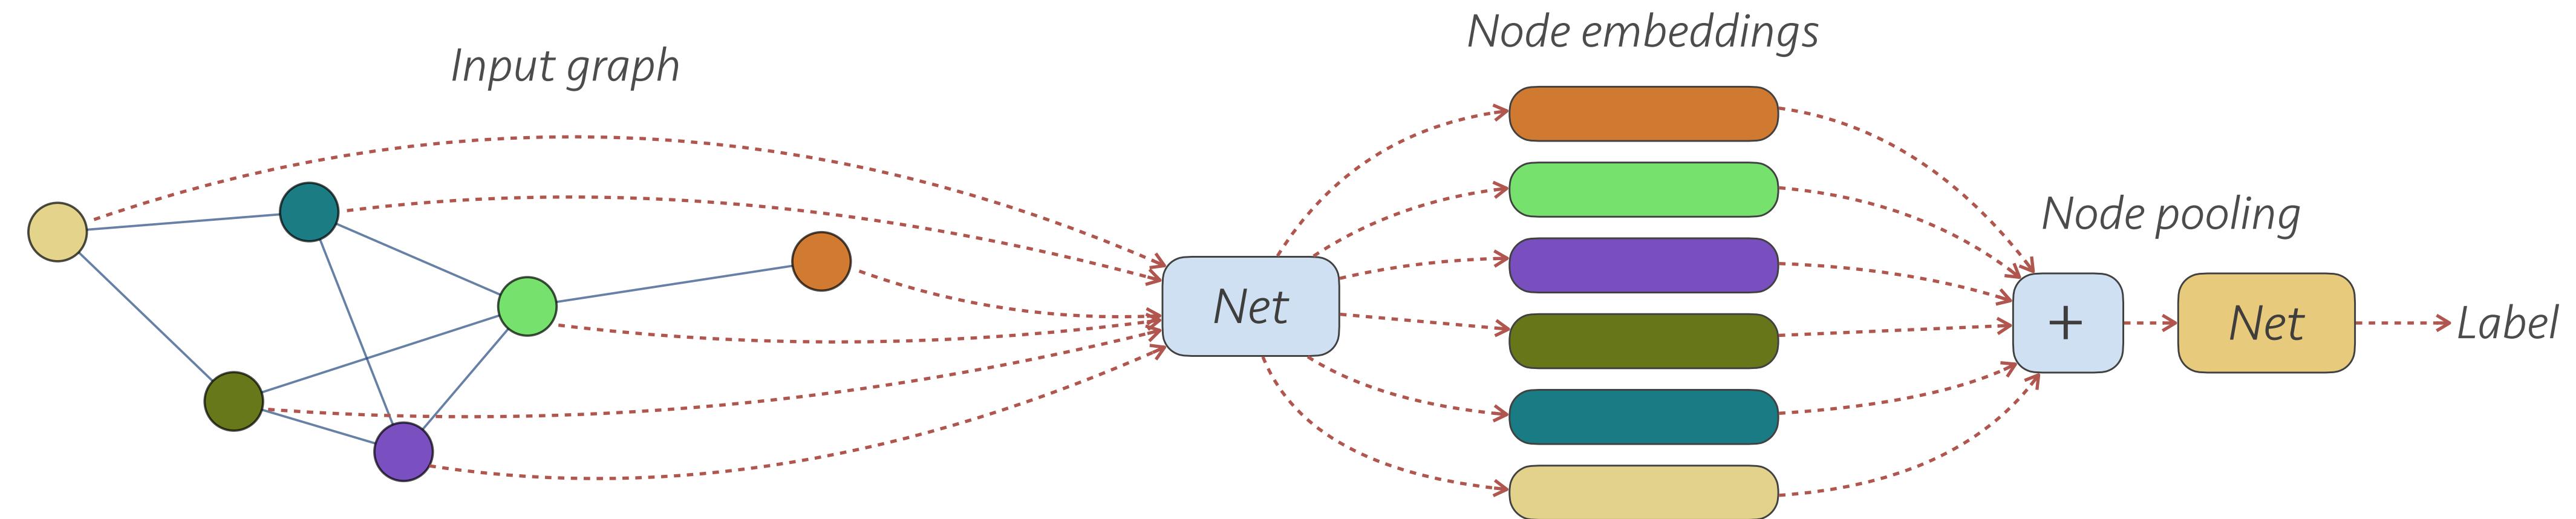
- Edge-level processing
 - Labelling/predicting each edge
 - E.g. action prediction



Classifying edges

Towards a graph network

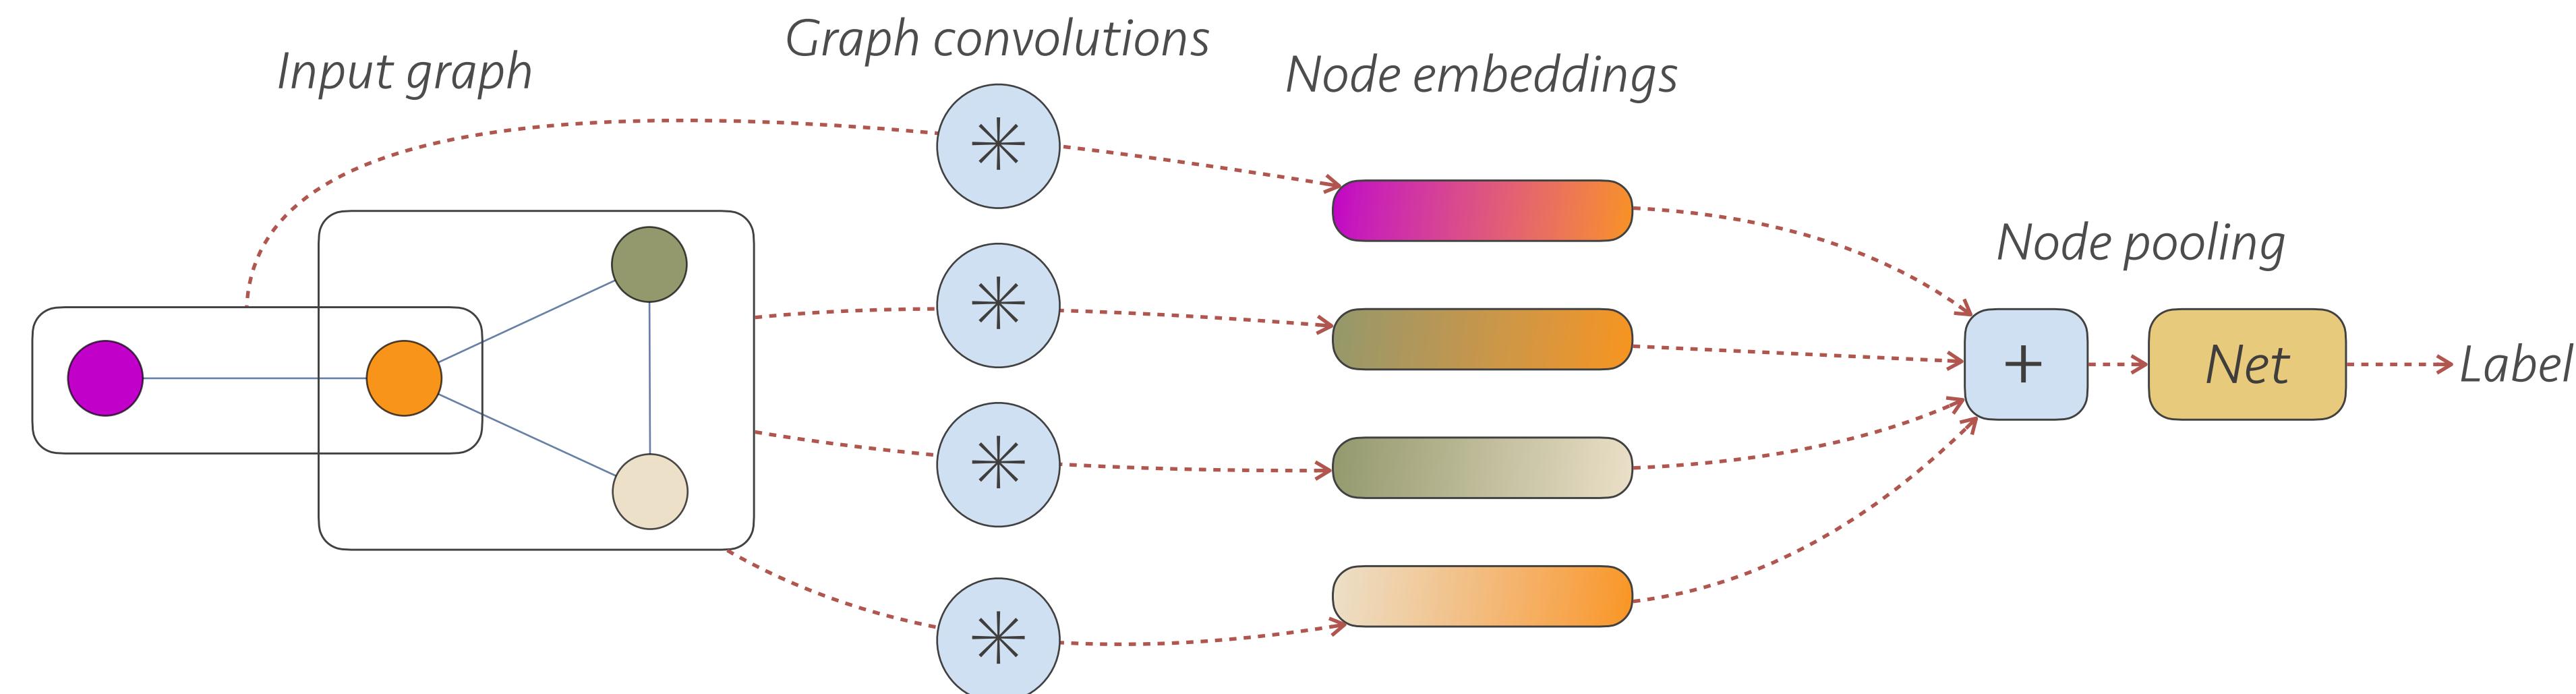
- Embedding the nodes only
 - Apply the same network on each node to obtain an embedding



- Depending on task
 - Label nodes → embedding is the label
 - Label graph → pool information by combining node embeddings
 - Usually take their sum and feed into a regular network

Adding the graph structure

- Use graph convolutions to combine neighboring nodes when computing node embeddings
 - Explicitly includes local structure of input graph
 - We can also add a “master node” to get global structure

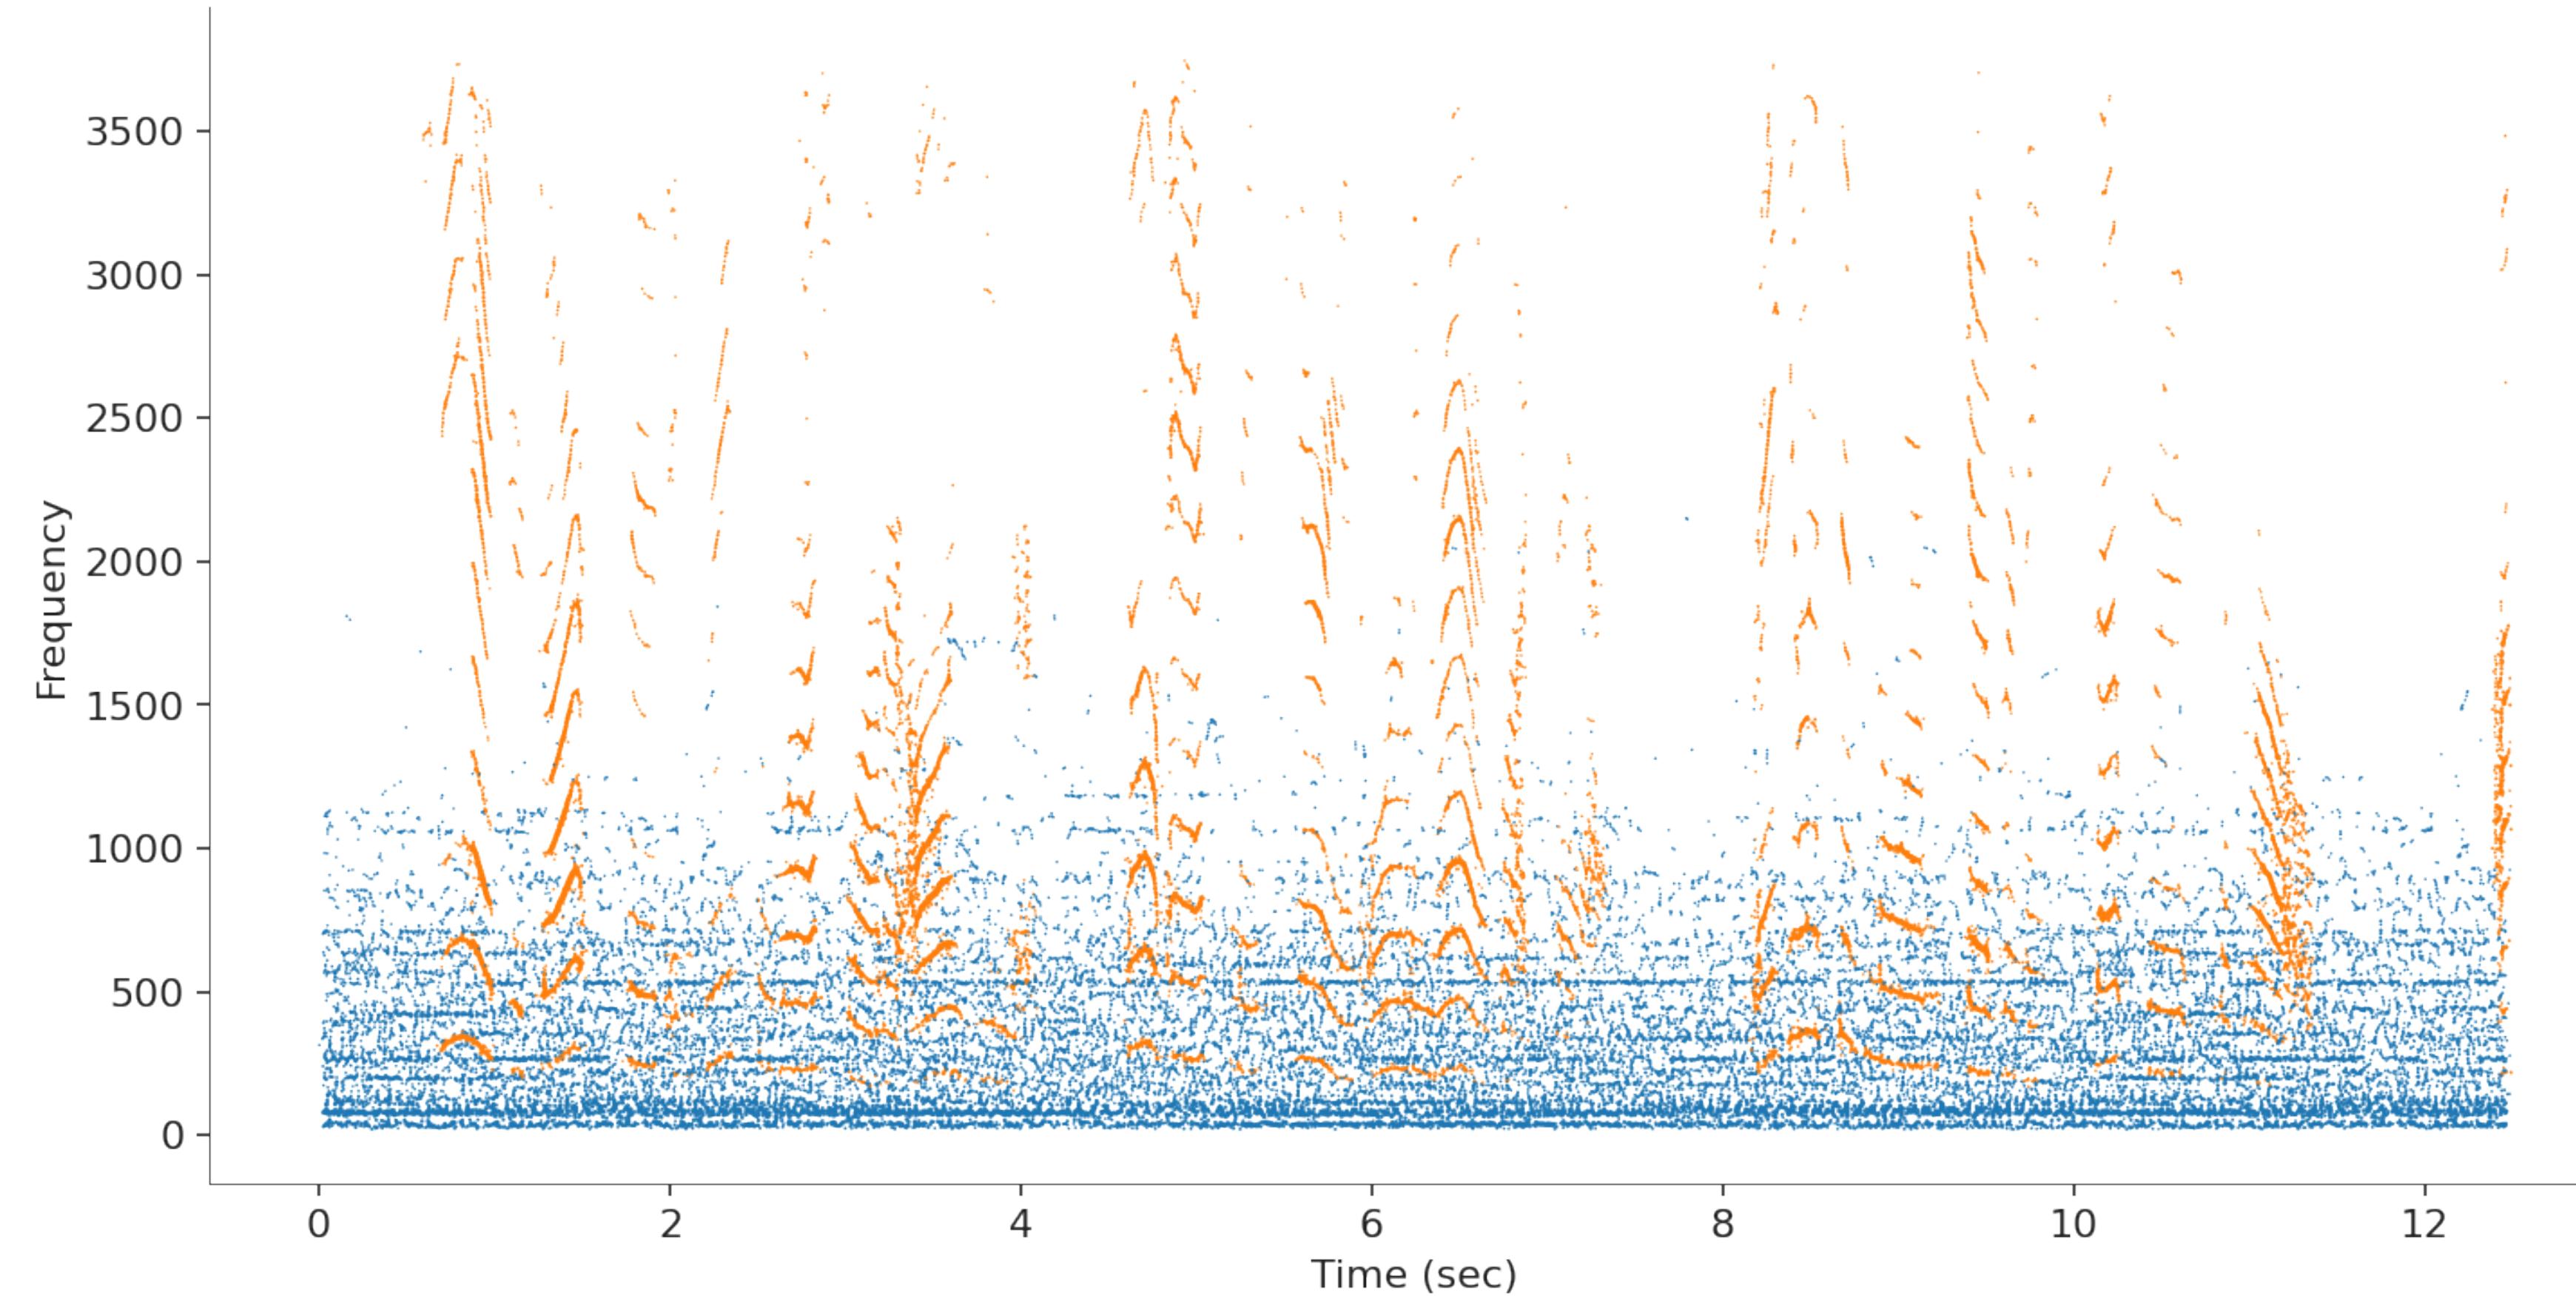


Lots of variations here

- Different sorts of Graph CNN types
 - Mostly least-publishable unit variations ☹
- Attention/transformer models
 - Instead of convolving, attend to relevant nodes
- Autoregressive models
 - GraphRNN, recursively passes graph messages

Example: Point cloud audio processing

- Using point values to describe spectral content
 - New representation that requires novel processing

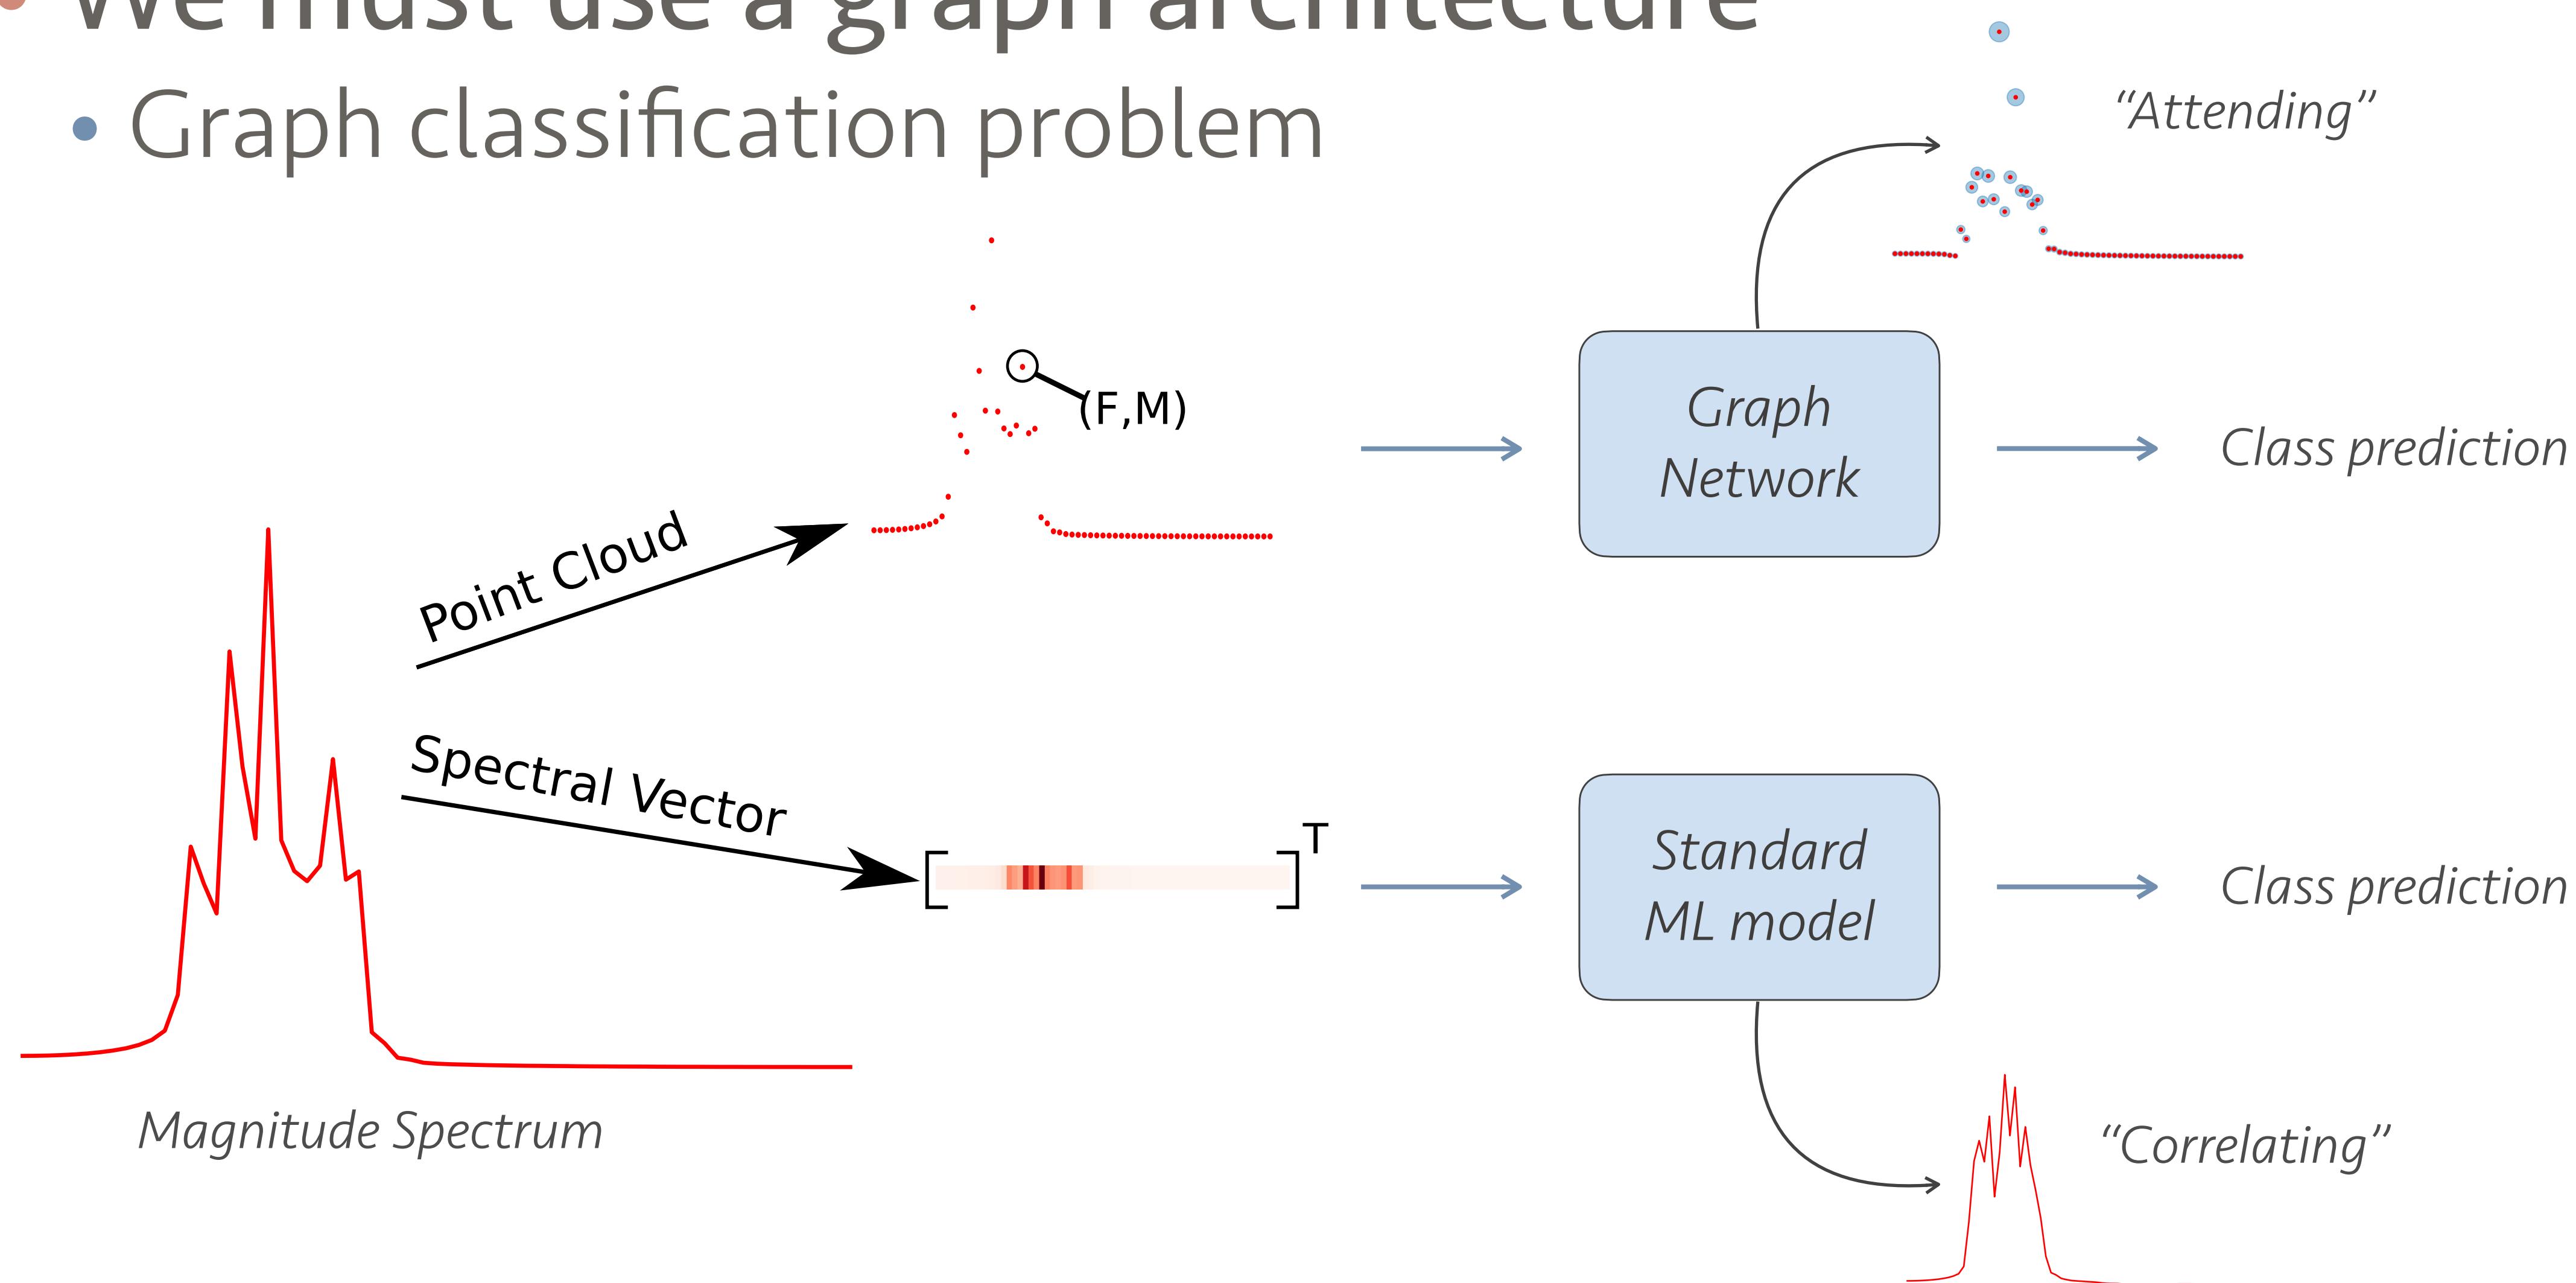


The advantage of point cloud data

- With point clouds we only use as much data as we can
 - We can pick and choose depending on the hardware
 - Or we can sample as needed
- A different representation:
 - Instead of e.g. a spectral vector: $\mathbf{f} = [f_{\omega_1}, f_{\omega_2}, f_{\omega_3}, \dots, f_{\omega_N}]$
 - Use nodes: $\mathbf{t} = [(\omega_1, v), (\omega_4, v), (\omega_{10.2}, v), \dots]$
 - Tuples \mathbf{t} can be a subset of \mathbf{f} , or even resampled values
- Frees us from fixed sizes and sample rates (ω is in Hz)

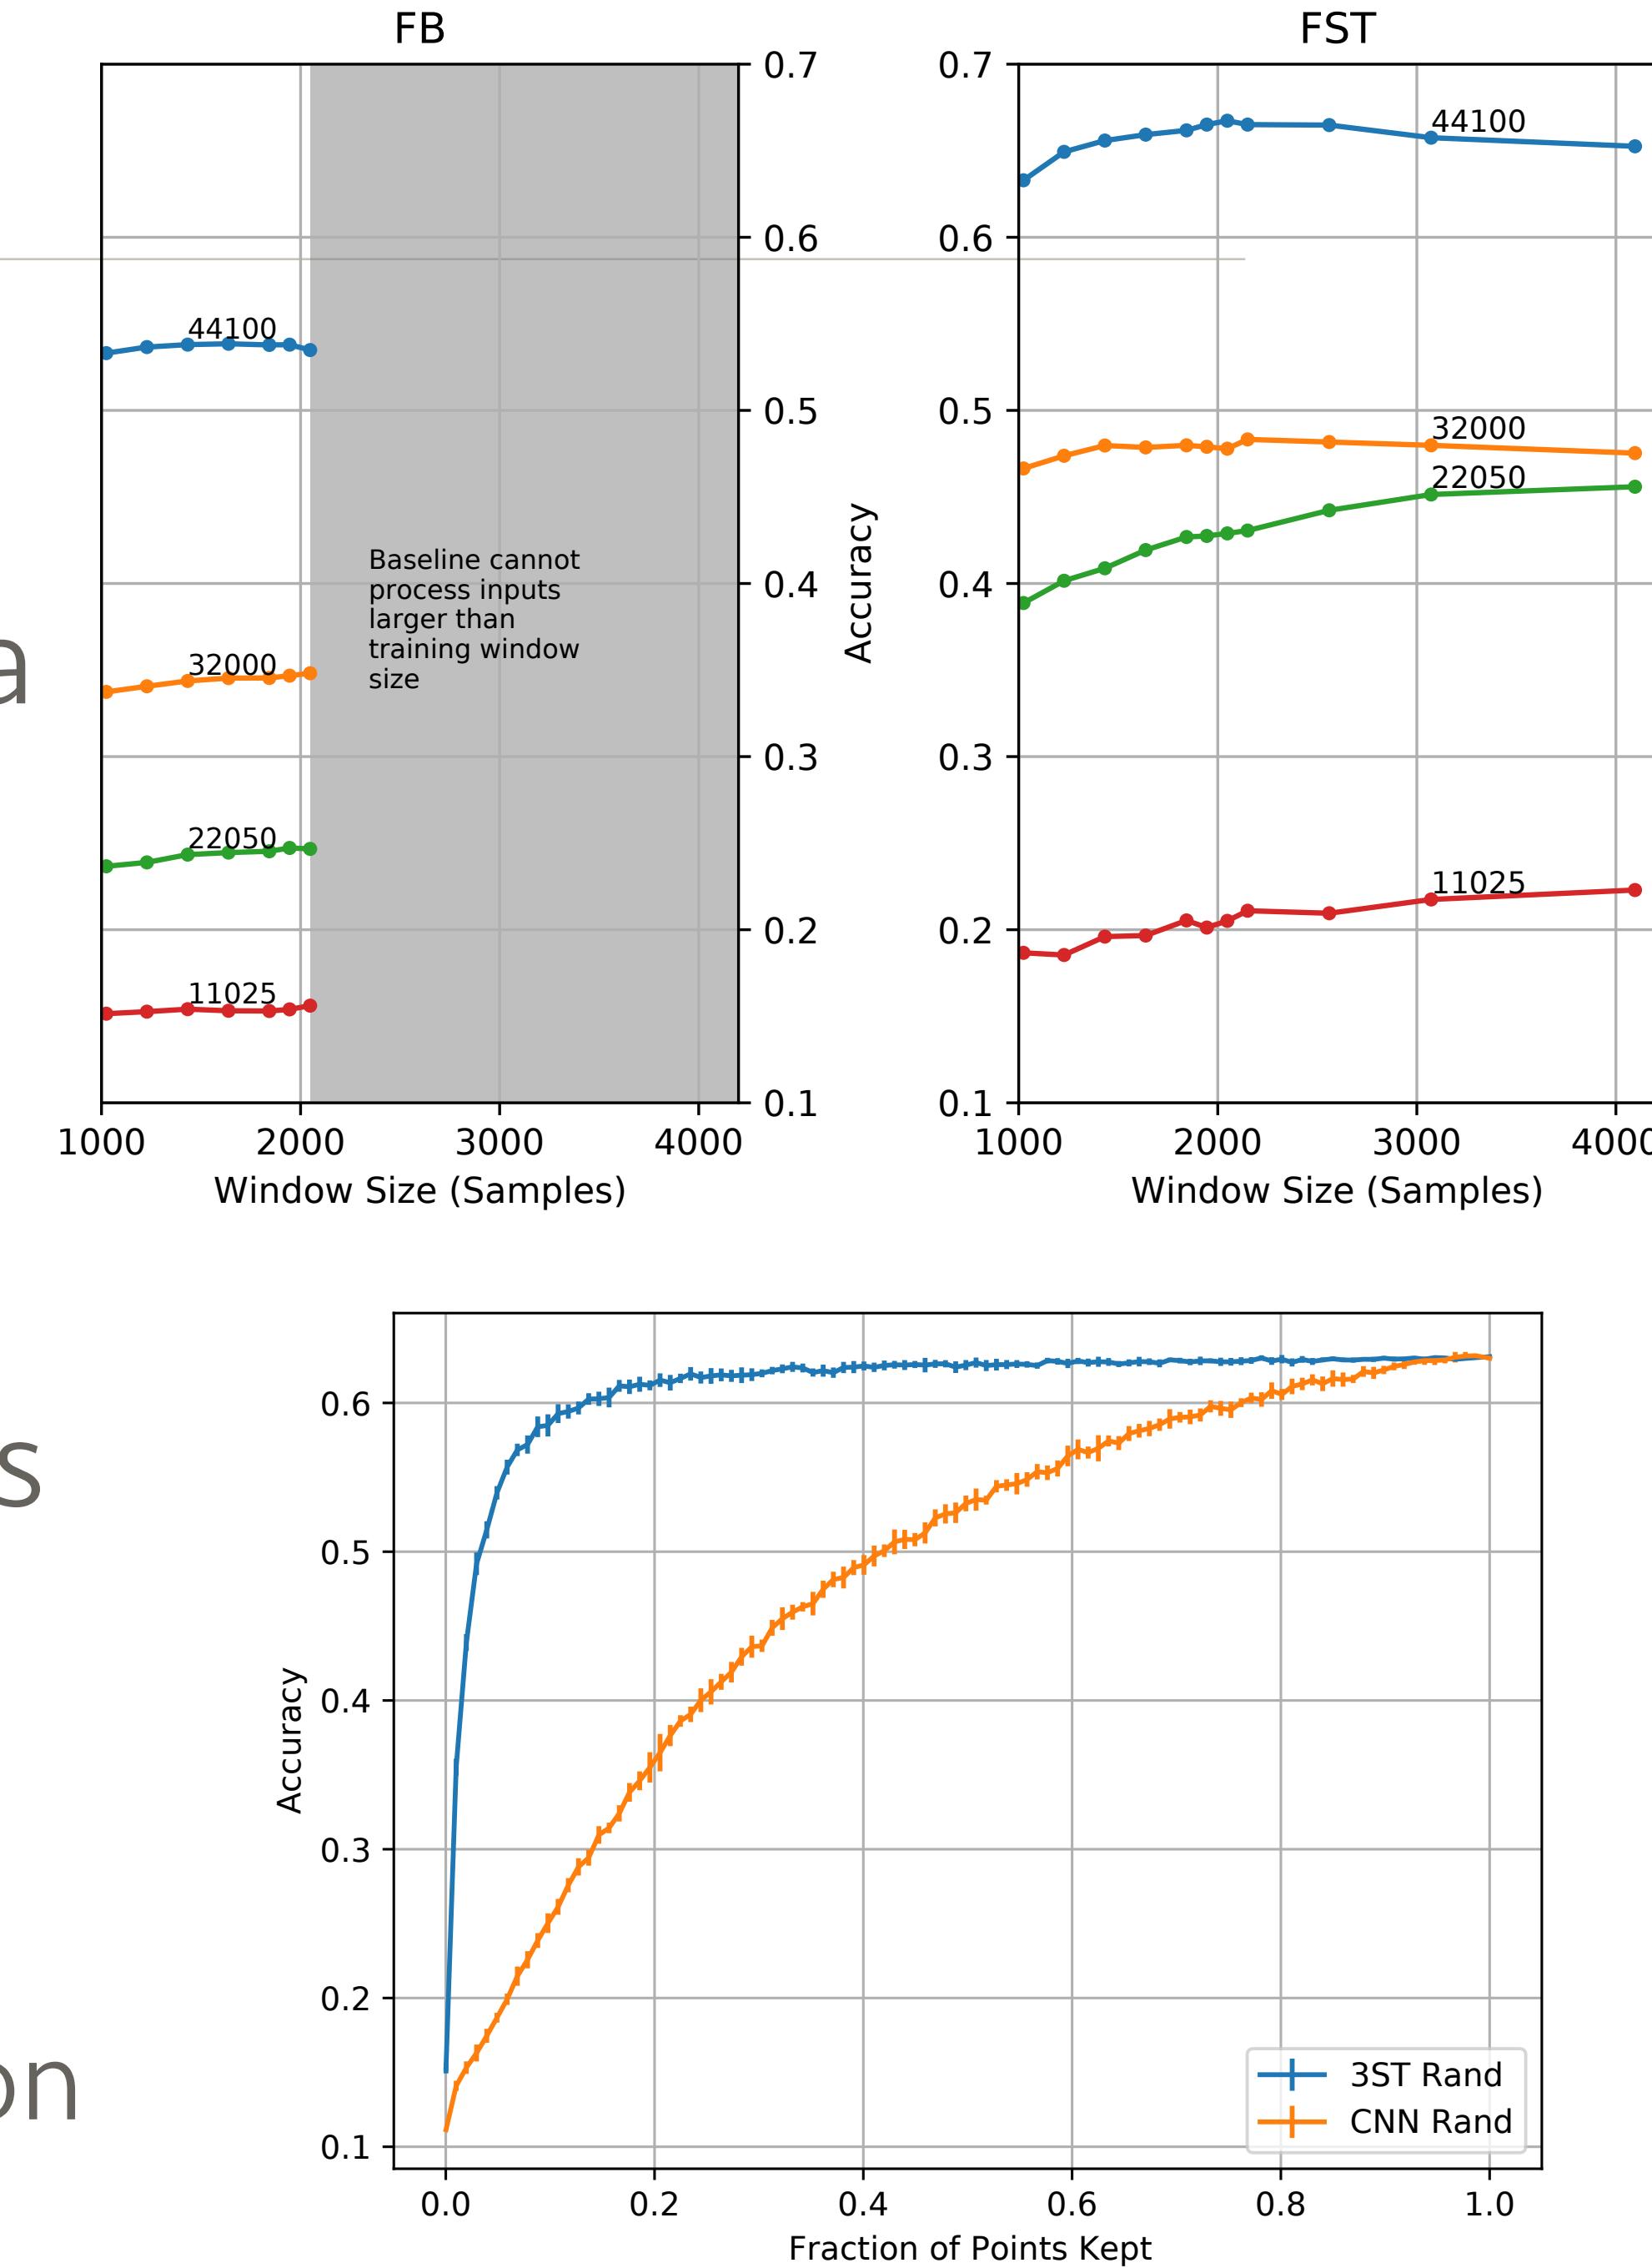
Processing point data

- We must use a graph architecture
 - Graph classification problem



Example application

- Sound recognition task (ESC-10)
 - Classifying from spectro-temporal data
- Win-win-win situation
 - Better performance overall
 - Largely invariant to sampling rate
 - Smaller model: 80K vs 160K parameters
 - Works well with only 20% of input!
 - Faster training
 - Using only 1/16th of training data when using subsampling of input representation



Recap

- Interpreting signals as a graph has benefits
 - Graphs can capture underlying structure
- Signal Processing on Graph signals
 - Generalizing classical DSP approaches to graph signals
 - The Graph Fourier Transform
 - Applications: Filtering and Denoising
 - E.g., Edge-preserving filters
- Graph networks
 - Generalize passage of information through graphs
 - Several novel applications now possible

Reading material

- Graph Signal Processing
 - <https://arxiv.org/pdf/1712.00468.pdf>
- Graph convolutions and graph nets:
 - <https://distill.pub/2021/understanding-gnns/>
 - <https://distill.pub/2021/gnn-intro/>
- Point cloud audio processing
 - <https://arxiv.org/pdf/2105.02469.pdf>

Upcoming schedule

- On travel Nov 27-30, next lecture will be our last
 - Secure Multiparty processing
 - Rest of generative models
 - We might have some extra time left
 - Send me ideas on subjects to discuss briefly
- Projects
 - We will hold a poster session on Dec 5th noon-2pm
 - This won't be your final report, fine if you have no results yet
 - Tell us what you are doing and how you plan to do it
 - We have 29 groups → ~4min/group!