

Privacy-Preserving Machine Learning

14 November 2023

Today's lecture


- Secure Multiparty Computations
 - Some cryptography basics
 - Applying SMC to MLSP
- Other pending tidbits

A problem

- Machine learning can be a great service
 - Give me your data, I'll give you some insight
- But, you can't ask everyone for their data
 - Sending your voice mail for transcription?
 - Camera feed from your house for security?
 - Sharing medical data with your phone?

What would be nice


- In a perfect world there would be absolute trust
 - But do you really trust anyone with your data?



- In the real world
 - We want to ensure the privacy of our data
 - Others want to ensure privacy of their algorithms & data
 - Can these constraints co-exist?


Starting simple

- Doing some statistics with private data
- E.g. finding the average person's salary in the room?
 - The unacceptable way: Ask everyone for their income



A more socially acceptable approach


- Obfuscate dollar amounts with noise
 - Give a random number to next person, ask them to add their salary, pass the sum to the next person, and repeat
 - I get back sum of salaries plus my known random number
 - Remove random number and divide by number of people!
 - No private data has been shared!




A signal-friendly example

- Your government wants you to examine the differences between satellite images over time
 - But the images are classified

Before



After



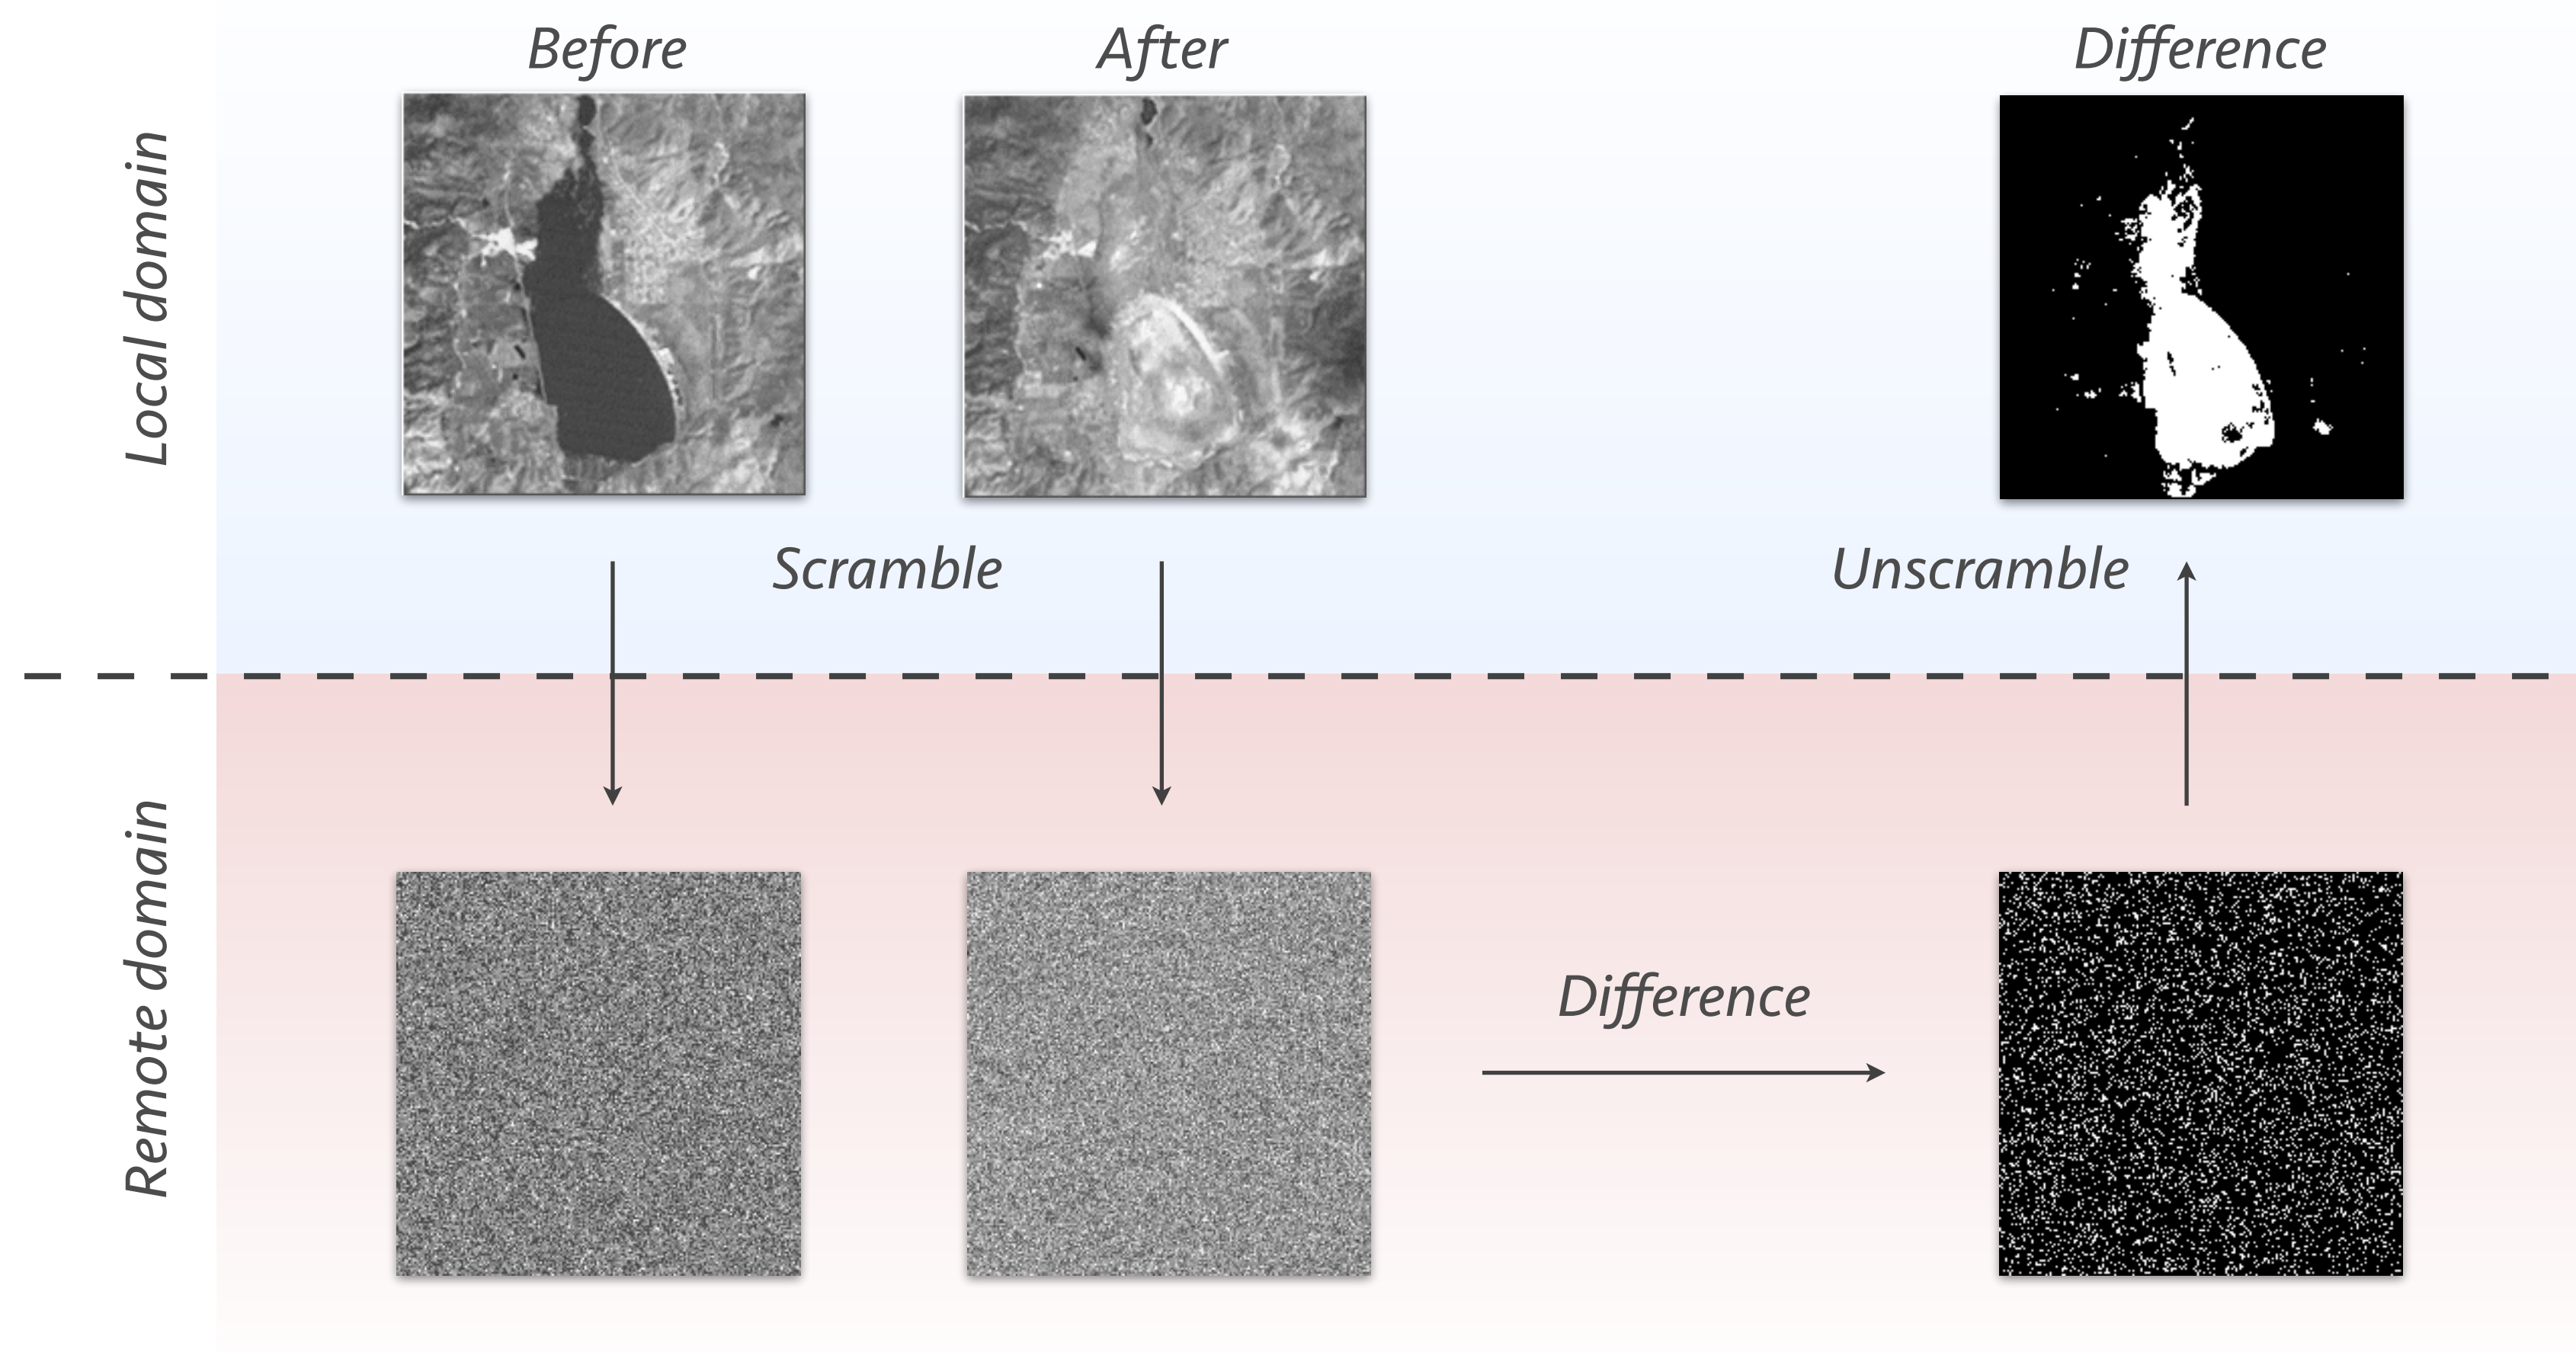
Difference



- How do we do this one?

Using noise to obfuscate the data

- Permute the pixels before sending them for processing!



Basic processing primitives

- The vast majority of MLSP uses only a few operations
 - Dot product, greater than, maxarg, ...
- If we define secure versions of these, then we can construct more complex MLSP operations
 - While maintaining data privacy

The setup

- We will use two collaborators: Alice and Bob
 - They are not malicious, but curious
 - i.e. Won't go out of their way to cheat, but are happy to inspect the data
 - There is no trusted 3rd party
 - i.e. they cannot trust a third person to mediate
- Both have private data that they do not want to share
 - But they need each other's data to perform a computation
- They do not have infinite computing resources
 - Otherwise they could easily cheat

Multiple scenario cases

- **Blind transaction**
 - Alice wants to protect her data; Bob cannot see the data, but can see the process output
 - Even if Bob or a third party hacker is malicious they can't deduce the data
- **Double-blind transaction**
 - Alice wants to protect her data and the identity of the results; Bob cannot see the data nor the process output
 - Even if Bob or a third party hacker is malicious he can't deduce anything
- **But there's more**
 - Alice is malicious and wants to send specific data to reverse engineer Bob's algorithm
 - Alice has data which needs processing, and Bob's result can be sent over to Carl who analyses blindly to sent to Darren who then ...

A simple exchange for now

- Alice has vector \mathbf{x} , Bob has vector \mathbf{w} and threshold θ
 - They do not want to share their data with each other
 - They want to compute whether $\mathbf{w}^\top \mathbf{x} > \theta$
- Why this operation?
 - It is the core operation for most classifiers (elaborations later)

A cryptography diversion

- The Paillier cryptosystem
 - Provides an encoding: $y = E[x]$ and its inverse: $x = E^{-1}[y]$
 - x is referred to as the *plaintext* and y as the *ciphertext*
- Has a hugely useful property
 - $E[x]E[y] = E[x+y]$
 - This is known as *homomorphic encryption*
 - An operation on ciphertext corresponds to an operation in plaintext
 - Many other cryptosystems with similar properties

Some details

- Encryption: $c = E[x] = g^x r^n \bmod n^2$, $x \in \mathbb{Z}_n$, $r \in \mathbb{Z}_n^*$
 - g is a random integer $\mathbb{Z}_{n^2}^*$, r is a random integer \mathbb{Z}_n^*
 - $n = p q$, where p, q are equal length primes
 - Public key: $\{n, g\}$
- Decryption: $x = E[c] = L(c^\lambda \bmod n^2)m \bmod n$
 - $\lambda = \text{lcd}(p-1, q-1)$, $m = L(g^\lambda \bmod n^2) - 1 \bmod n$, $L(x) = (x-1)/n$
 - Private key is: $\{\lambda, m\}$

Homomorphic properties

- Adding plaintexts via ciphertexts
 - $E^{-1}[E[x_1] E[x_2] \bmod n^2] = x_1 + x_2 \bmod n$
 - $E^{-1}[E[x_1] g^{x_2} \bmod n^2] = x_1 + x_2 \bmod n$
- Multiplying plaintexts via ciphertexts
 - $E^{-1}[E[x_1]^{x_2} \bmod n^2] = x_1 x_2 \bmod n$
 - $E^{-1}[E[x_2]^{x_1} \bmod n^2] = x_1 x_2 \bmod n$

Using Paillier for a secure inner product

- Desired operation:
 - Alice has vector $\mathbf{x} = \{x_1, x_2, \dots, x_N\}$
 - Bob has vector $\mathbf{w} = \{w_1, w_2, \dots, w_N\}$
 - Wanted outcome: Bob obtains $E[\mathbf{w}^\top \mathbf{x}]$
- In the end, Bob should have result but cannot see it

The Secure Inner Product (SIP) protocol

- Alice:
 - Generates an $E[]$ and $E^{-1}[]$, sends $E[]$ to Bob
 - Encrypts \mathbf{x} and sends all $E[x_i]$ to Bob
- Bob:
 - Computes homomorphic element-wise multiplication
 - $E[x_i]^{w_i} = E[x_i w_i]$
 - Bob computes homomorphic summation
 - $\prod E[x_i w_i] = E[\sum x_i w_i] = E[\mathbf{w}^\top \mathbf{x}]$

But we need some more work

- A potential leak of Bob's data
 - Bob has $E[\mathbf{w}^\top \mathbf{x}]$ and he can't see the result
 - That's a feature not a bug!
 - Alice can decrypt it, but then she will know \mathbf{w}
 - That would be a bug, not a feature!
- To protect Bob's data we need to "mask" the result

Masking

- In *masking/blinding* we obfuscate our data by adding random numbers (as we did with earlier examples)
- In our case we can perform *additive masking*
 - Bob can add a secret random number ρ to $\mathbf{w}^\top \mathbf{x}$
 - $E[\mathbf{w}^\top \mathbf{x}]E[\rho] = E[\mathbf{w}^\top \mathbf{x} + \rho]$

Inching towards a classifier

- Use the SIP protocol:
 - Bob masks the obtained dot product: $E[\mathbf{w}^\top \mathbf{x}]$
 - Bob sends the result $E[\mathbf{w}^\top \mathbf{x} + \rho]$ to Alice
- Outcome:
 - Alice can obtain $\mathbf{w}^\top \mathbf{x} + \rho$, but has not clue what \mathbf{w} is
 - Bob has no clue what \mathbf{x} is
- To classify we need to compare $\mathbf{w}^\top \mathbf{x} + \rho$ with $\theta + \rho$

Performing a secure comparison

- Yao's millionaires' problem
 - Alice and Bob want to compare their assets, but not to reveal them
- More formally
 - Alice has x dollars
 - Bob has y dollars
 - Alice and Bob only need to find out if $x > y$
 - But Alice can't learn y and Bob can't learn x
- In our case instead of dollars we compare $\mathbf{w}^\top \mathbf{x} + \rho$ with $\theta + \rho$

Representing the values to compare

- Alice & Bob decide on a range of numbers to compare
 - e.g. \$5M, \$10M, ..., \$50M, each value represented by $i = \{1, \dots, 10\}$
- Alice and Bob then have i_a and i_b dollars
 - They need to know if $i_a > i_b$
- A minor issue: This is a discrete and bound set
 - Bad for us since the dot product and threshold are real
 - But we can always quantize

Secure comparison algorithm

Alice's data is red
Bob's data is blue
Shared data is yellow

- Alice sends her public key to Bob
 - Bob can thus compute $E[]$
- Bob computes $c = E[x]$
 - x being a random integer of N bits
- Bob transmits to Alice $v = c + 1 - i_b$
 - i_b is Bob's representation of assets

← This ensures that Bob's number remains private

Secure comparison algorithm

Alice's data is red
Bob's data is blue
Shared data is yellow

- Alice generates a set of 10 numbers ← Because we are considering 10 possible values
 - $y_{i=1,\dots,10} = E^{-1}[v + i - 1]$
 - The number corresponding to y_{i_b} will be equal to x
- Alice generates a random prime p ← This ensures that we don't leak Alice's private key
 - And computes $z_i = y_i \bmod p$
 - p is $N/2$ -bits and must result in z_i 's, that are apart by at least 2
- Alice sends p and $u = \{z_1, \dots, z_{i_a-1}, z_{i_a}, 1 + z_{i_a+1}, \dots, 1 + z_{10}\}$

Secure comparison algorithm

Alice's data is red
Bob's data is blue
Shared data is yellow

- Bob computes $g = x \bmod p$
 - x was Bob's original random number, p is Alice's prime
- Bob compares g with the i_b 'th element of u
 - if $u_{i_b} = g$, then $i_a \geq i_b$
 - Otherwise $i_a < i_b$
- Neither party gets to share their original number!

Back to our original problem

- Alice has data \mathbf{x} , Bob has classifier $\{\mathbf{w}, \theta\}$
 - Step 1. Perform secure inner product
 - Bob obtains $E[\mathbf{w}^\top \mathbf{x}]$
 - Bob sends to Alice $E[\mathbf{w}^\top \mathbf{x}]E[\rho] = E[\mathbf{w}^\top \mathbf{x} + \rho]$
 - Step 2. Perform secure comparison
 - Alice compares $\mathbf{w}^\top \mathbf{x} + \rho$ with Bob's $\theta + \rho$, gets inequality result
- Alice keeps both data and classification outcome private
 - At the expense of extra overhead ...

So what we do with this?

- Linear classifiers!
 - These are essentially the previous formula!
- What about more complex classifiers?
 - e.g. a Gaussian likelihood classifier
 - This isn't a linear operation

Making the Gaussian a dot product

- Gaussian log likelihood:

$$\log P(\mathbf{x}; \mu, \Sigma) = -\frac{1}{2}(\mathbf{x} - \mu)^\top \Sigma^{-1}(\mathbf{x} - \mu) - \frac{d}{2} \log 2\pi - \frac{1}{2} \log |\Sigma|$$

- Equivalent to: $g(\mathbf{x}) = \mathbf{x}^\top \cdot \mathbf{W} \cdot \mathbf{x} + \mathbf{w}^\top \cdot \mathbf{x} + w$

- Where: $\mathbf{W} = -\frac{1}{2}\Sigma^{-1}, \quad \mathbf{w} = \Sigma^{-1} \cdot \mu, \quad w = -\frac{1}{2}\mu^\top \cdot \Sigma^{-1} \cdot \mu - \frac{1}{2} \log |\Sigma| - \frac{1}{2} \log 2\pi$

- Which can be simplified to: $g(\mathbf{x}) = \tilde{\mathbf{x}}^\top \cdot \tilde{\mathbf{W}} \cdot \tilde{\mathbf{x}}$

- Concatenate 1 to \mathbf{x} and include \mathbf{w}, w into \mathbf{W}

One more step

- But that's not a single product (we can do better)
- Instead we can do: $g(\mathbf{x}) = \tilde{\mathbf{W}} \cdot \tilde{\mathbf{x}}$

- with:

$$\tilde{\mathbf{x}} = [1, x_1, x_2, \dots, x_1 x_1, x_1 x_2, x_1 x_3, \dots, x_1 x_N, x_2 x_1, x_2 x_2, x_2 x_3, \dots, x_N x_N]$$

$$\tilde{\mathbf{W}} = \begin{bmatrix} w, & \Sigma^{-1} \cdot \mu, & -\frac{1}{2} \text{vec} \Sigma^{-1} \end{bmatrix}, \quad w = -\frac{1}{2} \mu^\top \Sigma^{-1} \mu - \frac{1}{2} \log |\Sigma| - \frac{1}{2} \log 2\pi$$

- Which only needs a simple secure product operation

Stringing protocols together

- It is often easier to modularize a process
 - Coming up with a secure HMM is hard
 - Coming up with a secure Gaussian, followed by a secure sum, followed by a secure transition regularizer, followed by ..., is easier
- To do so we can use *additive shares*
 - At each step the output is additively distributed between parties
 - E.g. for SIP: $z + v = \text{SIP}(x, y)$, Alice has x and z , Bob has y and v
 - We can keep processing keeping all intermediate results hidden

But there are some issues ...

- Data needs to be integer-valued
 - Not a huge problem, we can quantize data to desired accuracy
 - Can be a problem with, e.g. secure comparison though
- Encryption/decryption is computationally intensive
 - There is work on specialized hardware for this
- Are these worthwhile ideas?
 - Maybe later, email encryption was just as hopeless a while back

More reading material

- Yao's Millionaires' problem:
 - <http://research.cs.wisc.edu/areas/sec/Yao1982.pdf>
- Secure Multiparty Computations and Data Mining
 - <https://eprint.iacr.org/2008/197.pdf>
- Full-blown secure HMM implementation for speech:
 - <http://paris.cs.illinois.edu/pubs/smaragdis-tasl07-3.pdf>